

# Week\_2\_Data\_Ingestion

September 15, 2024

```
[1]: import zipfile
import os
import pandas as pd
import matplotlib.pyplot as plt
import h5py
import cv2
import numpy as np
```

## 0.1 Extract data from zipfile

```
[2]: import zipfile
import os

def extract_zip(zip_file_path: str, extract_to_directory: str):
    """
    Extracts a ZIP file to a specified directory if the ZIP file exists and the
    ↪files
    do not already exist.

    Parameters:
        zip_file_path (str): Path to the ZIP file.
        extract_to_directory (str): Directory where the ZIP file will be
    ↪extracted.
    """
    # Check if the ZIP file exists
    if not os.path.exists(zip_file_path):
        print(f"ZIP file {zip_file_path} does not exist. Skipping extraction.")
        return

    try:
        # Check if the extraction directory contains files
        if not os.path.exists(extract_to_directory) or not os.
    ↪listdir(extract_to_directory):
            with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
                zip_ref.extractall(extract_to_directory)
            print(f"Extracted {zip_file_path} to {extract_to_directory}")
        else:
```

```

        print(f"Files already exist in {extract_to_directory}. Skipping
↳extraction.")
    except zipfile.BadZipFile:
        print(f"Error: {zip_file_path} is not a valid ZIP file.")
    except Exception as e:
        print(f"An error occurred during extraction: {e}")

def remove_file(file_path: str):
    """
    Removes a file from the filesystem if it exists.

    Parameters:
        file_path (str): Path to the file to be removed.
    """
    if os.path.exists(file_path):
        os.remove(file_path)
        print(f"Removed {file_path}")
    else:
        print(f"File {file_path} does not exist.")

# Paths
zip_file_path = "../data/raw/train-metadata.zip"
extract_to_directory = "../data/raw/"

# Extract the ZIP file if it exists and the files are not already extracted
extract_zip(zip_file_path, extract_to_directory)

# Optionally, remove the ZIP file after extraction
remove_file(zip_file_path)

```

ZIP file ../data/raw/train-metadata.zip does not exist. Skipping extraction.  
File ../data/raw/train-metadata.zip does not exist.

## 0.2 Load Metadata

```
[3]: df_metadata = pd.read_csv("../data/raw/train-metadata.csv")
df_metadata.head(10)
```

/tmp/ipykernel\_1668636/1128991036.py:1: DtypeWarning: Columns (51,52) have mixed types. Specify dtype option on import or set low\_memory=False.

```
df_metadata = pd.read_csv("../data/raw/train-metadata.csv")
```

```
[3]:
```

	isic_id	target	patient_id	age_approx	sex	anatom_site_general	\
0	ISIC_0015670	0	IP_1235828	60.0	male	lower extremity	
1	ISIC_0015845	0	IP_8170065	60.0	male	head/neck	
2	ISIC_0015864	0	IP_6724798	60.0	male	posterior torso	
3	ISIC_0015902	0	IP_4111386	65.0	male	anterior torso	

4	ISIC_0024200	0	IP_8313778	55.0	male	anterior torso
5	ISIC_0035502	0	IP_3026693	75.0	female	head/neck
6	ISIC_0051648	0	IP_0218255	65.0	male	upper extremity
7	ISIC_0051665	0	IP_7734648	50.0	male	posterior torso
8	ISIC_0051710	0	IP_1307115	50.0	male	anterior torso
9	ISIC_0051758	0	IP_2180091	60.0	female	upper extremity

	clin_size_long_diam_mm	image_type	tbp_tile_type	tbp_lv_A	...	\
0	3.04	TBP tile: close-up	3D: white	20.244422	...	
1	1.10	TBP tile: close-up	3D: white	31.712570	...	
2	3.40	TBP tile: close-up	3D: XP	22.575830	...	
3	3.22	TBP tile: close-up	3D: XP	14.242329	...	
4	2.73	TBP tile: close-up	3D: white	24.725520	...	
5	2.54	TBP tile: close-up	3D: white	22.129183	...	
6	3.74	TBP tile: close-up	3D: XP	14.319188	...	
7	4.31	TBP tile: close-up	3D: XP	20.102610	...	
8	3.17	TBP tile: close-up	3D: white	14.166805	...	
9	3.97	TBP tile: close-up	3D: XP	16.251490	...	

	lesion_id	iddx_full	iddx_1	iddx_2	iddx_3	iddx_4	iddx_5	\
0	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
1	IL_6727506	Benign	Benign	NaN	NaN	NaN	NaN	
2	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
3	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
4	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
5	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
6	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
7	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
8	NaN	Benign	Benign	NaN	NaN	NaN	NaN	
9	NaN	Benign	Benign	NaN	NaN	NaN	NaN	

	mel_mitotic_index	mel_thick_mm	tbp_lv_dnn_lesion_confidence
0	NaN	NaN	97.517282
1	NaN	NaN	3.141455
2	NaN	NaN	99.804040
3	NaN	NaN	99.989998
4	NaN	NaN	70.442510
5	NaN	NaN	99.619603
6	NaN	NaN	99.918133
7	NaN	NaN	99.972390
8	NaN	NaN	99.818963
9	NaN	NaN	99.999690

[10 rows x 55 columns]

### 0.3 Dataset Stats

```
[4]: def df_stats(df: pd.DataFrame, include_all: bool = False):  
    """  
    Print statistics and null value counts for a pandas DataFrame.  
  
    Parameters:  
        df (pd.DataFrame): The DataFrame to analyze.  
        include_all (bool): If True, include all columns in the descriptive_  
↪statistics; otherwise, include only numeric columns.  
  
    Returns:  
        None  
    """  
    if df.empty:  
        print("The DataFrame is empty.")  
        return  
  
    # Print descriptive statistics  
    print("Descriptive Statistics:")  
    if include_all:  
        print(df.describe(include='all'))  
    else:  
        print(df.describe(include=[np.number]))  
    print("\n" + "-"*50 + "\n") # Separator for clarity  
  
    # Print the number of null values per column  
    print("Null Value Counts:")  
    print(df.isnull().sum())  
    print("\n" + "-"*50 + "\n") # Separator for clarity  
  
    # Additional information: Percentage of null values per column  
    print("Percentage of Null Values:")  
    print(df.isnull().mean() * 100)  
    print("\n" + "-"*50 + "\n") # Separator for clarity  
  
    # Number of rows and columns  
    print(f"Number of rows: {df.shape[0]}")  
    print(f"Number of columns: {df.shape[1]}")  
    print("\n" + "-"*50 + "\n") # Separator for clarity
```

```
[5]: df_stats(df_metadata)
```

Descriptive Statistics:

	target	age_approx	clin_size_long_diam_mm	tbp_lv_A \
count	401059.000000	398261.000000	401059.000000	401059.000000
mean	0.000980	58.012986	3.930827	19.974007
std	0.031288	13.596165	1.743068	3.999489

min	0.000000	5.000000	1.000000	-2.487115
25%	0.000000	50.000000	2.840000	17.330821
50%	0.000000	60.000000	3.370000	19.801910
75%	0.000000	70.000000	4.380000	22.304628
max	1.000000	85.000000	28.400000	48.189610

	tbp_lv_Aext	tbp_lv_B	tbp_lv_Bext	tbp_lv_C \
count	401059.000000	401059.000000	401059.000000	401059.000000
mean	14.919247	28.281706	26.913015	34.786341
std	3.529384	5.278676	4.482994	5.708469
min	-9.080269	-0.730989	9.237066	3.054228
25%	12.469740	24.704372	23.848125	31.003148
50%	14.713930	28.171570	26.701704	34.822580
75%	17.137175	31.637429	29.679913	38.430298
max	37.021680	54.306900	48.372700	58.765170

	tbp_lv_Cext	tbp_lv_H ...	tbp_lv_radial_color_std_max \
count	401059.000000	401059.000000 ...	401059.000000
mean	30.921279	54.653689 ...	1.016459
std	4.829345	5.520849 ...	0.734631
min	11.846520	-1.574164 ...	0.000000
25%	27.658285	51.566273 ...	0.563891
50%	30.804893	55.035632 ...	0.902281
75%	33.963868	58.298184 ...	1.334523
max	54.305290	105.875784 ...	11.491140

	tbp_lv_stdL	tbp_lv_stdLExt	tbp_lv_symm_2axis \
count	401059.000000	401059.000000	401059.000000
mean	2.715190	2.238605	0.306823
std	1.738165	0.623884	0.125038
min	0.268160	0.636247	0.052034
25%	1.456570	1.834745	0.211429
50%	2.186693	2.149758	0.282297
75%	3.474565	2.531443	0.382022
max	17.563650	25.534791	0.977055

	tbp_lv_symm_2axis_angle	tbp_lv_x	tbp_lv_y	tbp_lv_z \
count	401059.000000	401059.000000	401059.000000	401059.000000
mean	86.332073	-3.091862	1039.598221	55.823389
std	52.559511	197.257995	409.819653	87.968245
min	0.000000	-624.870728	-1052.134000	-291.890442
25%	40.000000	-147.022125	746.519673	-8.962647
50%	90.000000	-5.747253	1172.803000	67.957947
75%	130.000000	140.474835	1342.131540	126.611567
max	175.000000	614.471700	1887.766846	319.407000

	mel_thick_mm	tbp_lv_dnn_lesion_confidence
count	63.000000	4.010590e+05

mean	0.670952	9.716220e+01
std	0.792798	8.995782e+00
min	0.200000	1.261082e-16
25%	0.300000	9.966882e+01
50%	0.400000	9.999459e+01
75%	0.600000	9.999996e+01
max	5.000000	1.000000e+02

[8 rows x 37 columns]

-----

Null Value Counts:

isic_id	0
target	0
patient_id	0
age_approx	2798
sex	11517
anatom_site_general	5756
clin_size_long_diam_mm	0
image_type	0
tbp_tile_type	0
tbp_lv_A	0
tbp_lv_Aext	0
tbp_lv_B	0
tbp_lv_Bext	0
tbp_lv_C	0
tbp_lv_Cext	0
tbp_lv_H	0
tbp_lv_Hext	0
tbp_lv_L	0
tbp_lv_Lext	0
tbp_lv_areaMM2	0
tbp_lv_area_perim_ratio	0
tbp_lv_color_std_mean	0
tbp_lv_deltaA	0
tbp_lv_deltaB	0
tbp_lv_deltaL	0
tbp_lv_deltaLB	0
tbp_lv_deltaLBnorm	0
tbp_lv_eccentricity	0
tbp_lv_location	0
tbp_lv_location_simple	0
tbp_lv_minorAxisMM	0
tbp_lv_nevi_confidence	0
tbp_lv_norm_border	0
tbp_lv_norm_color	0
tbp_lv_perimeterMM	0

tbp_lv_radial_color_std_max	0
tbp_lv_stdL	0
tbp_lv_stdLExt	0
tbp_lv_symm_2axis	0
tbp_lv_symm_2axis_angle	0
tbp_lv_x	0
tbp_lv_y	0
tbp_lv_z	0
attribution	0
copyright_license	0
lesion_id	379001
iddx_full	0
iddx_1	0
iddx_2	399991
iddx_3	399994
iddx_4	400508
iddx_5	401058
mel_mitotic_index	401006
mel_thick_mm	400996
tbp_lv_dnn_lesion_confidence	0
dtype: int64	

-----

Percentage of Null Values:

isic_id	0.000000
target	0.000000
patient_id	0.000000
age_approx	0.697653
sex	2.871647
anatom_site_general	1.435200
clin_size_long_diam_mm	0.000000
image_type	0.000000
tbp_tile_type	0.000000
tbp_lv_A	0.000000
tbp_lv_Aext	0.000000
tbp_lv_B	0.000000
tbp_lv_Bext	0.000000
tbp_lv_C	0.000000
tbp_lv_Cext	0.000000
tbp_lv_H	0.000000
tbp_lv_Hext	0.000000
tbp_lv_L	0.000000
tbp_lv_Lext	0.000000
tbp_lv_areaMM2	0.000000
tbp_lv_area_perim_ratio	0.000000
tbp_lv_color_std_mean	0.000000
tbp_lv_deltaA	0.000000

tbp_lv_deltaB	0.000000
tbp_lv_deltaL	0.000000
tbp_lv_deltaLB	0.000000
tbp_lv_deltaLBnorm	0.000000
tbp_lv_eccentricity	0.000000
tbp_lv_location	0.000000
tbp_lv_location_simple	0.000000
tbp_lv_minorAxisMM	0.000000
tbp_lv_nevi_confidence	0.000000
tbp_lv_norm_border	0.000000
tbp_lv_norm_color	0.000000
tbp_lv_perimeterMM	0.000000
tbp_lv_radial_color_std_max	0.000000
tbp_lv_stdL	0.000000
tbp_lv_stdLExt	0.000000
tbp_lv_symm_2axis	0.000000
tbp_lv_symm_2axis_angle	0.000000
tbp_lv_x	0.000000
tbp_lv_y	0.000000
tbp_lv_z	0.000000
attribution	0.000000
copyright_license	0.000000
lesion_id	94.500061
ididx_full	0.000000
ididx_1	0.000000
ididx_2	99.733705
ididx_3	99.734453
ididx_4	99.862614
ididx_5	99.999751
mel_mitotic_index	99.986785
mel_thick_mm	99.984292
tbp_lv_dnn_lesion_confidence	0.000000

dtype: float64

-----

Number of rows: 401059  
Number of columns: 55

-----

## Load Image Byte String

```
[29]: import h5py

def load_image_from_hdf5(isic_id: str,
                        file_path: str = "../data/raw/train-image.hdf5",
```



```

        n_channels: int = 3):
    # Handle the case where the isic_id is passed incorrectly
    if not isic_id.lower().startswith("isic"):
        isic_id = f"ISIC_{int(str(isic_id).split('_', 1)[-1]):>07}"

    # Open the HDF5 file in read mode
    with h5py.File(file_path, 'r') as hf:

        # Retrieve the image data from the HDF5 dataset using the provided ISIC_ID
        ↪ID
        try:
            image_data = hf[isic_id][()]
        except KeyError:
            raise KeyError(f"ISIC ID {isic_id} not found in HDF5 file.")

        # Convert the binary data to a numpy array
        image_array = np.frombuffer(image_data, np.uint8)

        # Decode the image from the numpy array
        if n_channels == 3:
            # Load the image as a color image (BGR) and convert to RGB
            image = cv2.cvtColor(cv2.imdecode(image_array, cv2.IMREAD_COLOR), ↪
            ↪cv2.COLOR_BGR2RGB)
        else:
            # Load the image as a grayscale image
            image = cv2.imdecode(image_array, cv2.IMREAD_GRAYSCALE)

        # If the image failed to load for some reason (problems decoding) ...
        if image is None:
            raise ValueError(f"Could not decode image for ISIC ID: {isic_id}")

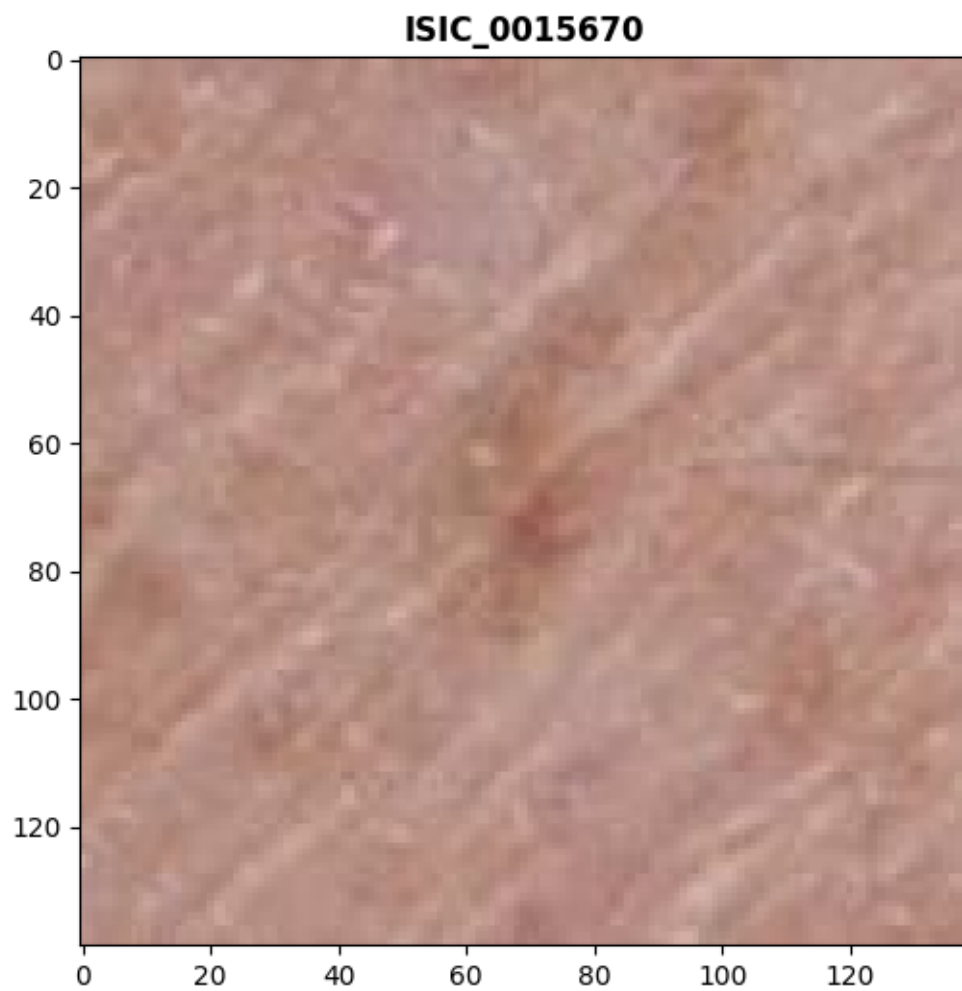
    return image

```

```

[30]: plt.figure(figsize=(6,6))
plt.title("ISIC_0015670", fontweight="bold")
plt.imshow(load_image_from_hdf5("ISIC_0015670"))
plt.show()

```



[ ]: