

# Week\_6\_First\_Modeling

October 14, 2024

```
[2]: # Standard Libraries
import io
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Deep Learning and PyTorch
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader

# Image Processing
from PIL import Image
from torchvision import transforms, models

# File Handling
import h5py

# Metrics and Evaluation
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, auc

# Progress Visualization
from tqdm import tqdm
```

## 0.1 Create Custom Dataset

```
[3]: class MultiInputDataset(Dataset):
    def __init__(self, hdf5_file, csv_file, transform=None):
        # Open the HDF5 file with error handling
        try:
            self.hdf5_file = h5py.File(hdf5_file, 'r') # Read-only mode
        except Exception as e:
            raise IOError(f"Could not open HDF5 file: {hdf5_file}. Error: {e}")

        # Read the CSV file containing image labels and additional features
```

```

try:
    self.labels_df = pd.read_csv(csv_file)
except Exception as e:
    raise IOError(f"Could not read CSV file: {csv_file}. Error: {e}")

# Ensure that all image IDs from the CSV are present in the HDF5 file
self.image_ids = self.labels_df['isic_id'].values
for image_id in self.image_ids:
    if str(image_id) not in self.hdf5_file.keys():
        raise ValueError(f"Image id {image_id} not found in HDF5 file.")

# Store any transformations to be applied to the images
self.transform = transform

def __len__(self):
    # Return the total number of samples in the dataset
    return len(self.labels_df)

def __getitem__(self, idx):
    # Get the image ID from the CSV file based on index
    image_id = str(self.labels_df.iloc[idx]['isic_id'])

    # Load the image data from the HDF5 file
    image_bytes = self.hdf5_file[image_id][()]

    # Convert the image bytes to a PIL Image
    image = Image.open(io.BytesIO(image_bytes))

    # Apply any specified transformations to the image
    if self.transform:
        image = self.transform(image)

    # Retrieve the label
    label = torch.tensor(self.labels_df.iloc[idx]['target'], dtype=torch.
↪long) # Adjust dtype if needed

    # Retrieve other features, excluding 'isic_id' and 'target'
    other_variables = self.labels_df.iloc[idx].drop(['isic_id', 'target']).
↪values.astype(float)

    # Convert other variables (metadata) to a tensor
    metadata_tensor = torch.tensor(other_variables, dtype=torch.float32)

    # Return the image, metadata, and label
    return image, metadata_tensor, label

```

```
[4]: # Define any necessary transformations for the image dataset

# Transformations for training set (with data augmentation)
train_transform = transforms.Compose([
    transforms.Resize((128, 128)), # Resize to 225x225
    transforms.RandomResizedCrop(128, scale=(0.8, 1.0)), # Random crop to
    ↪ 225x225 with scale
    transforms.RandomRotation(10), # Randomly rotate images by 10 degrees
    transforms.ToTensor(), # Convert image to PyTorch tensor
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# Transformations for validation/test set (no data augmentation)
normal_transform = transforms.Compose([
    transforms.Resize((128, 128)), # Ensure the image is resized to 225x225
    transforms.ToTensor(), # Convert image to PyTorch tensor
])
```

## 0.2 Train DataLoader

```
[5]: # Initialize the dataset
train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.hdf5',
    ↪ csv_file='../data/processed/processed-train-metadata1.csv',
    ↪ transform=normal_transform)
val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.hdf5',
    ↪ csv_file='../data/processed/processed-validation-metadata1.csv',
    ↪ transform=normal_transform)

# Create a DataLoader
train_dataloader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=64, shuffle=True)

[6]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

## 0.3 Model Building

```
[7]: class CustomImageFeatureCNN2(nn.Module):
    def __init__(self, feature_input_size, input_image_size=(128, 128)):
        super(CustomImageFeatureCNN2, self).__init__()

        # Image CNN with Batch Normalization
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3,
    ↪ padding=1)
        self.bn1 = nn.BatchNorm2d(32) # Batch normalization after conv1

        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
```

```

self.bn2 = nn.BatchNorm2d(64) # Batch normalization after conv2

self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
self.bn3 = nn.BatchNorm2d(128) # Batch normalization after conv3

self.pool = nn.MaxPool2d(kernel_size=2, stride=2) # 2x2 Max pooling

# Dynamically calculate the flattened size of the feature map
self.flattened_size = self._get_flattened_size(input_image_size)

# Fully connected layer after the CNN layers
self.fc_image = nn.Linear(self.flattened_size, 512)

# Fully connected layer for metadata (feature data)
self.fc_metadata = nn.Linear(feature_input_size, 128)

# Dropout layer to prevent overfitting
self.dropout = nn.Dropout(0.5) # 50% dropout

# Final fully connected layer for binary classification (combined image
↪+ feature input)
self.fc_combined = nn.Linear(512 + 128, 1) # Change 2 to 1 for binary
↪classification

def _get_flattened_size(self, input_image_size):
    # Forward pass a dummy image to get the size of the flattened features
    dummy_image = torch.zeros(1, 3, *input_image_size) # Batch size of 1,
↪3 channels (RGB), and input size
    x = self.pool(F.relu(self.bn1(self.conv1(dummy_image))))
    x = self.pool(F.relu(self.bn2(self.conv2(x))))
    x = self.pool(F.relu(self.bn3(self.conv3(x))))
    return x.view(-1).shape[0] # Flatten and return the size

def forward(self, image, metadata):
    # Forward pass for the image through the CNN
    x = self.pool(F.relu(self.bn1(self.conv1(image)))) # Conv layer 1 with
↪ReLU, BatchNorm, MaxPool
    x = self.pool(F.relu(self.bn2(self.conv2(x)))) # Conv layer 2 with
↪ReLU, BatchNorm, MaxPool
    x = self.pool(F.relu(self.bn3(self.conv3(x)))) # Conv layer 3 with
↪ReLU, BatchNorm, MaxPool

    # Flatten the feature map to feed into fully connected layer
    x = x.view(x.size(0), -1) # Flatten feature maps into a 1D vector
    image_features = F.relu(self.fc_image(x))

```

```

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            f"Batch sizes do not match! Image batch size: {image_features.
→shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
→dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

        # If you're using BCELoss, uncomment the next line to apply sigmoid
        output = torch.sigmoid(output)

        return output

```

## 0.4 Model Training

```

[8]: # Function to compute partial AUC-above-TPR
def score(solution: np.array, submission: np.array, min_tpr: float = 0.80) ->
→float:
    """
    Compute the partial AUC by focusing on a specific range of true positive
→rates (TPR).

    Args:
        solution (np.array): Ground truth binary labels.
        submission (np.array): Model predictions.
        min_tpr (float): Minimum true positive rate to calculate partial AUC.

    Returns:
        float: The calculated partial AUC.

    Raises:
        ValueError: If the min_tpr is not within a valid range.
    """
    # Rescale the target to handle sklearn limitations and flip the predictions
    v_gt = abs(solution - 1)
    v_pred = -1.0 * submission
    max_fpr = abs(1 - min_tpr)

```

```

# Compute ROC curve using sklearn
fpr, tpr, _ = roc_curve(v_gt, v_pred)
if max_fpr is None or max_fpr == 1:
    return auc(fpr, tpr)
if max_fpr <= 0 or max_fpr > 1:
    raise ValueError(f"Expected min_tpr in range [0, 1), got: {min_tpr}")

# Interpolate for partial AUC
stop = np.searchsorted(fpr, max_fpr, "right")
x_interp = [fpr[stop - 1], fpr[stop]]
y_interp = [tpr[stop - 1], tpr[stop]]
tpr = np.append(tpr[:stop], np.interp(max_fpr, x_interp, y_interp))
fpr = np.append(fpr[:stop], max_fpr)
partial_auc = auc(fpr, tpr)

return partial_auc

# Training and validation loop function
def train_and_validate(
    model: nn.Module,
    train_dataloader: torch.utils.data.DataLoader,
    val_dataloader: torch.utils.data.DataLoader,
    criterion: nn.Module,
    optimizer: torch.optim.Optimizer,
    epochs: int,
    device: torch.device,
    early_stopping_patience: int = 5,
    min_tpr: float = 0.80
) -> nn.Module:
    """
    Train and validate a PyTorch model with early stopping, AUROC, partial AUC,
    and error handling.

    Args:
        model (nn.Module): The model to be trained and validated.
        train_dataloader (torch.utils.data.DataLoader): Dataloader for training
        data.
        val_dataloader (torch.utils.data.DataLoader): Dataloader for validation
        data.
        criterion (nn.Module): Loss function.
        optimizer (torch.optim.Optimizer): Optimizer to update the model.
        epochs (int): Number of training epochs.
        device (torch.device): The device (CPU or GPU) to use.
        early_stopping_patience (int): Early stopping patience.
        min_tpr (float): The minimum true positive rate for calculating partial
        AUC.

```

```

Returns:
    nn.Module: The trained model.
"""
# Initialize tracking variables
best_val_loss = float('inf')
best_epoch = 0
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []
early_stopping_counter = 0

# Start the training and validation loop
for epoch in range(epochs):
    print(f'Epoch {epoch + 1}/{epochs}')

    # Training phase
    model.train()
    running_train_loss = 0.0
    correct_train = 0
    total_train = 0
    all_train_labels = []
    all_train_probs = []

    progress_bar = tqdm(train_dataloader, desc=f'Training Epoch {epoch + 1}')

    try:
        # Loop through the training batches
        for i, (image, metadata, labels) in enumerate(progress_bar):
            image, metadata, labels = image.to(device), metadata.
            to(device), labels.float().to(device)
            labels = labels.unsqueeze(1) # Adjust labels to have the right
            shape for binary classification

            optimizer.zero_grad()

            # Forward pass
            probs = model(image, metadata)

            if probs.shape != labels.shape:
                raise ValueError(f"Shape mismatch: Predictions shape {probs.
            shape} does not match labels shape {labels.shape}")

            # Calculate loss and backpropagate
            loss = criterion(probs, labels)
            loss.backward()

```

```

optimizer.step()

# Update running loss
running_train_loss += loss.item()

# Store labels and predictions for accuracy calculations
all_train_labels.extend(labels.cpu().detach().numpy())
all_train_probs.extend(probs.cpu().detach().numpy())

# Calculate binary predictions for training accuracy
predicted_train = (probs >= 0.5).float()
total_train += labels.size(0)
correct_train += (predicted_train == labels).sum().item()

# Update progress bar
progress_bar.set_postfix(train_loss=running_train_loss / (i + 1))

# Calculate training accuracy and loss
train_accuracy = 100 * correct_train / total_train
train_losses.append(running_train_loss / len(train_dataloader))
train_accuracies.append(train_accuracy)

except ValueError as ve:
    print(f"Error during training loop: {ve}")
    break

# Validation phase
model.eval()
running_val_loss = 0.0
correct = 0
total = 0
all_labels = []
all_probs = []

progress_bar = tqdm(val_dataloader, desc=f'Validating Epoch {epoch + 1}')

with torch.no_grad():
    try:
        # Loop through the validation batches
        for i, (images, metadata, labels) in enumerate(progress_bar):
            images, metadata, labels = images.to(device), metadata.
            labels = labels.float().to(device)
            labels = labels.unsqueeze(1)

            probs = model(images, metadata)

```



```

        loss = criterion(probs, labels)
        running_val_loss += loss.item()

        all_labels.extend(labels.cpu().detach().numpy())
        all_probs.extend(probs.cpu().detach().numpy())

        # Calculate binary predictions for validation accuracy
        predicted = (probs >= 0.5).float()
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        progress_bar.set_postfix(val_loss=running_val_loss / (i + 1))

    val_accuracy = 100 * correct / total
    val_loss = running_val_loss / len(val_dataloader)
    val_accurrencies.append(val_accuracy)
    val_losses.append(val_loss)

    # Calculate AUROC
    try:
        valid_auroc = roc_auc_score(all_labels, all_probs)
    except ValueError as ve:
        print(f"AUROC Calculation Error: {ve}")
        valid_auroc = 0.0

    # Calculate partial AUC-above-TPR
    try:
        partial_auroc = score(np.array(all_labels), np.
        array(all_probs), min_tpr=min_tpr)
    except ValueError as ve:
        print(f"Partial AUC Calculation Error: {ve}")
        partial_auroc = 0.0

    print(f'Epoch [{epoch + 1}/{epochs}], Train Loss: {train_losses[-1]:.4f}, Val Loss: {val_loss:.4f}, '
          f'Val Accuracy: {val_accuracy:.2f}%, Val AUROC: {valid_auroc:.4f}, Partial AUROC: {partial_auroc:.4f}')

    # Early stopping based on validation loss
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_epoch = epoch + 1
        early_stopping_counter = 0
        torch.save(model.state_dict(), 'best_model.pth')
    else:

```

```

        early_stopping_counter += 1

        if early_stopping_counter >= early_stopping_patience:
            print(f"Early stopping triggered at epoch {epoch + 1}")
            break

    except Exception as e:
        print(f"Error during validation loop: {e}")
        break

    print(f"Best Epoch: {best_epoch}, Best Validation Loss: {best_val_loss:.4f}")
    print('Training Complete')

    # Plot training and validation loss
    plt.figure(figsize=(10, 5))
    plt.plot(train_losses, label='Train Loss')
    plt.plot(val_losses, label='Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
    plt.show()

    # Plot training and validation accuracy
    plt.figure(figsize=(10, 5))
    plt.plot(train_accuracies, label='Train Accuracy')
    plt.plot(val_accuracies, label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.title('Training and Validation Accuracy')
    plt.legend()
    plt.show()

    # Generate classification report
    try:
        print("Classification Report:")
        print(classification_report(all_labels, (np.array(all_probs) >= 0.5).
        ↪astype(int), target_names=['Class 0', 'Class 1']))
    except Exception as e:
        print(f"Error generating classification report: {e}")

    return model

```

## 0.5 Model 1

```
[8]: model1 = CustomImageFeatureCNN2(feature_input_size=9) # Assuming 9 features
    ↪for metadata
model1.to(device)
# Initialize optimizer
optimizer = optim.Adam(model1.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20

[9]: train_and_validate(model1, train_dataloader, val_dataloader, criterion,
    ↪optimizer, epochs, device )
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [01:27<00:00, 2.64s/it,  
train\_loss=31.7]

Validating Epoch 1: 100%| | 24/24 [00:24<00:00, 1.02s/it,  
val\_loss=3.84]

Epoch [1/20], Train Loss: 31.7326, Val Loss: 3.8411, Val Accuracy: 96.04%, Val  
AUROC: 0.5000, Partial AUROC: 0.0200

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [01:20<00:00, 2.45s/it, train\_loss=33]

Validating Epoch 2: 100%| | 24/24 [00:25<00:00, 1.07s/it,  
val\_loss=4.01]

Epoch [2/20], Train Loss: 32.9726, Val Loss: 4.0075, Val Accuracy: 96.04%, Val  
AUROC: 0.5000, Partial AUROC: 0.0200

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [01:22<00:00, 2.50s/it,  
train\_loss=32.9]

Validating Epoch 3: 100%| | 24/24 [00:26<00:00, 1.09s/it,  
val\_loss=3.84]

Epoch [3/20], Train Loss: 32.9230, Val Loss: 3.8411, Val Accuracy: 96.04%, Val  
AUROC: 0.5000, Partial AUROC: 0.0200

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [01:24<00:00, 2.55s/it,  
train\_loss=32.8]

Validating Epoch 4: 100%| | 24/24 [00:26<00:00, 1.10s/it,  
val\_loss=4.17]

Epoch [4/20], Train Loss: 32.8238, Val Loss: 4.1739, Val Accuracy: 96.04%, Val  
AUROC: 0.5000, Partial AUROC: 0.0200


Epoch 5/20


```
Training Epoch 5: 100%|      | 33/33 [01:22<00:00,  2.49s/it,
train_loss=32.9]
Validating Epoch 5: 100%|      | 24/24 [00:24<00:00,  1.01s/it,
val_loss=3.84]

Epoch [5/20], Train Loss: 32.9230, Val Loss: 3.8411, Val Accuracy: 96.04%, Val
AUROC: 0.5000, Partial AUROC: 0.0200
Epoch 6/20

Training Epoch 6: 100%|      | 33/33 [01:21<00:00,  2.48s/it,
train_loss=32.9]
Validating Epoch 6: 100%|      | 24/24 [00:23<00:00,  1.00it/s,
val_loss=4.01]

Epoch [6/20], Train Loss: 32.8734, Val Loss: 4.0075, Val Accuracy: 96.04%, Val
AUROC: 0.5000, Partial AUROC: 0.0200
Early stopping triggered at epoch 6
Best Epoch: 1, Best Validation Loss: 3.8411
Training Complete
```





#### Classification Report:

	precision	recall	f1-score	support
Class 0	0.96	1.00	0.98	1431
Class 1	0.00	0.00	0.00	59
accuracy			0.96	1490
macro avg	0.48	0.50	0.49	1490
weighted avg	0.92	0.96	0.94	1490

```

/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```
[9]: CustomImageFeatureCNN2(
    (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc_image): Linear(in_features=32768, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)
```

## 0.6 Model 2

```
[10]: model2 = CustomImageFeatureCNN2(feature_input_size=9)  # Assuming 9 features
      ↪for metadata
model2.to(device)
      # Initialize optimizer
optimizer = optim.SGD(model2.parameters(), lr=0.001)
      # Define the loss function with the class weights
criterion = nn.BCELoss()  # Binary classification loss
      # Set the number of epochs
epochs = 20
```

```
[11]: train_and_validate(model2,train_dataloader, val_dataloader, criterion,
      ↪optimizer, epochs, device )
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [01:18<00:00, 2.39s/it,  
train\_loss=0.615]

Validating Epoch 1: 100%| | 24/24 [00:25<00:00, 1.08s/it,  
val\_loss=0.619]

Epoch [1/20], Train Loss: 0.6151, Val Loss: 0.6193, Val Accuracy: 92.35%, Val  
AUROC: 0.6333, Partial AUROC: 0.0372

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [01:15<00:00, 2.28s/it,  
train\_loss=0.569]

Validating Epoch 2: 100%| | 24/24 [00:25<00:00, 1.05s/it,  
val\_loss=0.427]

Epoch [2/20], Train Loss: 0.5693, Val Loss: 0.4272, Val Accuracy: 95.44%, Val AUROC: 0.7560, Partial AUROC: 0.0553

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [01:18<00:00, 2.37s/it, train\_loss=0.541]

Validating Epoch 3: 100%| | 24/24 [00:25<00:00, 1.05s/it, val\_loss=0.336]

Epoch [3/20], Train Loss: 0.5412, Val Loss: 0.3361, Val Accuracy: 95.57%, Val AUROC: 0.7691, Partial AUROC: 0.0580

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [01:14<00:00, 2.26s/it, train\_loss=0.522]

Validating Epoch 4: 100%| | 24/24 [00:25<00:00, 1.07s/it, val\_loss=0.404]

Epoch [4/20], Train Loss: 0.5217, Val Loss: 0.4044, Val Accuracy: 89.53%, Val AUROC: 0.7813, Partial AUROC: 0.0611

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [01:17<00:00, 2.34s/it, train\_loss=0.5]

Validating Epoch 5: 100%| | 24/24 [00:25<00:00, 1.08s/it, val\_loss=0.325]

Epoch [5/20], Train Loss: 0.5001, Val Loss: 0.3252, Val Accuracy: 92.95%, Val AUROC: 0.7870, Partial AUROC: 0.0620

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [01:16<00:00, 2.31s/it, train\_loss=0.491]

Validating Epoch 6: 100%| | 24/24 [00:25<00:00, 1.07s/it, val\_loss=0.382]

Epoch [6/20], Train Loss: 0.4913, Val Loss: 0.3820, Val Accuracy: 89.33%, Val AUROC: 0.8013, Partial AUROC: 0.0698

Epoch 7/20

Training Epoch 7: 100%| | 33/33 [01:17<00:00, 2.34s/it, train\_loss=0.471]

Validating Epoch 7: 100%| | 24/24 [00:27<00:00, 1.16s/it, val\_loss=0.377]

Epoch [7/20], Train Loss: 0.4713, Val Loss: 0.3771, Val Accuracy: 88.26%, Val AUROC: 0.8055, Partial AUROC: 0.0716

Epoch 8/20

Training Epoch 8: 100%| | 33/33 [01:23<00:00, 2.52s/it, train\_loss=0.456]

Validating Epoch 8: 100%| | 24/24 [00:24<00:00, 1.03s/it, val\_loss=0.298]

Epoch [8/20], Train Loss: 0.4565, Val Loss: 0.2976, Val Accuracy: 91.81%, Val AUROC: 0.8065, Partial AUROC: 0.0696

Epoch 9/20

Training Epoch 9: 100%| | 33/33 [01:18<00:00, 2.37s/it, train\_loss=0.452]

Validating Epoch 9: 100%| | 24/24 [00:23<00:00, 1.01it/s, val\_loss=0.27]

Epoch [9/20], Train Loss: 0.4515, Val Loss: 0.2701, Val Accuracy: 93.15%, Val AUROC: 0.8119, Partial AUROC: 0.0729

Epoch 10/20

Training Epoch 10: 100%| | 33/33 [01:14<00:00, 2.25s/it, train\_loss=0.44]

Validating Epoch 10: 100%| | 24/24 [00:24<00:00, 1.03s/it, val\_loss=0.339]

Epoch [10/20], Train Loss: 0.4395, Val Loss: 0.3387, Val Accuracy: 89.60%, Val AUROC: 0.8151, Partial AUROC: 0.0739

Epoch 11/20

Training Epoch 11: 100%| | 33/33 [01:17<00:00, 2.35s/it, train\_loss=0.434]

Validating Epoch 11: 100%| | 24/24 [00:26<00:00, 1.12s/it, val\_loss=0.301]

Epoch [11/20], Train Loss: 0.4335, Val Loss: 0.3011, Val Accuracy: 90.87%, Val AUROC: 0.8142, Partial AUROC: 0.0716

Epoch 12/20

Training Epoch 12: 100%| | 33/33 [01:16<00:00, 2.31s/it, train\_loss=0.422]

Validating Epoch 12: 100%| | 24/24 [00:26<00:00, 1.12s/it, val\_loss=0.268]

Epoch [12/20], Train Loss: 0.4217, Val Loss: 0.2682, Val Accuracy: 93.02%, Val AUROC: 0.8211, Partial AUROC: 0.0758

Epoch 13/20

Training Epoch 13: 100%| | 33/33 [01:15<00:00, 2.29s/it, train\_loss=0.41]

Validating Epoch 13: 100%| | 24/24 [00:24<00:00, 1.02s/it, val\_loss=0.225]

Epoch [13/20], Train Loss: 0.4102, Val Loss: 0.2249, Val Accuracy: 94.70%, Val AUROC: 0.8217, Partial AUROC: 0.0751

Epoch 14/20

Training Epoch 14: 100%| | 33/33 [01:15<00:00, 2.29s/it, train\_loss=0.407]

Validating Epoch 14: 100%| | 24/24 [00:22<00:00, 1.08it/s, val\_loss=0.301]



Epoch [14/20], Train Loss: 0.4071, Val Loss: 0.3014, Val Accuracy: 90.87%, Val AUROC: 0.8245, Partial AUROC: 0.0768

Epoch 15/20

Training Epoch 15: 100%| | 33/33 [01:16<00:00, 2.31s/it, train\_loss=0.399]

Validating Epoch 15: 100%| | 24/24 [00:19<00:00, 1.23it/s, val\_loss=0.274]

Epoch [15/20], Train Loss: 0.3994, Val Loss: 0.2740, Val Accuracy: 92.15%, Val AUROC: 0.8313, Partial AUROC: 0.0818

Epoch 16/20

Training Epoch 16: 100%| | 33/33 [01:12<00:00, 2.19s/it, train\_loss=0.394]

Validating Epoch 16: 100%| | 24/24 [00:18<00:00, 1.27it/s, val\_loss=0.253]

Epoch [16/20], Train Loss: 0.3945, Val Loss: 0.2530, Val Accuracy: 93.62%, Val AUROC: 0.8298, Partial AUROC: 0.0800

Epoch 17/20

Training Epoch 17: 100%| | 33/33 [01:12<00:00, 2.19s/it, train\_loss=0.385]

Validating Epoch 17: 100%| | 24/24 [00:17<00:00, 1.35it/s, val\_loss=0.292]

Epoch [17/20], Train Loss: 0.3845, Val Loss: 0.2920, Val Accuracy: 90.81%, Val AUROC: 0.8347, Partial AUROC: 0.0837

Epoch 18/20

Training Epoch 18: 100%| | 33/33 [01:19<00:00, 2.41s/it, train\_loss=0.386]



Validating Epoch 18: 100%| | 24/24 [00:27<00:00, 1.13s/it, val\_loss=0.259]

Epoch [18/20], Train Loss: 0.3856, Val Loss: 0.2585, Val Accuracy: 92.68%, Val AUROC: 0.8346, Partial AUROC: 0.0821

Early stopping triggered at epoch 18

Best Epoch: 13, Best Validation Loss: 0.2249

Training Complete



Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.94	0.96	1431
Class 1	0.28	0.56	0.38	59
accuracy	0.93			1490

macro avg	0.63	0.75	0.67	1490
weighted avg	0.95	0.93	0.94	1490

```
[11]: CustomImageFeatureCNN2(
    (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc_image): Linear(in_features=32768, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)
```

## 0.7 Model 3

```
[9]: model3 = CustomImageFeatureCNN2(feature_input_size=9) # Assuming 9 features
    ↪for metadata
model3.to(device)
# Initialize optimizer
optimizer = optim.SGD(model3.parameters(), lr=0.0001, weight_decay=1e-4)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_size = 32
```

```
[10]: train_dataloader = DataLoader(train_dataset, batch_size=batch_size,
    ↪shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_size, shuffle=True)
```

```
[11]: train_and_validate(model3, train_dataloader, val_dataloader, criterion,
    ↪optimizer, epochs, device)
```

Epoch 1/20

Training Epoch 1: 100%| | 66/66 [01:09<00:00, 1.05s/it,  
train\_loss=0.631]

Validating Epoch 1: 100%| | 47/47 [00:20<00:00, 2.34it/s,  
val\_loss=0.429]

Epoch [1/20], Train Loss: 0.6314, Val Loss: 0.4291, Val Accuracy: 96.04%, Val AUROC: 0.5977, Partial AUROC: 0.0416

Epoch 2/20

Training Epoch 2: 100%| | 66/66 [01:10<00:00, 1.08s/it, train\_loss=0.615]

Validating Epoch 2: 100%| | 47/47 [00:18<00:00, 2.60it/s, val\_loss=0.441]

Epoch [2/20], Train Loss: 0.6152, Val Loss: 0.4406, Val Accuracy: 95.84%, Val AUROC: 0.6546, Partial AUROC: 0.0463

Epoch 3/20

Training Epoch 3: 100%| | 66/66 [01:12<00:00, 1.09s/it, train\_loss=0.608]

Validating Epoch 3: 100%| | 47/47 [00:17<00:00, 2.66it/s, val\_loss=0.416]

Epoch [3/20], Train Loss: 0.6081, Val Loss: 0.4158, Val Accuracy: 95.84%, Val AUROC: 0.6802, Partial AUROC: 0.0496

Epoch 4/20

Training Epoch 4: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.597]

Validating Epoch 4: 100%| | 47/47 [00:18<00:00, 2.54it/s, val\_loss=0.395]

Epoch [4/20], Train Loss: 0.5972, Val Loss: 0.3950, Val Accuracy: 95.77%, Val AUROC: 0.6962, Partial AUROC: 0.0520

Epoch 5/20

Training Epoch 5: 100%| | 66/66 [01:13<00:00, 1.12s/it, train\_loss=0.59]

Validating Epoch 5: 100%| | 47/47 [00:18<00:00, 2.58it/s, val\_loss=0.413]

Epoch [5/20], Train Loss: 0.5901, Val Loss: 0.4128, Val Accuracy: 95.84%, Val AUROC: 0.7199, Partial AUROC: 0.0534

Epoch 6/20

Training Epoch 6: 100%| | 66/66 [01:08<00:00, 1.04s/it, train\_loss=0.579]

Validating Epoch 6: 100%| | 47/47 [00:17<00:00, 2.64it/s, val\_loss=0.381]

Epoch [6/20], Train Loss: 0.5791, Val Loss: 0.3808, Val Accuracy: 95.91%, Val AUROC: 0.7299, Partial AUROC: 0.0544

Epoch 7/20

Training Epoch 7: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.572]

Validating Epoch 7: 100%| | 47/47 [00:18<00:00, 2.60it/s, val\_loss=0.415]

Epoch [7/20], Train Loss: 0.5719, Val Loss: 0.4152, Val Accuracy: 95.64%, Val AUROC: 0.7443, Partial AUROC: 0.0548

Epoch 8/20

Training Epoch 8: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.566]

Validating Epoch 8: 100%| | 47/47 [00:24<00:00, 1.94it/s, val\_loss=0.379]

Epoch [8/20], Train Loss: 0.5664, Val Loss: 0.3788, Val Accuracy: 95.91%, Val AUROC: 0.7505, Partial AUROC: 0.0555

Epoch 9/20

Training Epoch 9: 100%| | 66/66 [01:11<00:00, 1.09s/it, train\_loss=0.56]

Validating Epoch 9: 100%| | 47/47 [00:17<00:00, 2.75it/s, val\_loss=0.387]

Epoch [9/20], Train Loss: 0.5603, Val Loss: 0.3869, Val Accuracy: 95.77%, Val AUROC: 0.7552, Partial AUROC: 0.0564

Epoch 10/20

Training Epoch 10: 100%| | 66/66 [01:08<00:00, 1.04s/it, train\_loss=0.554]

Validating Epoch 10: 100%| | 47/47 [00:24<00:00, 1.89it/s, val\_loss=0.381]

Epoch [10/20], Train Loss: 0.5535, Val Loss: 0.3810, Val Accuracy: 95.50%, Val AUROC: 0.7604, Partial AUROC: 0.0574

Epoch 11/20

Training Epoch 11: 100%| | 66/66 [01:10<00:00, 1.06s/it, train\_loss=0.549]

Validating Epoch 11: 100%| | 47/47 [00:17<00:00, 2.64it/s, val\_loss=0.383]

Epoch [11/20], Train Loss: 0.5493, Val Loss: 0.3833, Val Accuracy: 95.17%, Val AUROC: 0.7665, Partial AUROC: 0.0574

Epoch 12/20

Training Epoch 12: 100%| | 66/66 [01:14<00:00, 1.13s/it, train\_loss=0.541]

Validating Epoch 12: 100%| | 47/47 [00:18<00:00, 2.51it/s, val\_loss=0.362]

Epoch [12/20], Train Loss: 0.5406, Val Loss: 0.3625, Val Accuracy: 95.30%, Val AUROC: 0.7689, Partial AUROC: 0.0579

Epoch 13/20

Training Epoch 13: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.541]

Validating Epoch 13: 100%| | 47/47 [00:18<00:00, 2.59it/s, val\_loss=0.353]

Epoch [13/20], Train Loss: 0.5414, Val Loss: 0.3531, Val Accuracy: 95.84%, Val AUROC: 0.7737, Partial AUROC: 0.0586

Epoch 14/20

Training Epoch 14: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.538]

Validating Epoch 14: 100%| | 47/47 [00:18<00:00, 2.56it/s, val\_loss=0.367]

Epoch [14/20], Train Loss: 0.5377, Val Loss: 0.3667, Val Accuracy: 95.17%, Val AUROC: 0.7777, Partial AUROC: 0.0588

Epoch 15/20

Training Epoch 15: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.533]

Validating Epoch 15: 100%| | 47/47 [00:18<00:00, 2.60it/s, val\_loss=0.346]

Epoch [15/20], Train Loss: 0.5327, Val Loss: 0.3462, Val Accuracy: 95.50%, Val AUROC: 0.7787, Partial AUROC: 0.0586

Epoch 16/20

Training Epoch 16: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.533]

Validating Epoch 16: 100%| | 47/47 [00:17<00:00, 2.62it/s, val\_loss=0.39]

Epoch [16/20], Train Loss: 0.5332, Val Loss: 0.3904, Val Accuracy: 94.23%, Val AUROC: 0.7843, Partial AUROC: 0.0600

Epoch 17/20

Training Epoch 17: 100%| | 66/66 [01:11<00:00, 1.09s/it, train\_loss=0.522]

Validating Epoch 17: 100%| | 47/47 [00:18<00:00, 2.60it/s, val\_loss=0.374]

Epoch [17/20], Train Loss: 0.5222, Val Loss: 0.3737, Val Accuracy: 94.16%, Val AUROC: 0.7845, Partial AUROC: 0.0596

Epoch 18/20

Training Epoch 18: 100%| | 66/66 [01:11<00:00, 1.09s/it, train\_loss=0.518]

Validating Epoch 18: 100%| | 47/47 [00:20<00:00, 2.28it/s, val\_loss=0.346]

Epoch [18/20], Train Loss: 0.5176, Val Loss: 0.3460, Val Accuracy: 94.83%, Val AUROC: 0.7861, Partial AUROC: 0.0597

Epoch 19/20

Training Epoch 19: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.513]

Validating Epoch 19: 100%| | 47/47 [00:17<00:00, 2.62it/s, val\_loss=0.348]

Epoch [19/20], Train Loss: 0.5134, Val Loss: 0.3483, Val Accuracy: 94.43%, Val AUROC: 0.7861, Partial AUROC: 0.0602

Epoch 20/20


Training Epoch 20: 100%| | 66/66 [01:12<00:00, 1.10s/it, train\_loss=0.511]


Validating Epoch 20: 100%| | 47/47 [00:19<00:00, 2.47it/s, val\_loss=0.367]

Epoch [20/20], Train Loss: 0.5113, Val Loss: 0.3666, Val Accuracy: 93.83%, Val AUROC: 0.7901, Partial AUROC: 0.0610

Best Epoch: 18, Best Validation Loss: 0.3460

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.97	0.96	0.97	1431
Class 1	0.29	0.37	0.32	59
accuracy			0.94	1490
macro avg	0.63	0.67	0.65	1490
weighted avg	0.95	0.94	0.94	1490

```
[11]: CustomImageFeatureCNN2(
    (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc_image): Linear(in_features=32768, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
```



)

[ ]: