

# Week\_7\_Second\_Modeling-Copy1

October 22, 2024

```
[1]: # Standard Libraries
import io
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Deep Learning and PyTorch
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import models

# Image Processing
from PIL import Image
from torchvision import transforms, models

# File Handling
import h5py

# Metrics and Evaluation
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, auc

# Progress Visualization
from tqdm import tqdm
```

## 0.1 Create Custom Dataset

```
[2]: class MultiInputDataset(Dataset):
    def __init__(self, hdf5_file, csv_file, transform=None):
        # Open the HDF5 file with error handling
        try:
            self.hdf5_file = h5py.File(hdf5_file, 'r') # Read-only mode
        except Exception as e:
            raise IOError(f"Could not open HDF5 file: {hdf5_file}. Error: {e}")

        # Read the CSV file containing image labels and additional features
```

```

try:
    self.labels_df = pd.read_csv(csv_file)
except Exception as e:
    raise IOError(f"Could not read CSV file: {csv_file}. Error: {e}")

# Ensure that all image IDs from the CSV are present in the HDF5 file
self.image_ids = self.labels_df['isic_id'].values
for image_id in self.image_ids:
    if str(image_id) not in self.hdf5_file.keys():
        raise ValueError(f"Image id {image_id} not found in HDF5 file.")

# Store any transformations to be applied to the images
self.transform = transform

def __len__(self):
    # Return the total number of samples in the dataset
    return len(self.labels_df)

def __getitem__(self, idx):
    # Get the image ID from the CSV file based on index
    image_id = str(self.labels_df.iloc[idx]['isic_id'])

    # Load the image data from the HDF5 file
    image_bytes = self.hdf5_file[image_id][()]

    # Convert the image bytes to a PIL Image
    image = Image.open(io.BytesIO(image_bytes))

    # Apply any specified transformations to the image
    if self.transform:
        image = self.transform(image)

    # Retrieve the label
    label = torch.tensor(self.labels_df.iloc[idx]['target'], dtype=torch.
↪long) # Adjust dtype if needed

    # Retrieve other features, excluding 'isic_id' and 'target'
    other_variables = self.labels_df.iloc[idx].drop(['isic_id', 'target']).
↪values.astype(float)

    # Convert other variables (metadata) to a tensor
    metadata_tensor = torch.tensor(other_variables, dtype=torch.float32)

    # Return the image, metadata, and label
    return image, metadata_tensor, label

```

```
[3]: # Define any necessary transformations for the image dataset

# Transformations for training set (with data augmentation)
train_transform = transforms.Compose([
    transforms.Resize((225, 225)), # Resize to 225x225
    transforms.RandomResizedCrop(128, scale=(0.8, 1.0)), # Random crop to
    ↪225x225 with scale
    transforms.RandomRotation(10), # Randomly rotate images by 10 degrees
    transforms.ToTensor(), # Convert image to PyTorch tensor
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# Transformations for validation/test set (no data augmentation)
normal_transform = transforms.Compose([
    transforms.Resize((225, 225)), # Ensure the image is resized to 225x225
    transforms.ToTensor(), # Convert image to PyTorch tensor
])
```

## 0.2 Train DataLoader

```
[4]: # Initialize the dataset
train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.hdf5', ↪
    ↪csv_file='../data/processed/processed-train-metadata1.csv', ↪
    ↪transform=normal_transform)
val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.hdf5', ↪
    ↪csv_file='../data/processed/processed-validation-metadata1.csv', ↪
    ↪transform=normal_transform)

# Create a DataLoader
train_dataloader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=64, shuffle=True)

[5]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

## 0.3 Model Building

```
[6]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models

class CustomImageFeatureResNet(nn.Module):
    def __init__(self, feature_input_size, pretrained=True):
        super(CustomImageFeatureResNet, self).__init__()
```

```

    # Load a pretrained ResNet model for image feature extraction (ResNet18,
    ↪ in this case)
    resnet = models.resnet18(pretrained=pretrained) # Change to resnet50,
    ↪ resnet101 as needed
    self.resnet = nn.Sequential(*list(resnet.children())[:-1]) # Remove
    ↪ the final classification layer

    # The output of ResNet18's last conv layer is 512-dimensional (for
    ↪ ResNet50, it would be 2048)
    self.fc_image = nn.Linear(resnet.fc.in_features, 512) # Adjust if
    ↪ using ResNet50

    # Fully connected layer for metadata (feature data)
    self.fc_metadata = nn.Linear(feature_input_size, 128)

    # Dropout layer to prevent overfitting
    self.dropout = nn.Dropout(0.5) # 50% dropout

    # Final fully connected layer for binary classification (combined image
    ↪ + feature input)
    self.fc_combined = nn.Linear(512 + 128, 1) # For binary classification

    def forward(self, image, metadata):
        # Forward pass for the image through the ResNet (without the final
        ↪ classification layer)
        x = self.resnet(image) # ResNet feature extraction
        x = x.view(x.size(0), -1) # Flatten the ResNet output
        image_features = F.relu(self.fc_image(x))

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            f"Batch sizes do not match! Image batch size: {image_features.
        ↪ shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
        ↪ dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

        # If you're using BCELoss, uncomment the next line to apply sigmoid

```

```

output = torch.sigmoid(output)

return output

```

## 0.4 Model Training

```

[7]: # Function to compute partial AUC-above-TPR
def score(solution: np.array, submission: np.array, min_tpr: float = 0.80) -> float:
    """
    Compute the partial AUC by focusing on a specific range of true positive
    rates (TPR).

    Args:
        solution (np.array): Ground truth binary labels.
        submission (np.array): Model predictions.
        min_tpr (float): Minimum true positive rate to calculate partial AUC.

    Returns:
        float: The calculated partial AUC.

    Raises:
        ValueError: If the min_tpr is not within a valid range.
    """
    # Rescale the target to handle sklearn limitations and flip the predictions
    v_gt = abs(solution - 1)
    v_pred = -1.0 * submission
    max_fpr = abs(1 - min_tpr)

    # Compute ROC curve using sklearn
    fpr, tpr, _ = roc_curve(v_gt, v_pred)
    if max_fpr is None or max_fpr == 1:
        return auc(fpr, tpr)
    if max_fpr <= 0 or max_fpr > 1:
        raise ValueError(f"Expected min_tpr in range [0, 1), got: {min_tpr}")

    # Interpolate for partial AUC
    stop = np.searchsorted(fpr, max_fpr, "right")
    x_interp = [fpr[stop - 1], fpr[stop]]
    y_interp = [tpr[stop - 1], tpr[stop]]
    tpr = np.append(tpr[:stop], np.interp(max_fpr, x_interp, y_interp))
    fpr = np.append(fpr[:stop], max_fpr)
    partial_auc = auc(fpr, tpr)

    return partial_auc

```

```

# Training and validation loop function
def train_and_validate(
    model: nn.Module,
    train_dataloader: torch.utils.data.DataLoader,
    val_dataloader: torch.utils.data.DataLoader,
    criterion: nn.Module,
    optimizer: torch.optim.Optimizer,
    epochs: int,
    device: torch.device,
    early_stopping_patience: int = 5,
    min_tpr: float = 0.80
) -> nn.Module:
    """
    Train and validate a PyTorch model with early stopping, AUROC, partial AUC,
    and error handling.

    Args:
        model (nn.Module): The model to be trained and validated.
        train_dataloader (torch.utils.data.DataLoader): Dataloader for training
        data.
        val_dataloader (torch.utils.data.DataLoader): Dataloader for validation
        data.
        criterion (nn.Module): Loss function.
        optimizer (torch.optim.Optimizer): Optimizer to update the model.
        epochs (int): Number of training epochs.
        device (torch.device): The device (CPU or GPU) to use.
        early_stopping_patience (int): Early stopping patience.
        min_tpr (float): The minimum true positive rate for calculating partial
        AUC.

    Returns:
        nn.Module: The trained model.
    """
    # Initialize tracking variables
    best_val_loss = float('inf')
    best_epoch = 0
    train_losses = []
    val_losses = []
    train_accuracies = []
    val_accuracies = []
    early_stopping_counter = 0

    # Start the training and validation loop
    for epoch in range(epochs):
        print(f'Epoch {epoch + 1}/{epochs}')

        # Training phase

```

```

model.train()
running_train_loss = 0.0
correct_train = 0
total_train = 0
all_train_labels = []
all_train_probs = []

progress_bar = tqdm(train_dataloader, desc=f'Training Epoch {epoch + 1}')

try:
    # Loop through the training batches
    for i, (image, metadata, labels) in enumerate(progress_bar):
        image, metadata, labels = image.to(device), metadata.
        to(device), labels.float().to(device)
        labels = labels.unsqueeze(1) # Adjust labels to have the right
        shape for binary classification

        optimizer.zero_grad()

        # Forward pass
        probs = model(image, metadata)

        if probs.shape != labels.shape:
            raise ValueError(f"Shape mismatch: Predictions shape {probs.
            shape} does not match labels shape {labels.shape}")

        # Calculate loss and backpropagate
        loss = criterion(probs, labels)
        loss.backward()
        optimizer.step()

        # Update running loss
        running_train_loss += loss.item()

        # Store labels and predictions for accuracy calculations
        all_train_labels.extend(labels.cpu().detach().numpy())
        all_train_probs.extend(probs.cpu().detach().numpy())

        # Calculate binary predictions for training accuracy
        predicted_train = (probs >= 0.5).float()
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

        # Update progress bar
        progress_bar.set_postfix(train_loss=running_train_loss / (i + 1))

```

```

        # Calculate training accuracy and loss
        train_accuracy = 100 * correct_train / total_train
        train_losses.append(running_train_loss / len(train_dataloader))
        train_accuracies.append(train_accuracy)

    except ValueError as ve:
        print(f"Error during training loop: {ve}")
        break

    # Validation phase
    model.eval()
    running_val_loss = 0.0
    correct = 0
    total = 0
    all_labels = []
    all_probs = []

    progress_bar = tqdm(val_dataloader, desc=f'Validating Epoch {epoch + 1}')

    with torch.no_grad():
        try:
            # Loop through the validation batches
            for i, (images, metadata, labels) in enumerate(progress_bar):
                images, metadata, labels = images.to(device), metadata.
                labels = labels.float().to(device)
                labels = labels.unsqueeze(1)

                probs = model(images, metadata)

                loss = criterion(probs, labels)
                running_val_loss += loss.item()

                all_labels.extend(labels.cpu().detach().numpy())
                all_probs.extend(probs.cpu().detach().numpy())

                # Calculate binary predictions for validation accuracy
                predicted = (probs >= 0.5).float()
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

                progress_bar.set_postfix(val_loss=running_val_loss / (i + 1))

        val_accuracy = 100 * correct / total
        val_loss = running_val_loss / len(val_dataloader)

```



```

        val_accuracies.append(val_accuracy)
        val_losses.append(val_loss)

        # Calculate AUROC
        try:
            valid_auroc = roc_auc_score(all_labels, all_probs)
        except ValueError as ve:
            print(f"AUROC Calculation Error: {ve}")
            valid_auroc = 0.0

        # Calculate partial AUC-above-TPR
        try:
            partial_auroc = score(np.array(all_labels), np.
↪array(all_probs), min_tpr=min_tpr)
        except ValueError as ve:
            print(f"Partial AUC Calculation Error: {ve}")
            partial_auroc = 0.0

        print(f'Epoch [{epoch + 1}/{epochs}], Train Loss:␣
↪{train_losses[-1]:.4f}, Val Loss: {val_loss:.4f}, '
              f'Val Accuracy: {val_accuracy:.2f}%, Val AUROC:␣
↪{valid_auroc:.4f}, Partial AUROC: {partial_auroc:.4f}')

        # Early stopping based on validation loss
        if val_loss < best_val_loss:
            best_val_loss = val_loss
            best_epoch = epoch + 1
            early_stopping_counter = 0
            torch.save(model.state_dict(), 'best_model.pth')
        else:
            early_stopping_counter += 1

        if early_stopping_counter >= early_stopping_patience:
            print(f"Early stopping triggered at epoch {epoch + 1}")
            break

    except Exception as e:
        print(f"Error during validation loop: {e}")
        break

    print(f"Best Epoch: {best_epoch}, Best Validation Loss: {best_val_loss:.
↪4f}")
    print('Training Complete')

    # Plot training and validation loss
    plt.figure(figsize=(10, 5))
    plt.plot(train_losses, label='Train Loss')

```

```

plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.show()

# Plot training and validation accuracy
plt.figure(figsize=(10, 5))
plt.plot(train_accuracies, label='Train Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy (%)')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.show()

# Generate classification report
try:
    print("Classification Report:")
    print(classification_report(all_labels, (np.array(all_probs) >= 0.5).
→astype(int), target_names=['Class 0', 'Class 1']))
except Exception as e:
    print(f"Error generating classification report: {e}")

return model

```

## 0.5 Model 1

```

[8]: model1 = CustomImageFeatureResNet(feature_input_size=9) # Assuming 9 features
    ↪ for metadata
model1.to(device)
# Initialize optimizer
optimizer = optim.Adam(model1.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed

```

in the future. The current behavior is equivalent to passing  
`weights=ResNet18\_Weights.IMAGENET1K\_V1`. You can also use  
`weights=ResNet18\_Weights.DEFAULT` to get the most up-to-date weights.  
warnings.warn(msg)

```
[9]: train_and_validate(model1,train_dataloader, val_dataloader, criterion,optimizer, epochs, device )
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [05:20<00:00, 9.70s/it,  
train\_loss=0.487]

Validating Epoch 1: 100%| | 24/24 [01:10<00:00, 2.93s/it,  
val\_loss=0.359]

Epoch [1/20], Train Loss: 0.4871, Val Loss: 0.3585, Val Accuracy: 85.57%, Val  
AUROC: 0.7102, Partial AUROC: 0.0499

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [05:13<00:00, 9.50s/it,  
train\_loss=0.37]

Validating Epoch 2: 100%| | 24/24 [01:09<00:00, 2.88s/it,  
val\_loss=0.458]

Epoch [2/20], Train Loss: 0.3699, Val Loss: 0.4577, Val Accuracy: 82.21%, Val  
AUROC: 0.8521, Partial AUROC: 0.1003

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [05:16<00:00, 9.58s/it,  
train\_loss=0.31]

Validating Epoch 3: 100%| | 24/24 [01:09<00:00, 2.88s/it,  
val\_loss=0.288]

Epoch [3/20], Train Loss: 0.3105, Val Loss: 0.2877, Val Accuracy: 83.22%, Val  
AUROC: 0.8656, Partial AUROC: 0.1226

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [05:07<00:00, 9.30s/it,  
train\_loss=0.248]

Validating Epoch 4: 100%| | 24/24 [01:11<00:00, 2.97s/it,  
val\_loss=0.301]

Epoch [4/20], Train Loss: 0.2482, Val Loss: 0.3012, Val Accuracy: 87.18%, Val  
AUROC: 0.8861, Partial AUROC: 0.1248

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [05:04<00:00, 9.24s/it,  
train\_loss=0.18]

Validating Epoch 5: 100%| | 24/24 [01:06<00:00, 2.77s/it,  
val\_loss=0.25]

Epoch [5/20], Train Loss: 0.1802, Val Loss: 0.2495, Val Accuracy: 90.27%, Val  
AUROC: 0.8954, Partial AUROC: 0.1271

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [05:06<00:00, 9.28s/it,  
train\_loss=0.152]

Validating Epoch 6: 100%| | 24/24 [01:06<00:00, 2.77s/it,  
val\_loss=0.206]

Epoch [6/20], Train Loss: 0.1522, Val Loss: 0.2059, Val Accuracy: 94.23%, Val  
AUROC: 0.7653, Partial AUROC: 0.0532

Epoch 7/20

Training Epoch 7: 100%| | 33/33 [05:02<00:00, 9.16s/it,  
train\_loss=0.145]

Validating Epoch 7: 100%| | 24/24 [01:07<00:00, 2.79s/it,  
val\_loss=1.37]

Epoch [7/20], Train Loss: 0.1453, Val Loss: 1.3696, Val Accuracy: 67.38%, Val  
AUROC: 0.8349, Partial AUROC: 0.1053

Epoch 8/20

Training Epoch 8: 100%| | 33/33 [05:04<00:00, 9.23s/it,  
train\_loss=0.125]

Validating Epoch 8: 100%| | 24/24 [01:14<00:00, 3.10s/it,  
val\_loss=0.193]

Epoch [8/20], Train Loss: 0.1246, Val Loss: 0.1926, Val Accuracy: 93.15%, Val  
AUROC: 0.8817, Partial AUROC: 0.1245

Epoch 9/20

Training Epoch 9: 100%| | 33/33 [05:04<00:00, 9.21s/it,  
train\_loss=0.0645]

Validating Epoch 9: 100%| | 24/24 [01:12<00:00, 3.00s/it,  
val\_loss=0.146]

Epoch [9/20], Train Loss: 0.0645, Val Loss: 0.1464, Val Accuracy: 96.51%, Val  
AUROC: 0.8796, Partial AUROC: 0.1200

Epoch 10/20

Training Epoch 10: 100%| | 33/33 [04:57<00:00, 9.02s/it,  
train\_loss=0.039]

Validating Epoch 10: 100%| | 24/24 [01:08<00:00, 2.87s/it,  
val\_loss=0.183]

Epoch [10/20], Train Loss: 0.0390, Val Loss: 0.1834, Val Accuracy: 95.03%, Val  
AUROC: 0.8746, Partial AUROC: 0.1210

Epoch 11/20

Training Epoch 11: 100%| | 33/33 [05:08<00:00, 9.34s/it,  
train\_loss=0.031]

Validating Epoch 11: 100%| | 24/24 [01:14<00:00, 3.10s/it,  
val\_loss=0.254]

Epoch [11/20], Train Loss: 0.0310, Val Loss: 0.2539, Val Accuracy: 92.55%, Val  
AUROC: 0.8574, Partial AUROC: 0.1095

Epoch 12/20

Training Epoch 12: 100%| | 33/33 [05:08<00:00, 9.36s/it,  
train\_loss=0.0615]

Validating Epoch 12: 100%| | 24/24 [01:11<00:00, 2.98s/it,  
val\_loss=0.189]

Epoch [12/20], Train Loss: 0.0615, Val Loss: 0.1888, Val Accuracy: 94.03%, Val  
AUROC: 0.8672, Partial AUROC: 0.1175

Epoch 13/20

Training Epoch 13: 100%| | 33/33 [04:57<00:00, 9.03s/it,  
train\_loss=0.0328]

Validating Epoch 13: 100%| | 24/24 [01:07<00:00, 2.80s/it,  
val\_loss=0.244]

Epoch [13/20], Train Loss: 0.0328, Val Loss: 0.2440, Val Accuracy: 93.76%, Val  
AUROC: 0.8565, Partial AUROC: 0.1048

Epoch 14/20

Training Epoch 14: 100%| | 33/33 [05:06<00:00, 9.27s/it,  
train\_loss=0.0311]


Validating Epoch 14: 100%| | 24/24 [01:08<00:00, 2.86s/it,  
val\_loss=0.275]


Epoch [14/20], Train Loss: 0.0311, Val Loss: 0.2752, Val Accuracy: 94.97%, Val  
AUROC: 0.8445, Partial AUROC: 0.1057

Early stopping triggered at epoch 14

Best Epoch: 9, Best Validation Loss: 0.1464

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.97	0.98	0.97	1431
Class 1	0.35	0.31	0.32	59
accuracy			0.95	1490
macro avg	0.66	0.64	0.65	1490
weighted avg	0.95	0.95	0.95	1490

```
[9]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(

```

```

        (0): BasicBlock(
          (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    (7): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```



```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (8): AdaptiveAvgPool2d(output_size=(1, 1))
    )
    (fc_image): Linear(in_features=512, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )

```

## 0.6 Model 2

```

[10]: model2 = CustomImageFeatureResNet(feature_input_size=9)  # Assuming 9 features
      ↪for metadata
model2.to(device)
# Initialize optimizer
optimizer = optim.SGD(model2.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss()  # Binary classification loss
# Set the number of epochs
epochs = 20

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[11]: train_and_validate(model2, train_dataloader, val_dataloader, criterion,
      ↪optimizer, epochs, device )

```

Epoch 1/20


Training Epoch 1: 100%| | 33/33 [05:10<00:00, 9.41s/it,  
 train\_loss=0.649]  
 Validating Epoch 1: 100%| | 24/24 [01:07<00:00, 2.80s/it,  
 val\_loss=0.484]  
 Epoch [1/20], Train Loss: 0.6494, Val Loss: 0.4835, Val Accuracy: 96.04%, Val  
 AUROC: 0.5763, Partial AUROC: 0.0208  
 Epoch 2/20  
 Training Epoch 2: 100%| | 33/33 [05:09<00:00, 9.37s/it,  
 train\_loss=0.637]  
 Validating Epoch 2: 100%| | 24/24 [01:08<00:00, 2.86s/it,  
 val\_loss=0.44]  
 Epoch [2/20], Train Loss: 0.6371, Val Loss: 0.4399, Val Accuracy: 96.04%, Val  
 AUROC: 0.5812, Partial AUROC: 0.0224  
 Epoch 3/20  
 Training Epoch 3: 100%| | 33/33 [05:13<00:00, 9.51s/it,  
 train\_loss=0.632]  
 Validating Epoch 3: 100%| | 24/24 [01:09<00:00, 2.88s/it,  
 val\_loss=0.427]  
 Epoch [3/20], Train Loss: 0.6323, Val Loss: 0.4268, Val Accuracy: 96.04%, Val  
 AUROC: 0.6200, Partial AUROC: 0.0318  
 Epoch 4/20  
 Training Epoch 4: 100%| | 33/33 [05:12<00:00, 9.46s/it,  
 train\_loss=0.622]  
 Validating Epoch 4: 100%| | 24/24 [01:09<00:00, 2.88s/it,  
 val\_loss=0.421]  
 Epoch [4/20], Train Loss: 0.6218, Val Loss: 0.4208, Val Accuracy: 96.04%, Val  
 AUROC: 0.6634, Partial AUROC: 0.0398  
 Epoch 5/20  
 Training Epoch 5: 100%| | 33/33 [05:13<00:00, 9.49s/it,  
 train\_loss=0.613]  
 Validating Epoch 5: 100%| | 24/24 [01:09<00:00, 2.88s/it,  
 val\_loss=0.429]  
 Epoch [5/20], Train Loss: 0.6126, Val Loss: 0.4288, Val Accuracy: 96.04%, Val  
 AUROC: 0.6879, Partial AUROC: 0.0439  
 Epoch 6/20  
 Training Epoch 6: 100%| | 33/33 [05:01<00:00, 9.13s/it,  
 train\_loss=0.606]  
 Validating Epoch 6: 100%| | 24/24 [01:07<00:00, 2.81s/it,  
 val\_loss=0.42]  
 Epoch [6/20], Train Loss: 0.6055, Val Loss: 0.4201, Val Accuracy: 96.04%, Val  
 AUROC: 0.7171, Partial AUROC: 0.0529  
 Epoch 7/20

Training Epoch 7: 100%|        | 33/33 [04:58<00:00, 9.05s/it,  
 train\_loss=0.599]  
 Validating Epoch 7: 100%|        | 24/24 [01:07<00:00, 2.79s/it,  
 val\_loss=0.407]  
  
 Epoch [7/20], Train Loss: 0.5989, Val Loss: 0.4068, Val Accuracy: 96.04%, Val  
 AUROC: 0.7272, Partial AUROC: 0.0528  
 Epoch 8/20  
  
 Training Epoch 8: 100%|        | 33/33 [05:12<00:00, 9.48s/it,  
 train\_loss=0.591]  
 Validating Epoch 8: 100%|        | 24/24 [01:08<00:00, 2.86s/it,  
 val\_loss=0.407]  
  
 Epoch [8/20], Train Loss: 0.5911, Val Loss: 0.4073, Val Accuracy: 96.04%, Val  
 AUROC: 0.7430, Partial AUROC: 0.0560  
 Epoch 9/20  
  
 Training Epoch 9: 100%|        | 33/33 [05:07<00:00, 9.32s/it,  
 train\_loss=0.587]  
 Validating Epoch 9: 100%|        | 24/24 [01:10<00:00, 2.95s/it,  
 val\_loss=0.405]  
  
 Epoch [9/20], Train Loss: 0.5865, Val Loss: 0.4047, Val Accuracy: 95.97%, Val  
 AUROC: 0.7575, Partial AUROC: 0.0573  
 Epoch 10/20  
  
 Training Epoch 10: 100%|        | 33/33 [05:11<00:00, 9.43s/it,  
 train\_loss=0.577]  
 Validating Epoch 10: 100%|        | 24/24 [01:08<00:00, 2.87s/it,  
 val\_loss=0.399]  
  
 Epoch [10/20], Train Loss: 0.5770, Val Loss: 0.3985, Val Accuracy: 95.97%, Val  
 AUROC: 0.7674, Partial AUROC: 0.0635  
 Epoch 11/20  
  
 Training Epoch 11: 100%|        | 33/33 [05:11<00:00, 9.44s/it,  
 train\_loss=0.567]  
 Validating Epoch 11: 100%|        | 24/24 [01:08<00:00, 2.87s/it,  
 val\_loss=0.404]  
  
 Epoch [11/20], Train Loss: 0.5675, Val Loss: 0.4040, Val Accuracy: 96.04%, Val  
 AUROC: 0.7689, Partial AUROC: 0.0615  
 Epoch 12/20  
  
 Training Epoch 12: 100%|        | 33/33 [05:08<00:00, 9.35s/it,  
 train\_loss=0.563]  
 Validating Epoch 12: 100%|        | 24/24 [01:11<00:00, 2.97s/it,  
 val\_loss=0.389]  
  
 Epoch [12/20], Train Loss: 0.5634, Val Loss: 0.3887, Val Accuracy: 95.91%, Val  
 AUROC: 0.7855, Partial AUROC: 0.0679  
 Epoch 13/20

Training Epoch 13: 100%|        | 33/33 [05:00<00:00, 9.10s/it,  
 train\_loss=0.559]  
 Validating Epoch 13: 100%|        | 24/24 [01:09<00:00, 2.90s/it,  
 val\_loss=0.381]  
  
 Epoch [13/20], Train Loss: 0.5594, Val Loss: 0.3807, Val Accuracy: 95.91%, Val  
 AUROC: 0.7903, Partial AUROC: 0.0692  
 Epoch 14/20  
  
 Training Epoch 14: 100%|        | 33/33 [05:09<00:00, 9.37s/it,  
 train\_loss=0.551]  
 Validating Epoch 14: 100%|        | 24/24 [01:09<00:00, 2.89s/it,  
 val\_loss=0.381]  
  
 Epoch [14/20], Train Loss: 0.5513, Val Loss: 0.3811, Val Accuracy: 95.91%, Val  
 AUROC: 0.7868, Partial AUROC: 0.0667  
 Epoch 15/20  
  
 Training Epoch 15: 100%|        | 33/33 [05:08<00:00, 9.35s/it,  
 train\_loss=0.542]  
 Validating Epoch 15: 100%|        | 24/24 [01:08<00:00, 2.86s/it,  
 val\_loss=0.379]  
  
 Epoch [15/20], Train Loss: 0.5425, Val Loss: 0.3791, Val Accuracy: 95.91%, Val  
 AUROC: 0.7936, Partial AUROC: 0.0692  
 Epoch 16/20  
  
 Training Epoch 16: 100%|        | 33/33 [05:04<00:00, 9.23s/it,  
 train\_loss=0.535]  
 Validating Epoch 16: 100%|        | 24/24 [01:08<00:00, 2.85s/it,  
 val\_loss=0.37]  
  
 Epoch [16/20], Train Loss: 0.5355, Val Loss: 0.3700, Val Accuracy: 95.77%, Val  
 AUROC: 0.8040, Partial AUROC: 0.0754  
 Epoch 17/20  
  
 Training Epoch 17: 100%|        | 33/33 [05:03<00:00, 9.18s/it,  
 train\_loss=0.523]  
 Validating Epoch 17: 100%|        | 24/24 [01:11<00:00, 2.98s/it,  
 val\_loss=0.365]  
  
 Epoch [17/20], Train Loss: 0.5230, Val Loss: 0.3650, Val Accuracy: 95.91%, Val  
 AUROC: 0.8042, Partial AUROC: 0.0754  
 Epoch 18/20  
  
 Training Epoch 18: 100%|        | 33/33 [05:03<00:00, 9.20s/it,  
 train\_loss=0.514]  
 Validating Epoch 18: 100%|        | 24/24 [01:08<00:00, 2.86s/it,  
 val\_loss=0.375]  
  
 Epoch [18/20], Train Loss: 0.5138, Val Loss: 0.3750, Val Accuracy: 94.77%, Val  
 AUROC: 0.8019, Partial AUROC: 0.0730  
 Epoch 19/20

Training Epoch 19: 100%| | 33/33 [05:02<00:00, 9.16s/it,  
train\_loss=0.518]  
Validating Epoch 19: 100%| | 24/24 [01:08<00:00, 2.87s/it,  
val\_loss=0.354]  
  
Epoch [19/20], Train Loss: 0.5178, Val Loss: 0.3542, Val Accuracy: 94.90%, Val  
AUROC: 0.8084, Partial AUROC: 0.0757  
Epoch 20/20  
  
Training Epoch 20: 100%| | 33/33 [04:59<00:00, 9.08s/it,  
train\_loss=0.501]  
Validating Epoch 20: 100%| | 24/24 [01:07<00:00, 2.80s/it,  
val\_loss=0.352]  
  
Epoch [20/20], Train Loss: 0.5011, Val Loss: 0.3516, Val Accuracy: 94.77%, Val  
AUROC: 0.8129, Partial AUROC: 0.0767  
Best Epoch: 20, Best Validation Loss: 0.3516  
Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.97	0.98	0.97	1431
Class 1	0.31	0.27	0.29	59
accuracy			0.95	1490
macro avg	0.64	0.62	0.63	1490
weighted avg	0.94	0.95	0.95	1490

```
[11]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
    ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
        1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(

```

```

        (0): BasicBlock(
          (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    (7): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```



```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (8): AdaptiveAvgPool2d(output_size=(1, 1))
    )
    (fc_image): Linear(in_features=512, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )

```

## 0.7 Model 3

```

[12]: model3 = CustomImageFeatureResNet(feature_input_size=9) # Assuming 9 features
      ↪for metadata
model3.to(device)
# Initialize optimizer
optimizer = optim.SGD(model3.parameters(), lr=0.0001, weight_decay=1e-4)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_size = 32

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[13]: train_dataloader = DataLoader(train_dataset, batch_size=batch_size,
      ↪shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_size, shuffle=True)

```

```
[14]: train_and_validate(model3,train_dataloader, val_dataloader, criterion,optimizer, epochs, device )
```

Epoch 1/20

Training Epoch 1: 100%| | 66/66 [04:39<00:00, 4.23s/it,  
train\_loss=0.7]

Validating Epoch 1: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.659]

Epoch [1/20], Train Loss: 0.6999, Val Loss: 0.6592, Val Accuracy: 77.05%, Val  
AUROC: 0.3767, Partial AUROC: 0.0063

Epoch 2/20

Training Epoch 2: 100%| | 66/66 [04:35<00:00, 4.18s/it,  
train\_loss=0.685]

Validating Epoch 2: 100%| | 47/47 [01:07<00:00, 1.44s/it,  
val\_loss=0.62]

Epoch [2/20], Train Loss: 0.6847, Val Loss: 0.6203, Val Accuracy: 92.28%, Val  
AUROC: 0.3632, Partial AUROC: 0.0067

Epoch 3/20

Training Epoch 3: 100%| | 66/66 [04:42<00:00, 4.28s/it,  
train\_loss=0.674]

Validating Epoch 3: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.586]

Epoch [3/20], Train Loss: 0.6740, Val Loss: 0.5864, Val Accuracy: 95.57%, Val  
AUROC: 0.3735, Partial AUROC: 0.0074

Epoch 4/20

Training Epoch 4: 100%| | 66/66 [04:47<00:00, 4.36s/it,  
train\_loss=0.667]

Validating Epoch 4: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.558]

Epoch [4/20], Train Loss: 0.6675, Val Loss: 0.5584, Val Accuracy: 96.04%, Val  
AUROC: 0.3819, Partial AUROC: 0.0084

Epoch 5/20

Training Epoch 5: 100%| | 66/66 [04:45<00:00, 4.32s/it,  
train\_loss=0.663]

Validating Epoch 5: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.541]

Epoch [5/20], Train Loss: 0.6630, Val Loss: 0.5414, Val Accuracy: 96.04%, Val  
AUROC: 0.3721, Partial AUROC: 0.0083

Epoch 6/20

Training Epoch 6: 100%| | 66/66 [04:37<00:00, 4.21s/it,  
train\_loss=0.65]

Validating Epoch 6: 100%| | 47/47 [01:07<00:00, 1.44s/it,  
val\_loss=0.521]

Epoch [6/20], Train Loss: 0.6500, Val Loss: 0.5210, Val Accuracy: 96.04%, Val  
AUROC: 0.3960, Partial AUROC: 0.0105

Epoch 7/20

Training Epoch 7: 100%| | 66/66 [04:34<00:00, 4.16s/it,  
train\_loss=0.652]

Validating Epoch 7: 100%| | 47/47 [01:05<00:00, 1.39s/it,  
val\_loss=0.506]

Epoch [7/20], Train Loss: 0.6518, Val Loss: 0.5059, Val Accuracy: 96.04%, Val  
AUROC: 0.4008, Partial AUROC: 0.0120

Epoch 8/20

Training Epoch 8: 100%| | 66/66 [04:38<00:00, 4.22s/it,  
train\_loss=0.645]

Validating Epoch 8: 100%| | 47/47 [01:05<00:00, 1.39s/it,  
val\_loss=0.496]

Epoch [8/20], Train Loss: 0.6452, Val Loss: 0.4963, Val Accuracy: 96.04%, Val  
AUROC: 0.4022, Partial AUROC: 0.0107

Epoch 9/20

Training Epoch 9: 100%| | 66/66 [04:33<00:00, 4.14s/it,  
train\_loss=0.645]

Validating Epoch 9: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.483]

Epoch [9/20], Train Loss: 0.6446, Val Loss: 0.4831, Val Accuracy: 96.04%, Val  
AUROC: 0.4151, Partial AUROC: 0.0131

Epoch 10/20

Training Epoch 10: 100%| | 66/66 [04:34<00:00, 4.17s/it,  
train\_loss=0.636]

Validating Epoch 10: 100%| | 47/47 [01:04<00:00, 1.38s/it,  
val\_loss=0.477]

Epoch [10/20], Train Loss: 0.6357, Val Loss: 0.4769, Val Accuracy: 96.04%, Val  
AUROC: 0.4274, Partial AUROC: 0.0137

Epoch 11/20

Training Epoch 11: 100%| | 66/66 [04:40<00:00, 4.25s/it,  
train\_loss=0.634]

Validating Epoch 11: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.47]

Epoch [11/20], Train Loss: 0.6341, Val Loss: 0.4700, Val Accuracy: 96.04%, Val  
AUROC: 0.4331, Partial AUROC: 0.0140

Epoch 12/20

Training Epoch 12: 100%| | 66/66 [04:41<00:00, 4.26s/it,  
train\_loss=0.639]

Validating Epoch 12: 100%| | 47/47 [01:07<00:00, 1.44s/it,  
val\_loss=0.465]

Epoch [12/20], Train Loss: 0.6386, Val Loss: 0.4655, Val Accuracy: 96.04%, Val  
AUROC: 0.4561, Partial AUROC: 0.0164

Epoch 13/20

Training Epoch 13: 100%| | 66/66 [04:45<00:00, 4.32s/it,  
train\_loss=0.638]

Validating Epoch 13: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.46]

Epoch [13/20], Train Loss: 0.6376, Val Loss: 0.4599, Val Accuracy: 96.04%, Val  
AUROC: 0.4592, Partial AUROC: 0.0170

Epoch 14/20

Training Epoch 14: 100%| | 66/66 [04:44<00:00, 4.30s/it,  
train\_loss=0.633]

Validating Epoch 14: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.454]

Epoch [14/20], Train Loss: 0.6331, Val Loss: 0.4540, Val Accuracy: 96.04%, Val  
AUROC: 0.4620, Partial AUROC: 0.0180

Epoch 15/20

Training Epoch 15: 100%| | 66/66 [04:44<00:00, 4.30s/it,  
train\_loss=0.632]

Validating Epoch 15: 100%| | 47/47 [01:07<00:00, 1.44s/it,  
val\_loss=0.452]

Epoch [15/20], Train Loss: 0.6315, Val Loss: 0.4521, Val Accuracy: 96.04%, Val  
AUROC: 0.4847, Partial AUROC: 0.0194

Epoch 16/20

Training Epoch 16: 100%| | 66/66 [04:44<00:00, 4.31s/it,  
train\_loss=0.63]

Validating Epoch 16: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.447]

Epoch [16/20], Train Loss: 0.6296, Val Loss: 0.4473, Val Accuracy: 96.04%, Val  
AUROC: 0.4830, Partial AUROC: 0.0192

Epoch 17/20

Training Epoch 17: 100%| | 66/66 [04:42<00:00, 4.29s/it,  
train\_loss=0.633]

Validating Epoch 17: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.444]

Epoch [17/20], Train Loss: 0.6327, Val Loss: 0.4435, Val Accuracy: 96.04%, Val  
AUROC: 0.4997, Partial AUROC: 0.0210

Epoch 18/20

Training Epoch 18: 100%| | 66/66 [04:45<00:00, 4.32s/it,  
train\_loss=0.629]

Validating Epoch 18: 100%| | 47/47 [01:07<00:00, 1.43s/it,  
val\_loss=0.443]

Epoch [18/20], Train Loss: 0.6285, Val Loss: 0.4429, Val Accuracy: 96.04%, Val  
AUROC: 0.5151, Partial AUROC: 0.0219

Epoch 19/20

Training Epoch 19: 100%| | 66/66 [04:41<00:00, 4.27s/it,  
train\_loss=0.625]

Validating Epoch 19: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.442]

Epoch [19/20], Train Loss: 0.6254, Val Loss: 0.4423, Val Accuracy: 96.04%, Val  
AUROC: 0.5279, Partial AUROC: 0.0251

Epoch 20/20


Training Epoch 20: 100%| | 66/66 [04:42<00:00, 4.29s/it,  
train\_loss=0.619]


Validating Epoch 20: 100%| | 47/47 [01:06<00:00, 1.42s/it,  
val\_loss=0.44]

Epoch [20/20], Train Loss: 0.6192, Val Loss: 0.4397, Val Accuracy: 96.04%, Val  
AUROC: 0.5377, Partial AUROC: 0.0247

Best Epoch: 20, Best Validation Loss: 0.4397

Training Complete





#### Classification Report:

	precision	recall	f1-score	support
Class 0	0.96	1.00	0.98	1431
Class 1	0.00	0.00	0.00	59
accuracy			0.96	1490
macro avg	0.48	0.50	0.49	1490
weighted avg	0.92	0.96	0.94	1490

```

/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/tljh/user/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

[14]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (5): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (6): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (7): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,

```



```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
  (8): AdaptiveAvgPool2d(output_size=(1, 1))
)
(fc_image): Linear(in_features=512, out_features=512, bias=True)
(fc_metadata): Linear(in_features=9, out_features=128, bias=True)
(dropout): Dropout(p=0.5, inplace=False)
(fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

[ ]:

[ ]: