

Week_8_Third_Modeling

October 29, 2024

```
[1]: # Standard Libraries
import io
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Deep Learning and PyTorch
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import models

# Image Processing
from PIL import Image
from torchvision import transforms, models

# File Handling
import h5py

# Metrics and Evaluation
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, auc

# Progress Visualization
from tqdm import tqdm
```

0.1 Create Custom Dataset

```
[2]: class MultiInputDataset(Dataset):
    def __init__(self, hdf5_file, csv_file, transform=None):
        # Open the HDF5 file with error handling
        try:
            self.hdf5_file = h5py.File(hdf5_file, 'r') # Read-only mode
        except Exception as e:
            raise IOError(f"Could not open HDF5 file: {hdf5_file}. Error: {e}")

        # Read the CSV file containing image labels and additional features
```

```

try:
    self.labels_df = pd.read_csv(csv_file)
except Exception as e:
    raise IOError(f"Could not read CSV file: {csv_file}. Error: {e}")

# Ensure that all image IDs from the CSV are present in the HDF5 file
self.image_ids = self.labels_df['isic_id'].values
for image_id in self.image_ids:
    if str(image_id) not in self.hdf5_file.keys():
        raise ValueError(f"Image id {image_id} not found in HDF5 file.")

# Store any transformations to be applied to the images
self.transform = transform

def __len__(self):
    # Return the total number of samples in the dataset
    return len(self.labels_df)

def __getitem__(self, idx):
    # Get the image ID from the CSV file based on index
    image_id = str(self.labels_df.iloc[idx]['isic_id'])

    # Load the image data from the HDF5 file
    image_bytes = self.hdf5_file[image_id][()]

    # Convert the image bytes to a PIL Image
    image = Image.open(io.BytesIO(image_bytes))

    # Apply any specified transformations to the image
    if self.transform:
        image = self.transform(image)

    # Retrieve the label
    label = torch.tensor(self.labels_df.iloc[idx]['target'], dtype=torch.
↪long) # Adjust dtype if needed

    # Retrieve other features, excluding 'isic_id' and 'target'
    other_variables = self.labels_df.iloc[idx].drop(['isic_id', 'target']).
↪values.astype(float)

    # Convert other variables (metadata) to a tensor
    metadata_tensor = torch.tensor(other_variables, dtype=torch.float32)

    # Return the image, metadata, and label
    return image, metadata_tensor, label

```

```
[3]: # Define any necessary transformations for the image dataset

# Transformations for training set (with data augmentation)
train_transform = transforms.Compose([
    transforms.Resize((224, 224)), # Resize to 225x225
    transforms.RandomResizedCrop(128, scale=(0.8, 1.0)), # Random crop to
    ↪225x225 with scale
    transforms.RandomRotation(10), # Randomly rotate images by 10 degrees
    transforms.ToTensor(), # Convert image to PyTorch tensor
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# Transformations for validation/test set (no data augmentation)
normal_transform = transforms.Compose([
    transforms.Resize((224, 224)), # Ensure the image is resized to 225x225
    transforms.ToTensor(), # Convert image to PyTorch tensor
])
```

0.2 Train DataLoader

```
[4]: # Initialize the dataset
train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.hdf5',
    ↪csv_file='../data/processed/processed-train-metadata1.csv',
    ↪transform=normal_transform)
val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.hdf5',
    ↪csv_file='../data/processed/processed-validation-metadata1.csv',
    ↪transform=normal_transform)

# Create a DataLoader
train_dataloader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=64, shuffle=True)

[5]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

0.3 Model Building

```
[6]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models

class CustomImageFeatureEfficientNet(nn.Module):
    def __init__(self, feature_input_size, pretrained=True):
        super(CustomImageFeatureEfficientNet, self).__init__()
```

```

    # Load a pretrained EfficientNet model for image feature extraction
    ↪(EfficientNet-B0 in this case)
    efficientnet = models.efficientnet_b0(pretrained=pretrained) # You can
    ↪change this to another EfficientNet version like B1 or B7
    self.efficientnet = nn.Sequential(*list(efficientnet.children())[:-1])
    ↪# Remove the final classification layer

    # The output of EfficientNet-B0's last conv layer is 1280-dimensional
    self.fc_image = nn.Linear(1280, 512) # Reduce dimension to match your
    ↪custom architecture

    # Fully connected layer for metadata (feature data)
    self.fc_metadata = nn.Linear(feature_input_size, 128)

    # Dropout layer to prevent overfitting
    self.dropout = nn.Dropout(0.5) # 50% dropout

    # Final fully connected layer for binary classification (combined image
    ↪+ feature input)
    self.fc_combined = nn.Linear(512 + 128, 1) # For binary classification

    def forward(self, image, metadata):
        # Forward pass for the image through EfficientNet (without the final
        ↪classification layer)
        x = self.efficientnet(image) # EfficientNet feature extraction
        x = x.view(x.size(0), -1) # Flatten the EfficientNet output
        image_features = F.relu(self.fc_image(x))

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            f"Batch sizes do not match! Image batch size: {image_features.
            ↪shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
        ↪dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

        # If you're using BCELoss, uncomment the next line to apply sigmoid
        output = torch.sigmoid(output)

```

```
return output
```

0.4 Model Training

```
[7]: # Function to compute partial AUC-above-TPR
def score(solution: np.array, submission: np.array, min_tpr: float = 0.80) -> float:
    """
    Compute the partial AUC by focusing on a specific range of true positive
    rates (TPR).

    Args:
        solution (np.array): Ground truth binary labels.
        submission (np.array): Model predictions.
        min_tpr (float): Minimum true positive rate to calculate partial AUC.

    Returns:
        float: The calculated partial AUC.

    Raises:
        ValueError: If the min_tpr is not within a valid range.
    """
    # Rescale the target to handle sklearn limitations and flip the predictions
    v_gt = abs(solution - 1)
    v_pred = -1.0 * submission
    max_fpr = abs(1 - min_tpr)

    # Compute ROC curve using sklearn
    fpr, tpr, _ = roc_curve(v_gt, v_pred)
    if max_fpr is None or max_fpr == 1:
        return auc(fpr, tpr)
    if max_fpr <= 0 or max_fpr > 1:
        raise ValueError(f"Expected min_tpr in range [0, 1), got: {min_tpr}")

    # Interpolate for partial AUC
    stop = np.searchsorted(fpr, max_fpr, "right")
    x_interp = [fpr[stop - 1], fpr[stop]]
    y_interp = [tpr[stop - 1], tpr[stop]]
    tpr = np.append(tpr[:stop], np.interp(max_fpr, x_interp, y_interp))
    fpr = np.append(fpr[:stop], max_fpr)
    partial_auc = auc(fpr, tpr)

    return partial_auc

# Training and validation loop function
```

```

def train_and_validate(
    model: nn.Module,
    train_dataloader: torch.utils.data.DataLoader,
    val_dataloader: torch.utils.data.DataLoader,
    criterion: nn.Module,
    optimizer: torch.optim.Optimizer,
    epochs: int,
    device: torch.device,
    early_stopping_patience: int = 5,
    min_tpr: float = 0.80
) -> nn.Module:
    """
    Train and validate a PyTorch model with early stopping, AUROC, partial AUC,
    and error handling.

    Args:
        model (nn.Module): The model to be trained and validated.
        train_dataloader (torch.utils.data.DataLoader): Dataloader for training
        data.
        val_dataloader (torch.utils.data.DataLoader): Dataloader for validation
        data.
        criterion (nn.Module): Loss function.
        optimizer (torch.optim.Optimizer): Optimizer to update the model.
        epochs (int): Number of training epochs.
        device (torch.device): The device (CPU or GPU) to use.
        early_stopping_patience (int): Early stopping patience.
        min_tpr (float): The minimum true positive rate for calculating partial
        AUC.

    Returns:
        nn.Module: The trained model.
    """
    # Initialize tracking variables
    best_val_loss = float('inf')
    best_epoch = 0
    train_losses = []
    val_losses = []
    train_accuracies = []
    val_accuracies = []
    early_stopping_counter = 0

    # Start the training and validation loop
    for epoch in range(epochs):
        print(f'Epoch {epoch + 1}/{epochs}')

        # Training phase
        model.train()

```

```

        running_train_loss = 0.0
        correct_train = 0
        total_train = 0
        all_train_labels = []
        all_train_probs = []

        progress_bar = tqdm(train_dataloader, desc=f'Training Epoch {epoch + 1}')

    try:
        # Loop through the training batches
        for i, (image, metadata, labels) in enumerate(progress_bar):
            image, metadata, labels = image.to(device), metadata.
            to(device), labels.float().to(device)
            labels = labels.unsqueeze(1) # Adjust labels to have the right
            shape for binary classification

            optimizer.zero_grad()

            # Forward pass
            probs = model(image, metadata)

            if probs.shape != labels.shape:
                raise ValueError(f"Shape mismatch: Predictions shape {probs.
            shape} does not match labels shape {labels.shape}")

            # Calculate loss and backpropagate
            loss = criterion(probs, labels)
            loss.backward()
            optimizer.step()

            # Update running loss
            running_train_loss += loss.item()

            # Store labels and predictions for accuracy calculations
            all_train_labels.extend(labels.cpu().detach().numpy())
            all_train_probs.extend(probs.cpu().detach().numpy())

            # Calculate binary predictions for training accuracy
            predicted_train = (probs >= 0.5).float()
            total_train += labels.size(0)
            correct_train += (predicted_train == labels).sum().item()

            # Update progress bar
            progress_bar.set_postfix(train_loss=running_train_loss / (i + 1))

```

```

        # Calculate training accuracy and loss
        train_accuracy = 100 * correct_train / total_train
        train_losses.append(running_train_loss / len(train_dataloader))
        train_accuracies.append(train_accuracy)

    except ValueError as ve:
        print(f"Error during training loop: {ve}")
        break

    # Validation phase
    model.eval()
    running_val_loss = 0.0
    correct = 0
    total = 0
    all_labels = []
    all_probs = []

    progress_bar = tqdm(val_dataloader, desc=f'Validating Epoch {epoch + 1}')

    with torch.no_grad():
        try:
            # Loop through the validation batches
            for i, (images, metadata, labels) in enumerate(progress_bar):
                images, metadata, labels = images.to(device), metadata.
                labels = labels.float().to(device)
                labels = labels.unsqueeze(1)

                probs = model(images, metadata)

                loss = criterion(probs, labels)
                running_val_loss += loss.item()

                all_labels.extend(labels.cpu().detach().numpy())
                all_probs.extend(probs.cpu().detach().numpy())

                # Calculate binary predictions for validation accuracy
                predicted = (probs >= 0.5).float()
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

                progress_bar.set_postfix(val_loss=running_val_loss / (i + 1))

        val_accuracy = 100 * correct / total
        val_loss = running_val_loss / len(val_dataloader)
        val_accuracies.append(val_accuracy)

```



```

        val_losses.append(val_loss)

        # Calculate AUROC
        try:
            valid_auroc = roc_auc_score(all_labels, all_probs)
        except ValueError as ve:
            print(f"AUROC Calculation Error: {ve}")
            valid_auroc = 0.0

        # Calculate partial AUC-above-TPR
        try:
            partial_auroc = score(np.array(all_labels), np.
↪array(all_probs), min_tpr=min_tpr)
        except ValueError as ve:
            print(f"Partial AUC Calculation Error: {ve}")
            partial_auroc = 0.0

        print(f'Epoch [{epoch + 1}/{epochs}], Train Loss:␣
↪{train_losses[-1]:.4f}, Val Loss: {val_loss:.4f}, '
              f'Val Accuracy: {val_accuracy:.2f}%, Val AUROC:␣
↪{valid_auroc:.4f}, Partial AUROC: {partial_auroc:.4f}')

        # Early stopping based on validation loss
        if val_loss < best_val_loss:
            best_val_loss = val_loss
            best_epoch = epoch + 1
            early_stopping_counter = 0
            torch.save(model.state_dict(), 'best_model.pth')
        else:
            early_stopping_counter += 1

        if early_stopping_counter >= early_stopping_patience:
            print(f"Early stopping triggered at epoch {epoch + 1}")
            break

    except Exception as e:
        print(f"Error during validation loop: {e}")
        break

    print(f"Best Epoch: {best_epoch}, Best Validation Loss: {best_val_loss:.
↪4f}")
    print('Training Complete')

    # Plot training and validation loss
    plt.figure(figsize=(10, 5))
    plt.plot(train_losses, label='Train Loss')
    plt.plot(val_losses, label='Validation Loss')

```

```

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.show()

# Plot training and validation accuracy
plt.figure(figsize=(10, 5))
plt.plot(train_accuracies, label='Train Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy (%)')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.show()

# Generate classification report
try:
    print("Classification Report:")
    print(classification_report(all_labels, (np.array(all_probs) >= 0.5).
→astype(int), target_names=['Class 0', 'Class 1']))
except Exception as e:
    print(f"Error generating classification report: {e}")

return model

```

0.5 Model 1

```

[8]: model1 = CustomImageFeatureEfficientNet(feature_input_size=9) # Assuming 9
→ features for metadata
model1.to(device)
# Initialize optimizer
optimizer = optim.Adam(model1.parameters(), lr= 1.1621608010269284e-05)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_sizes = 16

train_dataloader = DataLoader(train_dataset, batch_size=batch_sizes,
→ shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_sizes, shuffle=True)

```

/home/jupyter-sohka/.local/lib/python3.10/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

```
warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
warnings.warn(msg)
```

```
[9]: train_and_validate(model1,train_dataloader, val_dataloader, criterion,
    ↪optimizer, epochs, device )
```

Epoch 1/20

```
Training Epoch 1: 100%|      | 131/131 [06:36<00:00,  3.02s/it,
train_loss=0.656]
```

```
Validating Epoch 1: 100%|      | 94/94 [01:07<00:00,  1.40it/s,
val_loss=0.55]
```

```
Epoch [1/20], Train Loss: 0.6557, Val Loss: 0.5498, Val Accuracy: 92.48%, Val
AUROC: 0.6143, Partial AUROC: 0.0387
```

Epoch 2/20

```
Training Epoch 2: 100%|      | 131/131 [07:20<00:00,  3.36s/it,
train_loss=0.574]
```

```
Validating Epoch 2: 100%|      | 94/94 [00:57<00:00,  1.64it/s,
val_loss=0.422]
```

```
Epoch [2/20], Train Loss: 0.5737, Val Loss: 0.4218, Val Accuracy: 92.01%, Val
AUROC: 0.7433, Partial AUROC: 0.0565
```

Epoch 3/20

```
Training Epoch 3: 100%|      | 131/131 [05:11<00:00,  2.38s/it,
train_loss=0.476]
```

```
Validating Epoch 3: 100%|      | 94/94 [00:40<00:00,  2.32it/s,
val_loss=0.317]
```

```
Epoch [3/20], Train Loss: 0.4761, Val Loss: 0.3170, Val Accuracy: 92.21%, Val
AUROC: 0.8123, Partial AUROC: 0.0790
```

Epoch 4/20

```
Training Epoch 4: 100%|      | 131/131 [05:19<00:00,  2.44s/it,
train_loss=0.399]
```

```
Validating Epoch 4: 100%|      | 94/94 [00:40<00:00,  2.32it/s,
val_loss=0.263]
```

```
Epoch [4/20], Train Loss: 0.3991, Val Loss: 0.2633, Val Accuracy: 92.08%, Val
AUROC: 0.8400, Partial AUROC: 0.0917
```

Epoch 5/20

```
Training Epoch 5: 100%|      | 131/131 [05:17<00:00,  2.42s/it,
train_loss=0.339]
```

Validating Epoch 5: 100%| | 94/94 [00:42<00:00, 2.24it/s,
val_loss=0.251]

Epoch [5/20], Train Loss: 0.3393, Val Loss: 0.2508, Val Accuracy: 90.67%, Val
AUROC: 0.8569, Partial AUROC: 0.1019

Epoch 6/20

Training Epoch 6: 100%| | 131/131 [05:31<00:00, 2.53s/it,
train_loss=0.292]

Validating Epoch 6: 100%| | 94/94 [00:41<00:00, 2.24it/s,
val_loss=0.227]

Epoch [6/20], Train Loss: 0.2924, Val Loss: 0.2267, Val Accuracy: 90.07%, Val
AUROC: 0.8687, Partial AUROC: 0.1095

Epoch 7/20

Training Epoch 7: 100%| | 131/131 [05:32<00:00, 2.54s/it,
train_loss=0.259]

Validating Epoch 7: 100%| | 94/94 [00:51<00:00, 1.84it/s,
val_loss=0.193]

Epoch [7/20], Train Loss: 0.2592, Val Loss: 0.1929, Val Accuracy: 92.75%, Val
AUROC: 0.8605, Partial AUROC: 0.1050

Epoch 8/20

Training Epoch 8: 100%| | 131/131 [05:27<00:00, 2.50s/it,
train_loss=0.226]

Validating Epoch 8: 100%| | 94/94 [00:41<00:00, 2.25it/s,
val_loss=0.267]

Epoch [8/20], Train Loss: 0.2260, Val Loss: 0.2668, Val Accuracy: 86.24%, Val
AUROC: 0.8765, Partial AUROC: 0.1165

Epoch 9/20

Training Epoch 9: 100%| | 131/131 [05:34<00:00, 2.55s/it,
train_loss=0.195]

Validating Epoch 9: 100%| | 94/94 [00:49<00:00, 1.90it/s,
val_loss=0.2]

Epoch [9/20], Train Loss: 0.1948, Val Loss: 0.2004, Val Accuracy: 91.14%, Val
AUROC: 0.8652, Partial AUROC: 0.1062

Epoch 10/20

Training Epoch 10: 100%| | 131/131 [05:33<00:00, 2.55s/it,
train_loss=0.177]

Validating Epoch 10: 100%| | 94/94 [00:51<00:00, 1.83it/s,
val_loss=0.245]

Epoch [10/20], Train Loss: 0.1773, Val Loss: 0.2454, Val Accuracy: 88.39%, Val
AUROC: 0.8726, Partial AUROC: 0.1106

Epoch 11/20

Training Epoch 11: 100%| | 131/131 [05:37<00:00, 2.58s/it,
train_loss=0.148]

Validating Epoch 11: 100%| | 94/94 [00:42<00:00, 2.22it/s,
val_loss=0.211]

Epoch [11/20], Train Loss: 0.1485, Val Loss: 0.2108, Val Accuracy: 90.60%, Val
AUROC: 0.8619, Partial AUROC: 0.1064

Epoch 12/20

Training Epoch 12: 100%| | 131/131 [05:31<00:00, 2.53s/it,
train_loss=0.134]


Validating Epoch 12: 100%| | 94/94 [00:41<00:00, 2.25it/s,
val_loss=0.206]


Epoch [12/20], Train Loss: 0.1338, Val Loss: 0.2058, Val Accuracy: 91.01%, Val
AUROC: 0.8582, Partial AUROC: 0.1033

Early stopping triggered at epoch 12

Best Epoch: 7, Best Validation Loss: 0.1929

Training Complete





Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Class 0 | 0.98 | 0.93 | 0.95 | 1431 |
| Class 1 | 0.22 | 0.51 | 0.31 | 59 |
| accuracy | | | 0.91 | 1490 |
| macro avg | 0.60 | 0.72 | 0.63 | 1490 |
| weighted avg | 0.95 | 0.91 | 0.93 | 1490 |

```
[9]: CustomImageFeatureEfficientNet(
      (efficientnet): Sequential(
        (0): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Sequential(
            (0): MBConv(
              (block): Sequential(
                (0): Conv2dNormActivation(
                  (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
```

```

        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (2): Conv2dNormActivation(
        (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.0, mode=row)
    )
    )
    (2): Sequential(
        (0): MBConv(
            (block): Sequential(
                (0): Conv2dNormActivation(
                    (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(
                (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
                (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.025, mode=row)
)
(3): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
    )
  )

```



```

        (1): Conv2dNormActivation(
          (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
          (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
    )
    (1): MBCConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(

```

```

        (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
)
)
(4): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
  )
  (1): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
)
(2): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))

```

```

        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
)
)
(5): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
  )
  (1): MBConv(

```

```

        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
      )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
      )
    )
  )
)

```

```

        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.125, mode=row)
    )
  )
  (6): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )

```

```

    )
    (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)

```

```

        (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
)
(3): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,

```



```

track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
    )
    )
    (7): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
track_running_stats=True)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      )
      (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
    )
    )
    (8): Conv2dNormActivation(
      (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
  )
)

```

```

        (1): AdaptiveAvgPool2d(output_size=1)
    )
    (fc_image): Linear(in_features=1280, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

0.6 Model 2

```

[8]: model2 = CustomImageFeatureEfficientNet(feature_input_size=9) # Assuming 9
    ↪ features for metadata
model2.to(device)
# Initialize optimizer
optimizer = optim.SGD(model2.parameters(), lr=0.01)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_sizes = 16

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[9]: train_dataloader = DataLoader(train_dataset, batch_size=batch_sizes,
    ↪ shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_sizes, shuffle=True)

```

```

[10]: train_and_validate(model2, train_dataloader, val_dataloader, criterion,
    ↪ optimizer, epochs, device )

```

Epoch 1/20

```

Training Epoch 1: 100%|          | 131/131 [04:34<00:00,  2.10s/it,
train_loss=0.616]
Validating Epoch 1: 100%|          | 94/94 [00:43<00:00,  2.15it/s,
val_loss=0.43]

```

Epoch [1/20], Train Loss: 0.6158, Val Loss: 0.4299, Val Accuracy: 95.23%, Val AUROC: 0.6865, Partial AUROC: 0.0289

Epoch 2/20

Training Epoch 2: 100%| | 131/131 [04:30<00:00, 2.07s/it, train_loss=0.551]

Validating Epoch 2: 100%| | 94/94 [00:42<00:00, 2.21it/s, val_loss=0.394]

Epoch [2/20], Train Loss: 0.5508, Val Loss: 0.3940, Val Accuracy: 93.96%, Val AUROC: 0.7499, Partial AUROC: 0.0437

Epoch 3/20

Training Epoch 3: 100%| | 131/131 [04:26<00:00, 2.03s/it, train_loss=0.487]

Validating Epoch 3: 100%| | 94/94 [00:45<00:00, 2.07it/s, val_loss=0.374]

Epoch [3/20], Train Loss: 0.4873, Val Loss: 0.3735, Val Accuracy: 90.13%, Val AUROC: 0.7978, Partial AUROC: 0.0662

Epoch 4/20

Training Epoch 4: 100%| | 131/131 [04:33<00:00, 2.09s/it, train_loss=0.415]

Validating Epoch 4: 100%| | 94/94 [00:42<00:00, 2.23it/s, val_loss=0.37]

Epoch [4/20], Train Loss: 0.4153, Val Loss: 0.3697, Val Accuracy: 90.13%, Val AUROC: 0.8109, Partial AUROC: 0.0820

Epoch 5/20

Training Epoch 5: 100%| | 131/131 [04:26<00:00, 2.04s/it, train_loss=0.356]

Validating Epoch 5: 100%| | 94/94 [00:42<00:00, 2.21it/s, val_loss=0.314]

Epoch [5/20], Train Loss: 0.3559, Val Loss: 0.3136, Val Accuracy: 88.12%, Val AUROC: 0.8329, Partial AUROC: 0.0810

Epoch 6/20

Training Epoch 6: 100%| | 131/131 [04:30<00:00, 2.06s/it, train_loss=0.287]

Validating Epoch 6: 100%| | 94/94 [00:42<00:00, 2.20it/s, val_loss=0.258]

Epoch [6/20], Train Loss: 0.2868, Val Loss: 0.2583, Val Accuracy: 90.81%, Val AUROC: 0.8405, Partial AUROC: 0.0820

Epoch 7/20

Training Epoch 7: 100%| | 131/131 [04:26<00:00, 2.03s/it, train_loss=0.233]

Validating Epoch 7: 100%| | 94/94 [00:42<00:00, 2.23it/s, val_loss=0.242]

Epoch [7/20], Train Loss: 0.2332, Val Loss: 0.2420, Val Accuracy: 90.47%, Val AUROC: 0.8521, Partial AUROC: 0.0878

Epoch 8/20

Training Epoch 8: 100%| | 131/131 [04:31<00:00, 2.07s/it, train_loss=0.2]

Validating Epoch 8: 100%| | 94/94 [00:42<00:00, 2.19it/s, val_loss=0.397]

Epoch [8/20], Train Loss: 0.2000, Val Loss: 0.3967, Val Accuracy: 79.19%, Val AUROC: 0.8120, Partial AUROC: 0.0802

Epoch 9/20

Training Epoch 9: 100%| | 131/131 [04:15<00:00, 1.95s/it, train_loss=0.161]

Validating Epoch 9: 100%| | 94/94 [00:40<00:00, 2.31it/s, val_loss=0.425]

Epoch [9/20], Train Loss: 0.1610, Val Loss: 0.4249, Val Accuracy: 75.23%, Val AUROC: 0.8130, Partial AUROC: 0.0871

Epoch 10/20

Training Epoch 10: 100%| | 131/131 [04:20<00:00, 1.99s/it, train_loss=0.135]

Validating Epoch 10: 100%| | 94/94 [00:41<00:00, 2.28it/s, val_loss=0.244]

Epoch [10/20], Train Loss: 0.1353, Val Loss: 0.2439, Val Accuracy: 89.87%, Val AUROC: 0.8739, Partial AUROC: 0.1062

Epoch 11/20

Training Epoch 11: 100%| | 131/131 [04:17<00:00, 1.96s/it, train_loss=0.11]

Validating Epoch 11: 100%| | 94/94 [00:41<00:00, 2.24it/s, val_loss=0.642]

Epoch [11/20], Train Loss: 0.1099, Val Loss: 0.6419, Val Accuracy: 70.87%, Val AUROC: 0.7527, Partial AUROC: 0.0498

Epoch 12/20

Training Epoch 12: 100%| | 131/131 [04:20<00:00, 1.99s/it, train_loss=0.105]

Validating Epoch 12: 100%| | 94/94 [00:41<00:00, 2.28it/s, val_loss=0.241]

Epoch [12/20], Train Loss: 0.1052, Val Loss: 0.2414, Val Accuracy: 91.28%, Val AUROC: 0.8585, Partial AUROC: 0.1001

Epoch 13/20

Training Epoch 13: 100%| | 131/131 [04:14<00:00, 1.94s/it, train_loss=0.0961]

Validating Epoch 13: 100%| | 94/94 [00:40<00:00, 2.29it/s, val_loss=0.203]

Epoch [13/20], Train Loss: 0.0961, Val Loss: 0.2032, Val Accuracy: 92.75%, Val AUROC: 0.8503, Partial AUROC: 0.0886

Epoch 14/20

Training Epoch 14: 100%| | 131/131 [04:25<00:00, 2.02s/it, train_loss=0.078]

Validating Epoch 14: 100%| | 94/94 [00:41<00:00, 2.27it/s, val_loss=0.249]

Epoch [14/20], Train Loss: 0.0780, Val Loss: 0.2486, Val Accuracy: 91.14%, Val AUROC: 0.8618, Partial AUROC: 0.1037

Epoch 15/20

Training Epoch 15: 100%| | 131/131 [04:16<00:00, 1.96s/it, train_loss=0.0776]

Validating Epoch 15: 100%| | 94/94 [00:42<00:00, 2.21it/s, val_loss=0.269]

Epoch [15/20], Train Loss: 0.0776, Val Loss: 0.2688, Val Accuracy: 90.34%, Val AUROC: 0.8610, Partial AUROC: 0.1079

Epoch 16/20

Training Epoch 16: 100%| | 131/131 [04:20<00:00, 1.99s/it, train_loss=0.0518]

Validating Epoch 16: 100%| | 94/94 [00:41<00:00, 2.29it/s, val_loss=0.202]

Epoch [16/20], Train Loss: 0.0518, Val Loss: 0.2024, Val Accuracy: 92.55%, Val AUROC: 0.8864, Partial AUROC: 0.1228

Epoch 17/20

Training Epoch 17: 100%| | 131/131 [04:14<00:00, 1.94s/it, train_loss=0.06]

Validating Epoch 17: 100%| | 94/94 [00:40<00:00, 2.31it/s, val_loss=0.217]

Epoch [17/20], Train Loss: 0.0600, Val Loss: 0.2168, Val Accuracy: 91.68%, Val AUROC: 0.8306, Partial AUROC: 0.0952

Epoch 18/20

Training Epoch 18: 100%| | 131/131 [04:21<00:00, 1.99s/it, train_loss=0.0471]

Validating Epoch 18: 100%| | 94/94 [00:41<00:00, 2.27it/s, val_loss=0.223]

Epoch [18/20], Train Loss: 0.0471, Val Loss: 0.2227, Val Accuracy: 92.01%, Val AUROC: 0.8683, Partial AUROC: 0.1082

Epoch 19/20

Training Epoch 19: 100%| | 131/131 [04:16<00:00, 1.96s/it, train_loss=0.0453]

Validating Epoch 19: 100%| | 94/94 [00:40<00:00, 2.33it/s, val_loss=0.238]

Epoch [19/20], Train Loss: 0.0453, Val Loss: 0.2382, Val Accuracy: 92.68%, Val AUROC: 0.8362, Partial AUROC: 0.0833

Epoch 20/20

Training Epoch 20: 100%| | 131/131 [04:20<00:00, 1.99s/it, train_loss=0.0331]


Validating Epoch 20: 100%| | 94/94 [00:41<00:00, 2.27it/s, val_loss=0.241]

Epoch [20/20], Train Loss: 0.0331, Val Loss: 0.2410, Val Accuracy: 92.01%, Val AUROC: 0.8047, Partial AUROC: 0.0742

Best Epoch: 16, Best Validation Loss: 0.2024

Training Complete





Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Class 0 | 0.98 | 0.94 | 0.96 | 1431 |
| Class 1 | 0.24 | 0.47 | 0.32 | 59 |
| accuracy | | | 0.92 | 1490 |
| macro avg | 0.61 | 0.71 | 0.64 | 1490 |
| weighted avg | 0.95 | 0.92 | 0.93 | 1490 |

```
[10]: CustomImageFeatureEfficientNet(
  (efficientnet): Sequential(
    (0): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Sequential(
        (0): MBConv(
          (block): Sequential(
            (0): Conv2dNormActivation(
              (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
```

```

        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (2): Conv2dNormActivation(
        (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.0, mode=row)
    )
    (2): Sequential(
        (0): MBConv(
            (block): Sequential(
                (0): Conv2dNormActivation(
                    (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(
                (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
                (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```



```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.025, mode=row)
)
(3): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
    )
  )

```

```

        (1): Conv2dNormActivation(
          (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
          (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
    )
    (1): MBCConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(

```

```

        (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
)
(4): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
  )
  (1): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
)
(2): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))

```

```

        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
)
)
(5): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
  )
  (1): MBConv(

```

```

        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
      )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
      )
    )
  )
)

```

```

        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.125, mode=row)
    )
  )
  (6): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )

```

```

    )
    (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)

```



```

        (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
)
(3): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
    )
    )
    (7): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
track_running_stats=True)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      )
      (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
    )
    )
    (8): Conv2dNormActivation(
      (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
  )
)

```

```

        (1): AdaptiveAvgPool2d(output_size=1)
    )
    (fc_image): Linear(in_features=1280, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

0.7 Model 3

```

[8]: model3 = CustomImageFeatureEfficientNet(feature_input_size=9) # Assuming 9
    ↪ features for metadata
model3.to(device)
# Initialize optimizer
optimizer = optim.Adam(model3.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_sizes = 16

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[9]: train_dataloader = DataLoader(train_dataset, batch_size=batch_sizes,
    ↪ shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_sizes, shuffle=True)

```

```

[10]: train_and_validate(model3, train_dataloader, val_dataloader, criterion,
    ↪ optimizer, epochs, device )

```

Epoch 1/20

Training Epoch 1: 100%| | 131/131 [04:29<00:00, 2.06s/it,
train_loss=0.443]

Validating Epoch 1: 100%| | 94/94 [00:40<00:00, 2.30it/s,
val_loss=0.174]

Epoch [1/20], Train Loss: 0.4428, Val Loss: 0.1745, Val Accuracy: 92.15%, Val AUROC: 0.8655, Partial AUROC: 0.0987

Epoch 2/20

Training Epoch 2: 100%| | 131/131 [04:28<00:00, 2.05s/it, train_loss=0.347]

Validating Epoch 2: 100%| | 94/94 [00:42<00:00, 2.21it/s, val_loss=0.191]

Epoch [2/20], Train Loss: 0.3468, Val Loss: 0.1911, Val Accuracy: 93.76%, Val AUROC: 0.8673, Partial AUROC: 0.1102

Epoch 3/20

Training Epoch 3: 100%| | 131/131 [04:34<00:00, 2.09s/it, train_loss=0.222]

Validating Epoch 3: 100%| | 94/94 [00:41<00:00, 2.26it/s, val_loss=0.247]

Epoch [3/20], Train Loss: 0.2223, Val Loss: 0.2474, Val Accuracy: 92.01%, Val AUROC: 0.8157, Partial AUROC: 0.0739

Epoch 4/20

Training Epoch 4: 100%| | 131/131 [04:27<00:00, 2.04s/it, train_loss=0.192]

Validating Epoch 4: 100%| | 94/94 [00:42<00:00, 2.23it/s, val_loss=0.429]

Epoch [4/20], Train Loss: 0.1918, Val Loss: 0.4288, Val Accuracy: 82.89%, Val AUROC: 0.8685, Partial AUROC: 0.1105

Epoch 5/20

Training Epoch 5: 100%| | 131/131 [04:34<00:00, 2.10s/it, train_loss=0.152]

Validating Epoch 5: 100%| | 94/94 [00:41<00:00, 2.27it/s, val_loss=0.312]

Epoch [5/20], Train Loss: 0.1519, Val Loss: 0.3117, Val Accuracy: 85.70%, Val AUROC: 0.8659, Partial AUROC: 0.1042

Epoch 6/20

Training Epoch 6: 100%| | 131/131 [04:28<00:00, 2.05s/it, train_loss=0.116]

Validating Epoch 6: 100%| | 94/94 [00:42<00:00, 2.22it/s, val_loss=0.169]

Epoch [6/20], Train Loss: 0.1155, Val Loss: 0.1691, Val Accuracy: 95.50%, Val AUROC: 0.8651, Partial AUROC: 0.0967

Epoch 7/20

Training Epoch 7: 100%| | 131/131 [04:36<00:00, 2.11s/it, train_loss=0.0595]

Validating Epoch 7: 100%| | 94/94 [00:41<00:00, 2.25it/s, val_loss=0.214]

Epoch [7/20], Train Loss: 0.0595, Val Loss: 0.2139, Val Accuracy: 93.36%, Val AUROC: 0.8953, Partial AUROC: 0.1278

Epoch 8/20

Training Epoch 8: 100%| | 131/131 [04:29<00:00, 2.06s/it, train_loss=0.0747]

Validating Epoch 8: 100%| | 94/94 [00:41<00:00, 2.26it/s, val_loss=0.168]

Epoch [8/20], Train Loss: 0.0747, Val Loss: 0.1682, Val Accuracy: 95.17%, Val AUROC: 0.9203, Partial AUROC: 0.1443

Epoch 9/20

Training Epoch 9: 100%| | 131/131 [04:34<00:00, 2.09s/it, train_loss=0.0608]

Validating Epoch 9: 100%| | 94/94 [00:44<00:00, 2.12it/s, val_loss=0.272]

Epoch [9/20], Train Loss: 0.0608, Val Loss: 0.2721, Val Accuracy: 93.42%, Val AUROC: 0.8302, Partial AUROC: 0.0850

Epoch 10/20

Training Epoch 10: 100%| | 131/131 [04:27<00:00, 2.04s/it, train_loss=0.0518]

Validating Epoch 10: 100%| | 94/94 [00:41<00:00, 2.25it/s, val_loss=0.24]

Epoch [10/20], Train Loss: 0.0518, Val Loss: 0.2395, Val Accuracy: 92.08%, Val AUROC: 0.8827, Partial AUROC: 0.1185

Epoch 11/20

Training Epoch 11: 100%| | 131/131 [04:34<00:00, 2.10s/it, train_loss=0.0532]

Validating Epoch 11: 100%| | 94/94 [00:42<00:00, 2.20it/s, val_loss=0.302]

Epoch [11/20], Train Loss: 0.0532, Val Loss: 0.3017, Val Accuracy: 93.83%, Val AUROC: 0.8263, Partial AUROC: 0.0818

Epoch 12/20

Training Epoch 12: 100%| | 131/131 [04:29<00:00, 2.06s/it, train_loss=0.0844]

Validating Epoch 12: 100%| | 94/94 [00:43<00:00, 2.16it/s, val_loss=0.315]


Epoch [12/20], Train Loss: 0.0844, Val Loss: 0.3155, Val Accuracy: 96.11%, Val AUROC: 0.8433, Partial AUROC: 0.0862

Epoch 13/20

Training Epoch 13: 100%| | 131/131 [04:35<00:00, 2.11s/it, train_loss=0.0379]

Validating Epoch 13: 100%| | 94/94 [00:41<00:00, 2.24it/s, val_loss=0.391]

Epoch [13/20], Train Loss: 0.0379, Val Loss: 0.3914, Val Accuracy: 88.39%, Val AUCROC: 0.8864, Partial AUCROC: 0.1184
Early stopping triggered at epoch 13
Best Epoch: 8, Best Validation Loss: 0.1682
Training Complete



Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Class 0 | 0.99 | 0.89 | 0.94 | 1431 |
| Class 1 | 0.22 | 0.76 | 0.34 | 59 |
| accuracy | | | 0.88 | 1490 |
| macro avg | 0.60 | 0.83 | 0.64 | 1490 |
| weighted avg | 0.96 | 0.88 | 0.91 | 1490 |

```
[10]: CustomImageFeatureEfficientNet(
      (efficientnet): Sequential(
        (0): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Sequential(
            (0): MBConv(
              (block): Sequential(
                (0): Conv2dNormActivation(
                  (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
                  (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                  (2): SiLU(inplace=True)
                )
                (1): SqueezeExcitation(
                  (avgpool): AdaptiveAvgPool2d(output_size=1)
                  (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
                  (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
                  (activation): SiLU(inplace=True)
                  (scale_activation): Sigmoid()
                )
                (2): Conv2dNormActivation(
                  (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
                  (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
              )
            )
            (stochastic_depth): StochasticDepth(p=0.0, mode=row)
          )
        )
      )
      (2): Sequential(
```

```

(0): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
      (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(

```



```

        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.025, mode=row)
    )
    )
    (3): Sequential(
        (0): MBConv(
            (block): Sequential(
                (0): Conv2dNormActivation(
                    (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                    (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                    (2): SiLU(inplace=True)
                )
                (1): Conv2dNormActivation(
                    (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
                    (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                    (2): SiLU(inplace=True)
                )
                (2): SqueezeExcitation(
                    (avgpool): AdaptiveAvgPool2d(output_size=1)
                    (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
                    (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
                    (activation): SiLU(inplace=True)
                    (scale_activation): Sigmoid()
                )
                (3): Conv2dNormActivation(
                    (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                    (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
            )
        )
    )

```

```

        (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
    )
    (1): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
  )
  (4): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),

```

```

padding=(1, 1), groups=240, bias=False)
    (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
    (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
    )
    (1): MBConv(
    (block): Sequential(
    (0): Conv2dNormActivation(
    (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
    (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
    (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
    )
    (2): MBConv(
    (block): Sequential(
    (0): Conv2dNormActivation(
    (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
    (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
    (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
    )
    )
    (5): Sequential(
    (0): MBConv(
    (block): Sequential(
    (0): Conv2dNormActivation(
    (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
    )
)

```

```

        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
    )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.125, mode=row)
    )
  )
  (6): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(

```

```

        (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
        (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)

```

```

        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
)
(3): MBCConv(

```



```

        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
    )
  (7): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1152, bias=False)
          (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
  )
)
(8): Conv2dNormActivation(
  (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): SiLU(inplace=True)
)
)
(1): AdaptiveAvgPool2d(output_size=1)
)
(fc_image): Linear(in_features=1280, out_features=512, bias=True)
(fc_metadata): Linear(in_features=9, out_features=128, bias=True)
(dropout): Dropout(p=0.5, inplace=False)
(fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

[]:

[]: