

# Week\_9\_WinningModel

November 4, 2024

```
[1]: # Standard Libraries
import io
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Deep Learning and PyTorch
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import models

# Image Processing
from PIL import Image
from torchvision import transforms, models

# File Handling
import h5py

# Metrics and Evaluation
from sklearn.metrics import classification_report, roc_auc_score, roc_curve, auc

# Progress Visualization
from tqdm import tqdm
```

## 0.1 Create Custom Dataset

```
[2]: class MultiInputDataset(Dataset):
    def __init__(self, hdf5_file, csv_file, transform=None):
        # Open the HDF5 file with error handling
        try:
            self.hdf5_file = h5py.File(hdf5_file, 'r') # Read-only mode
        except Exception as e:
            raise IOError(f"Could not open HDF5 file: {hdf5_file}. Error: {e}")

        # Read the CSV file containing image labels and additional features
```

```

try:
    self.labels_df = pd.read_csv(csv_file)
except Exception as e:
    raise IOError(f"Could not read CSV file: {csv_file}. Error: {e}")

# Ensure that all image IDs from the CSV are present in the HDF5 file
self.image_ids = self.labels_df['isic_id'].values
for image_id in self.image_ids:
    if str(image_id) not in self.hdf5_file.keys():
        raise ValueError(f"Image id {image_id} not found in HDF5 file.")

# Store any transformations to be applied to the images
self.transform = transform

def __len__(self):
    # Return the total number of samples in the dataset
    return len(self.labels_df)

def __getitem__(self, idx):
    # Get the image ID from the CSV file based on index
    image_id = str(self.labels_df.iloc[idx]['isic_id'])

    # Load the image data from the HDF5 file
    image_bytes = self.hdf5_file[image_id][()]

    # Convert the image bytes to a PIL Image
    image = Image.open(io.BytesIO(image_bytes))

    # Apply any specified transformations to the image
    if self.transform:
        image = self.transform(image)

    # Retrieve the label
    label = torch.tensor(self.labels_df.iloc[idx]['target'], dtype=torch.
↪long) # Adjust dtype if needed

    # Retrieve other features, excluding 'isic_id' and 'target'
    other_variables = self.labels_df.iloc[idx].drop(['isic_id', 'target']).
↪values.astype(float)

    # Convert other variables (metadata) to a tensor
    metadata_tensor = torch.tensor(other_variables, dtype=torch.float32)

    # Return the image, metadata, and label
    return image, metadata_tensor, label

```

```
[3]: def get_train_transform(resize_size=(224, 224), crop_size=128,
    ↪rotation_degree=10, normalize_means=(0.5, 0.5, 0.5), normalize_stds=(0.5, 0.
    ↪5, 0.5)):
    """
    Returns the transformations for the training dataset, including data
    ↪augmentation.

    Args:
        resize_size (tuple): The size to resize the image before cropping.
        crop_size (int): The size of the random crop.
        rotation_degree (int): Maximum degree for random rotation.
        normalize_means (tuple): Means for normalization.
        normalize_stds (tuple): Standard deviations for normalization.

    Returns:
        transforms.Compose: The composed transformations for the training set.
    """
    return transforms.Compose([
        transforms.Resize(resize_size), # Resize to specified size
        transforms.RandomResizedCrop(crop_size, scale=(0.8, 1.0)), # Random
    ↪crop with scale
        transforms.RandomRotation(rotation_degree), # Randomly rotate images
        transforms.ToTensor(), # Convert image to PyTorch tensor
        transforms.Normalize(normalize_means, normalize_stds) # Normalize with
    ↪specified means and stds
    ])

def get_normal_transform(resize_size=(224, 224), normalize_means=(0.5, 0.5, 0.
    ↪5), normalize_stds=(0.5, 0.5, 0.5)):
    """
    Returns the transformations for the validation/test dataset (without data
    ↪augmentation).

    Args:
        resize_size (tuple): The size to resize the image.
        normalize_means (tuple): Means for normalization.
        normalize_stds (tuple): Standard deviations for normalization.

    Returns:
        transforms.Compose: The composed transformations for the validation/
    ↪test set.
    """
    return transforms.Compose([
        transforms.Resize(resize_size), # Resize to specified size
        transforms.ToTensor(), # Convert image to PyTorch tensor
```

```

        transforms.Normalize(normalize_means, normalize_stds) # Normalize with
↪specified means and stds
    ])

```

## 0.2 Train DataLoader

```

[4]: device = "cuda" if torch.cuda.is_available() else "cpu"

```

## 0.3 Model Building

```

[5]: # CNN Model
class CustomImageFeatureCNN2(nn.Module):
    def __init__(self, feature_input_size, input_image_size=(128, 128)):
        super(CustomImageFeatureCNN2, self).__init__()

        # Image CNN with Batch Normalization
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3,
↪padding=1)
        self.bn1 = nn.BatchNorm2d(32) # Batch normalization after conv1

        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(64) # Batch normalization after conv2

        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
        self.bn3 = nn.BatchNorm2d(128) # Batch normalization after conv3

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2) # 2x2 Max pooling

        # Dynamically calculate the flattened size of the feature map
        self.flattened_size = self._get_flattened_size(input_image_size)

        # Fully connected layer after the CNN layers
        self.fc_image = nn.Linear(self.flattened_size, 512)

        # Fully connected layer for metadata (feature data)
        self.fc_metadata = nn.Linear(feature_input_size, 128)

        # Dropout layer to prevent overfitting
        self.dropout = nn.Dropout(0.5) # 50% dropout

        # Final fully connected layer for binary classification (combined image
↪+ feature input)
        self.fc_combined = nn.Linear(512 + 128, 1) # Change 2 to 1 for binary
↪classification

    def _get_flattened_size(self, input_image_size):

```

```

        # Forward pass a dummy image to get the size of the flattened features
        dummy_image = torch.zeros(1, 3, *input_image_size) # Batch size of 1, 3
        ↪ channels (RGB), and input size
        x = self.pool(F.relu(self.bn1(self.conv1(dummy_image))))
        x = self.pool(F.relu(self.bn2(self.conv2(x))))
        x = self.pool(F.relu(self.bn3(self.conv3(x))))
        return x.view(-1).shape[0] # Flatten and return the size

    def forward(self, image, metadata):
        # Forward pass for the image through the CNN
        x = self.pool(F.relu(self.bn1(self.conv1(image)))) # Conv layer 1 with
        ↪ ReLU, BatchNorm, MaxPool
        x = self.pool(F.relu(self.bn2(self.conv2(x)))) # Conv layer 2 with
        ↪ ReLU, BatchNorm, MaxPool
        x = self.pool(F.relu(self.bn3(self.conv3(x)))) # Conv layer 3 with
        ↪ ReLU, BatchNorm, MaxPool

        # Flatten the feature map to feed into fully connected layer
        x = x.view(x.size(0), -1) # Flatten feature maps into a 1D vector
        image_features = F.relu(self.fc_image(x))

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            f"Batch sizes do not match! Image batch size: {image_features.
        ↪ shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
        ↪ dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

        # If you're using BCELoss, uncomment the next line to apply sigmoid
        output = torch.sigmoid(output)

        return output

```

```

[6]: # Resnet Model
class CustomImageFeatureResNet(nn.Module):
    def __init__(self, feature_input_size, pretrained=True):
        super(CustomImageFeatureResNet, self).__init__()

```

```

    # Load a pretrained ResNet model for image feature extraction (ResNet18,
    ↪ in this case)
    resnet = models.resnet18(pretrained=pretrained) # Change to resnet50,
    ↪ resnet101 as needed
    self.resnet = nn.Sequential(*list(resnet.children())[:-1]) # Remove
    ↪ the final classification layer

    # The output of ResNet18's last conv layer is 512-dimensional (for
    ↪ ResNet50, it would be 2048)
    self.fc_image = nn.Linear(resnet.fc.in_features, 512) # Adjust if
    ↪ using ResNet50

    # Fully connected layer for metadata (feature data)
    self.fc_metadata = nn.Linear(feature_input_size, 128)

    # Dropout layer to prevent overfitting
    self.dropout = nn.Dropout(0.5) # 50% dropout

    # Final fully connected layer for binary classification (combined image
    ↪ + feature input)
    self.fc_combined = nn.Linear(512 + 128, 1) # For binary classification

    def forward(self, image, metadata):
        # Forward pass for the image through the ResNet (without the final
        ↪ classification layer)
        x = self.resnet(image) # ResNet feature extraction
        x = x.view(x.size(0), -1) # Flatten the ResNet output
        image_features = F.relu(self.fc_image(x))

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            ↪ f"Batch sizes do not match! Image batch size: {image_features.
            ↪ shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
        ↪ dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

```

```

# If you're using BCELoss, uncomment the next line to apply sigmoid
output = torch.sigmoid(output)

return output

```

```

[7]: # EfficientNet Model
class CustomImageFeatureEfficientNet(nn.Module):
    def __init__(self, feature_input_size, pretrained=True):
        super(CustomImageFeatureEfficientNet, self).__init__()

        # Load a pretrained EfficientNet model for image feature extraction
        ↪(EfficientNet-B0 in this case)
        efficientnet = models.efficientnet_b0(pretrained=pretrained) # You can
        ↪change this to another EfficientNet version like B1 or B7
        self.efficientnet = nn.Sequential(*list(efficientnet.children())[:-1])
        ↪# Remove the final classification layer

        # The output of EfficientNet-B0's last conv layer is 1280-dimensional
        self.fc_image = nn.Linear(1280, 512) # Reduce dimension to match your
        ↪custom architecture

        # Fully connected layer for metadata (feature data)
        self.fc_metadata = nn.Linear(feature_input_size, 128)

        # Dropout layer to prevent overfitting
        self.dropout = nn.Dropout(0.5) # 50% dropout

        # Final fully connected layer for binary classification (combined image
        ↪+ feature input)
        self.fc_combined = nn.Linear(512 + 128, 1) # For binary classification

    def forward(self, image, metadata):
        # Forward pass for the image through EfficientNet (without the final
        ↪classification layer)
        x = self.efficientnet(image) # EfficientNet feature extraction
        x = x.view(x.size(0), -1) # Flatten the EfficientNet output
        image_features = F.relu(self.fc_image(x))

        # Process metadata (feature data)
        metadata_features = F.relu(self.fc_metadata(metadata))

        # Ensure the batch sizes are consistent
        assert image_features.shape[0] == metadata_features.shape[0], \
            f"Batch sizes do not match! Image batch size: {image_features.
        ↪shape[0]}, Metadata batch size: {metadata_features.shape[0]}"

```

```

        # Concatenate image features and metadata features
        combined_features = torch.cat((image_features, metadata_features),
        ↪dim=1)

        # Dropout and final classification layer
        combined_features = self.dropout(combined_features)
        output = self.fc_combined(combined_features)

        # If you're using BCELoss, uncomment the next line to apply sigmoid
        output = torch.sigmoid(output)

        return output

```

## 0.4 Model Training

```

[8]: # Function to compute partial AUC-above-TPR
def score(solution: np.array, submission: np.array, min_tpr: float = 0.80) ->
    ↪float:
    """
        Compute the partial AUC by focusing on a specific range of true positive
        ↪rates (TPR).

        Args:
            solution (np.array): Ground truth binary labels.
            submission (np.array): Model predictions.
            min_tpr (float): Minimum true positive rate to calculate partial AUC.

        Returns:
            float: The calculated partial AUC.

        Raises:
            ValueError: If the min_tpr is not within a valid range.
    """
    # Rescale the target to handle sklearn limitations and flip the predictions
    v_gt = abs(solution - 1)
    v_pred = -1.0 * submission
    max_fpr = abs(1 - min_tpr)

    # Compute ROC curve using sklearn
    fpr, tpr, _ = roc_curve(v_gt, v_pred)
    if max_fpr is None or max_fpr == 1:
        return auc(fpr, tpr)
    if max_fpr <= 0 or max_fpr > 1:
        raise ValueError(f"Expected min_tpr in range [0, 1), got: {min_tpr}")

    # Interpolate for partial AUC

```



```

stop = np.searchsorted(fpr, max_fpr, "right")
x_interp = [fpr[stop - 1], fpr[stop]]
y_interp = [tpr[stop - 1], tpr[stop]]
tpr = np.append(tpr[:stop], np.interp(max_fpr, x_interp, y_interp))
fpr = np.append(fpr[:stop], max_fpr)
partial_auc = auc(fpr, tpr)

return partial_auc

# Training and validation loop function
def train_and_validate(
    model: nn.Module,
    train_dataloader: torch.utils.data.DataLoader,
    val_dataloader: torch.utils.data.DataLoader,
    criterion: nn.Module,
    optimizer: torch.optim.Optimizer,
    epochs: int,
    device: torch.device,
    best_model_path: str,
    early_stopping_patience: int = 5,
    min_tpr: float = 0.80
) -> nn.Module:
    """
    Train and validate a PyTorch model with early stopping, AUROC, partial AUC,
    and error handling.

    Args:
        model (nn.Module): The model to be trained and validated.
        train_dataloader (torch.utils.data.DataLoader): Dataloader for training
        data.
        val_dataloader (torch.utils.data.DataLoader): Dataloader for validation
        data.
        criterion (nn.Module): Loss function.
        optimizer (torch.optim.Optimizer): Optimizer to update the model.
        epochs (int): Number of training epochs.
        device (torch.device): The device (CPU or GPU) to use.
        early_stopping_patience (int): Early stopping patience.
        min_tpr (float): The minimum true positive rate for calculating partial
        AUC.

    Returns:
        nn.Module: The trained model.
    """
    # Initialize tracking variables
    best_val_loss = float('inf')
    best_epoch = 0

```

```

train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []
early_stopping_counter = 0

# Start the training and validation loop
for epoch in range(epochs):
    print(f'Epoch {epoch + 1}/{epochs}')

    # Training phase
    model.train()
    running_train_loss = 0.0
    correct_train = 0
    total_train = 0
    all_train_labels = []
    all_train_probs = []

    progress_bar = tqdm(train_dataloader, desc=f'Training Epoch {epoch + 1}')

    try:
        # Loop through the training batches
        for i, (image, metadata, labels) in enumerate(progress_bar):
            image, metadata, labels = image.to(device), metadata.
            to(device), labels.float().to(device)
            labels = labels.unsqueeze(1) # Adjust labels to have the right
            shape for binary classification

            optimizer.zero_grad()

            # Forward pass
            probs = model(image, metadata)

            if probs.shape != labels.shape:
                raise ValueError(f"Shape mismatch: Predictions shape {probs.
            shape} does not match labels shape {labels.shape}")

            # Calculate loss and backpropagate
            loss = criterion(probs, labels)
            loss.backward()
            optimizer.step()

            # Update running loss
            running_train_loss += loss.item()

            # Store labels and predictions for accuracy calculations

```

```

        all_train_labels.extend(labels.cpu().detach().numpy())
        all_train_probs.extend(probs.cpu().detach().numpy())

        # Calculate binary predictions for training accuracy
        predicted_train = (probs >= 0.5).float()
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

        # Update progress bar
        progress_bar.set_postfix(train_loss=running_train_loss / (i + 1))

    # Calculate training accuracy and loss
    train_accuracy = 100 * correct_train / total_train
    train_losses.append(running_train_loss / len(train_dataloader))
    train_accuracies.append(train_accuracy)

except ValueError as ve:
    print(f"Error during training loop: {ve}")
    break

# Validation phase
model.eval()
running_val_loss = 0.0
correct = 0
total = 0
all_labels = []
all_probs = []

progress_bar = tqdm(val_dataloader, desc=f'Validating Epoch {epoch + 1}')

with torch.no_grad():
    try:
        # Loop through the validation batches
        for i, (images, metadata, labels) in enumerate(progress_bar):
            images, metadata, labels = images.to(device), metadata.
            to(device), labels.float().to(device)
            labels = labels.unsqueeze(1)

            probs = model(images, metadata)

            loss = criterion(probs, labels)
            running_val_loss += loss.item()

            all_labels.extend(labels.cpu().detach().numpy())
            all_probs.extend(probs.cpu().detach().numpy())

```

```

        # Calculate binary predictions for validation accuracy
        predicted = (probs >= 0.5).float()
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        progress_bar.set_postfix(val_loss=running_val_loss / (i + 1))

    val_accuracy = 100 * correct / total
    val_loss = running_val_loss / len(val_dataloader)
    val accuracies.append(val_accuracy)
    val_losses.append(val_loss)

    # Calculate AUROC
    try:
        valid_auroc = roc_auc_score(all_labels, all_probs)
    except ValueError as ve:
        print(f"AUROC Calculation Error: {ve}")
        valid_auroc = 0.0

    # Calculate partial AUC-above-TPR
    try:
        partial_auroc = score(np.array(all_labels), np.
        array(all_probs), min_tpr=min_tpr)
    except ValueError as ve:
        print(f"Partial AUC Calculation Error: {ve}")
        partial_auroc = 0.0

    print(f'Epoch [{epoch}/{epochs}], Train Loss: {train_losses[-1]:
    .4f}, Val Loss: {val_loss:.4f}, '
          f'Val Accuracy: {val_accuracy:.2f}%, Val AUROC:
    {valid_auroc:.4f}, Partial AUROC: {partial_auroc:.4f}')

    # Early stopping based on validation loss
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_epoch = epoch + 1
        early_stopping_counter = 0
        torch.save(model.state_dict(), best_model_path)
    else:
        early_stopping_counter += 1

    if early_stopping_counter >= early_stopping_patience:
        print(f"Early stopping triggered at epoch {epoch}")
        break

```

```

        except Exception as e:
            print(f"Error during validation loop: {e}")
            break

    print(f"Best Epoch: {best_epoch}, Best Validation Loss: {best_val_loss:.
↪4f}")
    print('Training Complete')

    # Plot training and validation loss
    plt.figure(figsize=(10, 5))
    plt.plot(train_losses, label='Train Loss')
    plt.plot(val_losses, label='Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
    plt.show()

    # Plot training and validation accuracy
    plt.figure(figsize=(10, 5))
    plt.plot(train_accuracies, label='Train Accuracy')
    plt.plot(val_accuracies, label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.title('Training and Validation Accuracy')
    plt.legend()
    plt.show()

    # Generate classification report
    try:
        print("Classification Report:")
        print(classification_report(all_labels, (np.array(all_probs) >= 0.5).
↪astype(int), target_names=['Class 0', 'Class 1']))
    except Exception as e:
        print(f"Error generating classification report: {e}")

    return model

```

```

[9]: # Initialize the dataset
CNN_train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.
↪hdf5', csv_file='../data/processed/processed-train-metadata1.csv',
↪transform=get_train_transform(resize_size=(128,128)))
CNN_val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.
↪hdf5', csv_file='../data/processed/processed-validation-metadata1.csv',
↪transform=get_normal_transform(resize_size=(128,128)))
# Create a DataLoader

```

```
CNN_train_dataloader = DataLoader(CNN_train_dataset, batch_size=64,
    ↪shuffle=True)
CNN_val_dataloader = DataLoader(CNN_val_dataset, batch_size=64, shuffle=True)
```

```
[10]: # Initialize the dataset
resnet_train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.
    ↪hdf5', csv_file='../data/processed/processed-train-metadata1.csv',
    ↪transform=get_train_transform(resize_size=(225,225)))
resnet_val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.
    ↪hdf5', csv_file='../data/processed/processed-validation-metadata1.csv',
    ↪transform=get_normal_transform(resize_size=(225,225)))
# Create a DataLoader
resnet_train_dataloader = DataLoader(resnet_train_dataset, batch_size=64,
    ↪shuffle=True)
resnet_val_dataloader = DataLoader(resnet_val_dataset, batch_size=64,
    ↪shuffle=True)
```

```
[11]: # Initialize the dataset
effnet_train_dataset = MultiInputDataset(hdf5_file='../data/raw/train_images.
    ↪hdf5', csv_file='../data/processed/processed-train-metadata1.csv',
    ↪transform=get_train_transform(resize_size=(224,224)))
effnet_val_dataset = MultiInputDataset(hdf5_file='../data/raw/validation_image.
    ↪hdf5', csv_file='../data/processed/processed-validation-metadata1.csv',
    ↪transform=get_normal_transform(resize_size=(224,224)))
# Create a DataLoader
effnet_train_dataloader = DataLoader(effnet_train_dataset, batch_size=64,
    ↪shuffle=True)
effnet_val_dataloader = DataLoader(effnet_val_dataset, batch_size=64,
    ↪shuffle=True)
```

## 0.5 Model 1

```
[12]: model1 = CustomImageFeatureCNN2(feature_input_size=9) # Assuming 9 features
    ↪for metadata
model1.to(device)
# Initialize optimizer
optimizer = optim.Adam(model1.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
best_model_path = "best_model1.pth"
```

```
[13]: train_and_validate(model1,CNN_train_dataloader, CNN_val_dataloader, criterion,
    ↪optimizer, epochs, device ,best_model_path)
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [01:27<00:00, 2.64s/it,  
train\_loss=4.18]

Validating Epoch 1: 100%| | 24/24 [00:25<00:00, 1.06s/it,  
val\_loss=0.423]

Epoch [0/20], Train Loss: 4.1848, Val Loss: 0.4227, Val Accuracy: 90.81%, Val  
AUROC: 0.8038, Partial AUROC: 0.0749

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [01:19<00:00, 2.40s/it,  
train\_loss=0.83]

Validating Epoch 2: 100%| | 24/24 [00:29<00:00, 1.23s/it,  
val\_loss=0.358]

Epoch [1/20], Train Loss: 0.8297, Val Loss: 0.3575, Val Accuracy: 88.93%, Val  
AUROC: 0.7938, Partial AUROC: 0.0599

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [01:45<00:00, 3.19s/it,  
train\_loss=0.458]

Validating Epoch 3: 100%| | 24/24 [00:24<00:00, 1.01s/it,  
val\_loss=0.368]

Epoch [2/20], Train Loss: 0.4581, Val Loss: 0.3682, Val Accuracy: 87.92%, Val  
AUROC: 0.8200, Partial AUROC: 0.0693

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [01:18<00:00, 2.38s/it,  
train\_loss=0.427]

Validating Epoch 4: 100%| | 24/24 [00:24<00:00, 1.04s/it,  
val\_loss=0.313]

Epoch [3/20], Train Loss: 0.4267, Val Loss: 0.3129, Val Accuracy: 88.52%, Val  
AUROC: 0.8165, Partial AUROC: 0.0715

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [01:24<00:00, 2.57s/it,  
train\_loss=0.441]

Validating Epoch 5: 100%| | 24/24 [00:25<00:00, 1.04s/it,  
val\_loss=0.363]

Epoch [4/20], Train Loss: 0.4409, Val Loss: 0.3627, Val Accuracy: 86.51%, Val  
AUROC: 0.8393, Partial AUROC: 0.0821

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [01:22<00:00, 2.50s/it,  
train\_loss=0.421]

Validating Epoch 6: 100%| | 24/24 [00:25<00:00, 1.04s/it,  
val\_loss=0.31]

Epoch [5/20], Train Loss: 0.4213, Val Loss: 0.3103, Val Accuracy: 88.19%, Val  
AUROC: 0.8442, Partial AUROC: 0.0844

Epoch 7/20

Training Epoch 7: 100%| | 33/33 [01:22<00:00, 2.49s/it,  
train\_loss=0.401]

Validating Epoch 7: 100%| | 24/24 [00:26<00:00, 1.10s/it,  
val\_loss=0.291]

Epoch [6/20], Train Loss: 0.4014, Val Loss: 0.2913, Val Accuracy: 88.93%, Val  
AUROC: 0.8538, Partial AUROC: 0.0917

Epoch 8/20

Training Epoch 8: 100%| | 33/33 [01:28<00:00, 2.69s/it,  
train\_loss=0.382]

Validating Epoch 8: 100%| | 24/24 [00:49<00:00, 2.06s/it,  
val\_loss=0.269]

Epoch [7/20], Train Loss: 0.3820, Val Loss: 0.2685, Val Accuracy: 88.86%, Val  
AUROC: 0.8516, Partial AUROC: 0.0890

Epoch 9/20

Training Epoch 9: 100%| | 33/33 [01:22<00:00, 2.50s/it,  
train\_loss=0.367]

Validating Epoch 9: 100%| | 24/24 [00:32<00:00, 1.37s/it,  
val\_loss=0.224]

Epoch [8/20], Train Loss: 0.3666, Val Loss: 0.2236, Val Accuracy: 91.14%, Val  
AUROC: 0.8624, Partial AUROC: 0.1020

Epoch 10/20

Training Epoch 10: 100%| | 33/33 [02:13<00:00, 4.04s/it,  
train\_loss=0.361]

Validating Epoch 10: 100%| | 24/24 [00:42<00:00, 1.77s/it,  
val\_loss=0.289]

Epoch [9/20], Train Loss: 0.3605, Val Loss: 0.2888, Val Accuracy: 88.79%, Val  
AUROC: 0.8273, Partial AUROC: 0.0790

Epoch 11/20

Training Epoch 11: 100%| | 33/33 [02:02<00:00, 3.72s/it,  
train\_loss=0.346]

Validating Epoch 11: 100%| | 24/24 [00:34<00:00, 1.45s/it,  
val\_loss=0.209]

Epoch [10/20], Train Loss: 0.3460, Val Loss: 0.2089, Val Accuracy: 92.75%, Val  
AUROC: 0.8748, Partial AUROC: 0.1060

Epoch 12/20

Training Epoch 12: 100%| | 33/33 [01:57<00:00, 3.55s/it,  
train\_loss=0.334]

Validating Epoch 12: 100%| | 24/24 [00:54<00:00, 2.26s/it,  
val\_loss=0.238]

Epoch [11/20], Train Loss: 0.3341, Val Loss: 0.2383, Val Accuracy: 91.28%, Val  
AUROC: 0.8747, Partial AUROC: 0.1048



Epoch 13/20

Training Epoch 13: 100%| | 33/33 [02:03<00:00, 3.73s/it,  
train\_loss=0.34]

Validating Epoch 13: 100%| | 24/24 [00:35<00:00, 1.49s/it,  
val\_loss=0.187]

Epoch [12/20], Train Loss: 0.3399, Val Loss: 0.1871, Val Accuracy: 93.36%, Val  
AUROC: 0.8728, Partial AUROC: 0.1066

Epoch 14/20

Training Epoch 14: 100%| | 33/33 [01:52<00:00, 3.42s/it,  
train\_loss=0.34]

Validating Epoch 14: 100%| | 24/24 [00:35<00:00, 1.48s/it,  
val\_loss=0.277]

Epoch [13/20], Train Loss: 0.3402, Val Loss: 0.2766, Val Accuracy: 89.13%, Val  
AUROC: 0.8791, Partial AUROC: 0.1095

Epoch 15/20

Training Epoch 15: 100%| | 33/33 [01:56<00:00, 3.55s/it,  
train\_loss=0.333]

Validating Epoch 15: 100%| | 24/24 [00:36<00:00, 1.51s/it,  
val\_loss=0.283]

Epoch [14/20], Train Loss: 0.3327, Val Loss: 0.2826, Val Accuracy: 87.92%, Val  
AUROC: 0.8737, Partial AUROC: 0.1035

Epoch 16/20

Training Epoch 16: 100%| | 33/33 [01:59<00:00, 3.62s/it,  
train\_loss=0.324]

Validating Epoch 16: 100%| | 24/24 [00:44<00:00, 1.84s/it,  
val\_loss=0.287]

Epoch [15/20], Train Loss: 0.3239, Val Loss: 0.2866, Val Accuracy: 88.05%, Val  
AUROC: 0.8761, Partial AUROC: 0.1093

Epoch 17/20

Training Epoch 17: 100%| | 33/33 [02:14<00:00, 4.08s/it,  
train\_loss=0.307]

Validating Epoch 17: 100%| | 24/24 [00:35<00:00, 1.50s/it,  
val\_loss=0.269]

Epoch [16/20], Train Loss: 0.3074, Val Loss: 0.2691, Val Accuracy: 88.72%, Val  
AUROC: 0.8842, Partial AUROC: 0.1106



Epoch 18/20

Training Epoch 18: 100%| | 33/33 [01:55<00:00, 3.51s/it,  
train\_loss=0.3]

Validating Epoch 18: 100%| | 24/24 [00:36<00:00, 1.52s/it,  
val\_loss=0.248]

Epoch [17/20], Train Loss: 0.2999, Val Loss: 0.2485, Val Accuracy: 90.81%, Val  
AUROC: 0.8738, Partial AUROC: 0.1037

Early stopping triggered at epoch 17  
Best Epoch: 13, Best Validation Loss: 0.1871  
Training Complete



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Class 0	0.99	0.92	0.95	1431
Class 1	0.25	0.68	0.37	59
accuracy			0.91	1490
macro avg	0.62	0.80	0.66	1490
weighted avg	0.96	0.91	0.93	1490

```
[13]: CustomImageFeatureCNN2(
      (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
      (fc_image): Linear(in_features=32768, out_features=512, bias=True)
      (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
      (dropout): Dropout(p=0.5, inplace=False)
      (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )
```

## 0.6 Model 2

```
[14]: model2 = CustomImageFeatureCNN2(feature_input_size=9)  # Assuming 9 features
      ↪for metadata
model2.to(device)
      # Initialize optimizer
optimizer = optim.SGD(model2.parameters(), lr=0.001)
      # Define the loss function with the class weights
criterion = nn.BCELoss()  # Binary classification loss
      # Set the number of epochs
epochs = 20
best_model_path = "best_model2.pth"
```

```
[15]: train_and_validate(model2,CNN_train_dataloader, CNN_val_dataloader, criterion,
      ↪optimizer, epochs, device,best_model_path )
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [01:57<00:00, 3.57s/it,  
train\_loss=0.628]

Validating Epoch 1: 100%| | 24/24 [00:41<00:00, 1.75s/it,  
val\_loss=0.57]

Epoch [0/20], Train Loss: 0.6278, Val Loss: 0.5701, Val Accuracy: 95.97%, Val AUROC: 0.7021, Partial AUROC: 0.0547

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [01:52<00:00, 3.41s/it, train\_loss=0.599]

Validating Epoch 2: 100%| | 24/24 [00:53<00:00, 2.22s/it, val\_loss=0.496]

Epoch [1/20], Train Loss: 0.5989, Val Loss: 0.4958, Val Accuracy: 94.43%, Val AUROC: 0.7610, Partial AUROC: 0.0619

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [02:07<00:00, 3.87s/it, train\_loss=0.577]

Validating Epoch 3: 100%| | 24/24 [00:38<00:00, 1.62s/it, val\_loss=0.512]

Epoch [2/20], Train Loss: 0.5775, Val Loss: 0.5117, Val Accuracy: 88.93%, Val AUROC: 0.7735, Partial AUROC: 0.0635

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [01:56<00:00, 3.52s/it, train\_loss=0.555]

Validating Epoch 4: 100%| | 24/24 [00:37<00:00, 1.58s/it, val\_loss=0.482]

Epoch [3/20], Train Loss: 0.5548, Val Loss: 0.4821, Val Accuracy: 88.19%, Val AUROC: 0.7844, Partial AUROC: 0.0660

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [02:03<00:00, 3.74s/it, train\_loss=0.54]

Validating Epoch 5: 100%| | 24/24 [00:38<00:00, 1.61s/it, val\_loss=0.495]

Epoch [4/20], Train Loss: 0.5395, Val Loss: 0.4947, Val Accuracy: 84.09%, Val AUROC: 0.7909, Partial AUROC: 0.0696

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [01:54<00:00, 3.47s/it, train\_loss=0.525]

Validating Epoch 6: 100%| | 24/24 [01:05<00:00, 2.74s/it, val\_loss=0.507]

Epoch [5/20], Train Loss: 0.5253, Val Loss: 0.5066, Val Accuracy: 82.21%, Val AUROC: 0.7962, Partial AUROC: 0.0693

Epoch 7/20

Training Epoch 7: 100%| | 33/33 [01:56<00:00, 3.53s/it, train\_loss=0.506]

Validating Epoch 7: 100%| | 24/24 [00:37<00:00, 1.56s/it, val\_loss=0.451]

Epoch [6/20], Train Loss: 0.5061, Val Loss: 0.4514, Val Accuracy: 84.90%, Val AUROC: 0.8026, Partial AUROC: 0.0705

Epoch 8/20

Training Epoch 8: 100%| | 33/33 [01:55<00:00, 3.49s/it, train\_loss=0.497]

Validating Epoch 8: 100%| | 24/24 [00:39<00:00, 1.64s/it, val\_loss=0.492]

Epoch [7/20], Train Loss: 0.4968, Val Loss: 0.4923, Val Accuracy: 81.41%, Val AUROC: 0.8021, Partial AUROC: 0.0676

Epoch 9/20

Training Epoch 9: 100%| | 33/33 [02:03<00:00, 3.75s/it, train\_loss=0.485]

Validating Epoch 9: 100%| | 24/24 [00:38<00:00, 1.59s/it, val\_loss=0.445]

Epoch [8/20], Train Loss: 0.4849, Val Loss: 0.4452, Val Accuracy: 84.43%, Val AUROC: 0.8091, Partial AUROC: 0.0722

Epoch 10/20

Training Epoch 10: 100%| | 33/33 [01:55<00:00, 3.51s/it, train\_loss=0.471]

Validating Epoch 10: 100%| | 24/24 [01:06<00:00, 2.79s/it, val\_loss=0.484]

Epoch [9/20], Train Loss: 0.4712, Val Loss: 0.4843, Val Accuracy: 81.48%, Val AUROC: 0.8107, Partial AUROC: 0.0731

Epoch 11/20

Training Epoch 11: 100%| | 33/33 [01:58<00:00, 3.58s/it, train\_loss=0.477]

Validating Epoch 11: 100%| | 24/24 [00:38<00:00, 1.61s/it, val\_loss=0.479]

Epoch [10/20], Train Loss: 0.4768, Val Loss: 0.4790, Val Accuracy: 81.14%, Val AUROC: 0.8137, Partial AUROC: 0.0748

Epoch 12/20

Training Epoch 12: 100%| | 33/33 [01:58<00:00, 3.59s/it, train\_loss=0.468]

Validating Epoch 12: 100%| | 24/24 [00:41<00:00, 1.72s/it, val\_loss=0.448]

Epoch [11/20], Train Loss: 0.4682, Val Loss: 0.4476, Val Accuracy: 83.83%, Val AUROC: 0.8137, Partial AUROC: 0.0701

Epoch 13/20

Training Epoch 13: 100%| | 33/33 [01:55<00:00, 3.50s/it, train\_loss=0.464]

Validating Epoch 13: 100%| | 24/24 [00:38<00:00, 1.61s/it, val\_loss=0.452]

Epoch [12/20], Train Loss: 0.4639, Val Loss: 0.4522, Val Accuracy: 82.75%, Val AUROC: 0.8171, Partial AUROC: 0.0733

Epoch 14/20

Training Epoch 14: 100%| | 33/33 [01:56<00:00, 3.54s/it, train\_loss=0.457]

Validating Epoch 14: 100%| | 24/24 [00:53<00:00, 2.21s/it, val\_loss=0.42]

Epoch [13/20], Train Loss: 0.4574, Val Loss: 0.4204, Val Accuracy: 84.83%, Val AUROC: 0.8185, Partial AUROC: 0.0722

Epoch 15/20

Training Epoch 15: 100%| | 33/33 [02:17<00:00, 4.18s/it, train\_loss=0.453]

Validating Epoch 15: 100%| | 24/24 [00:47<00:00, 1.99s/it, val\_loss=0.428]

Epoch [14/20], Train Loss: 0.4527, Val Loss: 0.4277, Val Accuracy: 85.44%, Val AUROC: 0.8213, Partial AUROC: 0.0800

Epoch 16/20

Training Epoch 16: 100%| | 33/33 [02:25<00:00, 4.42s/it, train\_loss=0.45]

Validating Epoch 16: 100%| | 24/24 [00:45<00:00, 1.91s/it, val\_loss=0.399]

Epoch [15/20], Train Loss: 0.4504, Val Loss: 0.3990, Val Accuracy: 85.97%, Val AUROC: 0.8235, Partial AUROC: 0.0766

Epoch 17/20

Training Epoch 17: 100%| | 33/33 [02:16<00:00, 4.14s/it, train\_loss=0.444]

Validating Epoch 17: 100%| | 24/24 [00:45<00:00, 1.90s/it, val\_loss=0.407]

Epoch [16/20], Train Loss: 0.4443, Val Loss: 0.4070, Val Accuracy: 85.23%, Val AUROC: 0.8227, Partial AUROC: 0.0734

Epoch 18/20

Training Epoch 18: 100%| | 33/33 [02:47<00:00, 5.07s/it, train\_loss=0.439]

Validating Epoch 18: 100%| | 24/24 [00:43<00:00, 1.83s/it, val\_loss=0.406]

Epoch [17/20], Train Loss: 0.4394, Val Loss: 0.4057, Val Accuracy: 85.70%, Val AUROC: 0.8271, Partial AUROC: 0.0771

Epoch 19/20

Training Epoch 19: 100%| | 33/33 [01:26<00:00, 2.64s/it, train\_loss=0.439]

Validating Epoch 19: 100%| | 24/24 [00:26<00:00, 1.12s/it, val\_loss=0.384]

Epoch [18/20], Train Loss: 0.4394, Val Loss: 0.3842, Val Accuracy: 87.32%, Val AUROC: 0.8297, Partial AUROC: 0.0800

Epoch 20/20

Training Epoch 20: 100%| | 33/33 [01:33<00:00, 2.84s/it,  
train\_loss=0.432]


Validating Epoch 20: 100%| | 24/24 [00:32<00:00, 1.34s/it,  
val\_loss=0.415]

Epoch [19/20], Train Loss: 0.4317, Val Loss: 0.4151, Val Accuracy: 84.63%, Val AUROC: 0.8273, Partial AUROC: 0.0768

Best Epoch: 19, Best Validation Loss: 0.3842

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.85	0.91	1431
Class 1	0.16	0.68	0.26	59
accuracy			0.85	1490
macro avg	0.57	0.77	0.59	1490
weighted avg	0.95	0.85	0.89	1490

```
[15]: CustomImageFeatureCNN2(
    (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc_image): Linear(in_features=32768, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
```



)

## 0.7 Model 3

```
[16]: model3 = CustomImageFeatureCNN2(feature_input_size=9) # Assuming 9 features,
      ↪for metadata
      model3.to(device)
      # Initialize optimizer
      optimizer = optim.SGD(model3.parameters(), lr=0.0001, weight_decay=1e-4)
      # Define the loss function with the class weights
      criterion = nn.BCELoss() # Binary classification loss
      # Set the number of epochs
      epochs = 20
      batch_size = 32
      best_model_path = "best_model3.pth"

[17]: CNN_train_dataloader = DataLoader(CNN_train_dataset, batch_size=batch_size,
      ↪shuffle=True)
      CNN_val_dataloader = DataLoader(CNN_val_dataset, batch_size=batch_size,
      ↪shuffle=True)

[18]: train_and_validate(model3, CNN_train_dataloader, CNN_val_dataloader, criterion,
      ↪optimizer, epochs, device, best_model_path )
```

Epoch 1/20

Training Epoch 1: 100%| | 66/66 [01:26<00:00, 1.31s/it,  
train\_loss=0.64]

Validating Epoch 1: 100%| | 47/47 [00:37<00:00, 1.25it/s,  
val\_loss=0.569]

Epoch [0/20], Train Loss: 0.6399, Val Loss: 0.5689, Val Accuracy: 95.91%, Val  
AUROC: 0.6144, Partial AUROC: 0.0319

Epoch 2/20

Training Epoch 2: 100%| | 66/66 [02:27<00:00, 2.24s/it,  
train\_loss=0.629]

Validating Epoch 2: 100%| | 47/47 [00:41<00:00, 1.12it/s,  
val\_loss=0.536]

Epoch [1/20], Train Loss: 0.6285, Val Loss: 0.5364, Val Accuracy: 96.04%, Val  
AUROC: 0.6355, Partial AUROC: 0.0391

Epoch 3/20

Training Epoch 3: 100%| | 66/66 [01:44<00:00, 1.58s/it,  
train\_loss=0.621]

Validating Epoch 3: 100%| | 47/47 [00:21<00:00, 2.23it/s,  
val\_loss=0.522]

Epoch [2/20], Train Loss: 0.6212, Val Loss: 0.5219, Val Accuracy: 96.04%, Val AUROC: 0.6689, Partial AUROC: 0.0431

Epoch 4/20

Training Epoch 4: 100%| | 66/66 [01:39<00:00, 1.51s/it, train\_loss=0.613]

Validating Epoch 4: 100%| | 47/47 [00:42<00:00, 1.11it/s, val\_loss=0.51]

Epoch [3/20], Train Loss: 0.6132, Val Loss: 0.5100, Val Accuracy: 96.11%, Val AUROC: 0.6878, Partial AUROC: 0.0471

Epoch 5/20

Training Epoch 5: 100%| | 66/66 [02:02<00:00, 1.86s/it, train\_loss=0.602]

Validating Epoch 5: 100%| | 47/47 [00:39<00:00, 1.19it/s, val\_loss=0.514]

Epoch [4/20], Train Loss: 0.6023, Val Loss: 0.5141, Val Accuracy: 96.17%, Val AUROC: 0.7038, Partial AUROC: 0.0465

Epoch 6/20

Training Epoch 6: 100%| | 66/66 [02:29<00:00, 2.26s/it, train\_loss=0.602]

Validating Epoch 6: 100%| | 47/47 [00:25<00:00, 1.88it/s, val\_loss=0.5]

Epoch [5/20], Train Loss: 0.6020, Val Loss: 0.4998, Val Accuracy: 96.17%, Val AUROC: 0.7137, Partial AUROC: 0.0486

Epoch 7/20

Training Epoch 7: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.602]

Validating Epoch 7: 100%| | 47/47 [00:21<00:00, 2.23it/s, val\_loss=0.497]

Epoch [6/20], Train Loss: 0.6018, Val Loss: 0.4968, Val Accuracy: 96.04%, Val AUROC: 0.7155, Partial AUROC: 0.0510

Epoch 8/20

Training Epoch 8: 100%| | 66/66 [01:16<00:00, 1.17s/it, train\_loss=0.599]

Validating Epoch 8: 100%| | 47/47 [00:23<00:00, 2.01it/s, val\_loss=0.491]

Epoch [7/20], Train Loss: 0.5989, Val Loss: 0.4915, Val Accuracy: 95.50%, Val AUROC: 0.7197, Partial AUROC: 0.0518

Epoch 9/20

Training Epoch 9: 100%| | 66/66 [01:12<00:00, 1.09s/it, train\_loss=0.593]

Validating Epoch 9: 100%| | 47/47 [00:25<00:00, 1.86it/s, val\_loss=0.484]

Epoch [8/20], Train Loss: 0.5926, Val Loss: 0.4844, Val Accuracy: 95.44%, Val AUROC: 0.7217, Partial AUROC: 0.0523

Epoch 10/20

Training Epoch 10: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.584]

Validating Epoch 10: 100%| | 47/47 [00:19<00:00, 2.40it/s, val\_loss=0.46]

Epoch [9/20], Train Loss: 0.5841, Val Loss: 0.4598, Val Accuracy: 95.97%, Val AUROC: 0.7264, Partial AUROC: 0.0528

Epoch 11/20

Training Epoch 11: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.584]

Validating Epoch 11: 100%| | 47/47 [00:22<00:00, 2.07it/s, val\_loss=0.48]

Epoch [10/20], Train Loss: 0.5844, Val Loss: 0.4802, Val Accuracy: 94.77%, Val AUROC: 0.7367, Partial AUROC: 0.0524

Epoch 12/20

Training Epoch 12: 100%| | 66/66 [01:11<00:00, 1.08s/it, train\_loss=0.585]

Validating Epoch 12: 100%| | 47/47 [00:18<00:00, 2.48it/s, val\_loss=0.485]

Epoch [11/20], Train Loss: 0.5852, Val Loss: 0.4847, Val Accuracy: 93.22%, Val AUROC: 0.7396, Partial AUROC: 0.0542

Epoch 13/20

Training Epoch 13: 100%| | 66/66 [01:37<00:00, 1.48s/it, train\_loss=0.58]

Validating Epoch 13: 100%| | 47/47 [00:25<00:00, 1.86it/s, val\_loss=0.459]

Epoch [12/20], Train Loss: 0.5796, Val Loss: 0.4590, Val Accuracy: 94.90%, Val AUROC: 0.7420, Partial AUROC: 0.0535

Epoch 14/20

Training Epoch 14: 100%| | 66/66 [01:17<00:00, 1.17s/it, train\_loss=0.573]

Validating Epoch 14: 100%| | 47/47 [00:22<00:00, 2.12it/s, val\_loss=0.465]

Epoch [13/20], Train Loss: 0.5730, Val Loss: 0.4654, Val Accuracy: 93.42%, Val AUROC: 0.7402, Partial AUROC: 0.0540

Epoch 15/20

Training Epoch 15: 100%| | 66/66 [01:23<00:00, 1.27s/it, train\_loss=0.574]

Validating Epoch 15: 100%| | 47/47 [00:17<00:00, 2.62it/s, val\_loss=0.456]

Epoch [14/20], Train Loss: 0.5740, Val Loss: 0.4562, Val Accuracy: 93.29%, Val AUROC: 0.7453, Partial AUROC: 0.0545

Epoch 16/20

Training Epoch 16: 100%| | 66/66 [01:13<00:00, 1.11s/it, train\_loss=0.57]

Validating Epoch 16: 100%| | 47/47 [00:18<00:00, 2.48it/s, val\_loss=0.464]

Epoch [15/20], Train Loss: 0.5700, Val Loss: 0.4640, Val Accuracy: 92.55%, Val AUROC: 0.7510, Partial AUROC: 0.0556

Epoch 17/20

Training Epoch 17: 100%| | 66/66 [01:10<00:00, 1.07s/it, train\_loss=0.571]

Validating Epoch 17: 100%| | 47/47 [00:22<00:00, 2.13it/s, val\_loss=0.453]

Epoch [16/20], Train Loss: 0.5710, Val Loss: 0.4531, Val Accuracy: 92.48%, Val AUROC: 0.7481, Partial AUROC: 0.0542

Epoch 18/20

Training Epoch 18: 100%| | 66/66 [01:16<00:00, 1.16s/it, train\_loss=0.569]

Validating Epoch 18: 100%| | 47/47 [00:18<00:00, 2.50it/s, val\_loss=0.451]

Epoch [17/20], Train Loss: 0.5693, Val Loss: 0.4507, Val Accuracy: 92.42%, Val AUROC: 0.7562, Partial AUROC: 0.0552

Epoch 19/20

Training Epoch 19: 100%| | 66/66 [01:37<00:00, 1.47s/it, train\_loss=0.562]

Validating Epoch 19: 100%| | 47/47 [00:17<00:00, 2.66it/s, val\_loss=0.437]

Epoch [18/20], Train Loss: 0.5619, Val Loss: 0.4367, Val Accuracy: 93.22%, Val AUROC: 0.7566, Partial AUROC: 0.0554

Epoch 20/20



Training Epoch 20: 100%| | 66/66 [01:16<00:00, 1.16s/it, train\_loss=0.556]

Validating Epoch 20: 100%| | 47/47 [00:25<00:00, 1.83it/s, val\_loss=0.445]

Epoch [19/20], Train Loss: 0.5557, Val Loss: 0.4446, Val Accuracy: 92.15%, Val AUROC: 0.7592, Partial AUROC: 0.0553

Best Epoch: 19, Best Validation Loss: 0.4367

Training Complete



Classification Report:

	precision	recall	f1-score	support
Class 0	0.97	0.95	0.96	1431
Class 1	0.19	0.31	0.24	59
accuracy	0.92			1490

macro avg	0.58	0.63	0.60	1490
weighted avg	0.94	0.92	0.93	1490

```
[18]: CustomImageFeatureCNN2(
      (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
      (fc_image): Linear(in_features=32768, out_features=512, bias=True)
      (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
      (dropout): Dropout(p=0.5, inplace=False)
      (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )
```

## 0.8 Model 4

```
[19]: model4 = CustomImageFeatureResNet(feature_input_size=9) # Assuming 9 features
      ↪for metadata
model4.to(device)
# Initialize optimizer
optimizer = optim.Adam(model4.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
best_model_path = "best_model4.pth"
```

```
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
[20]: train_and_validate(model4,resnet_train_dataloader, resnet_val_dataloader,↵
      ↵criterion, optimizer, epochs, device, best_model_path )
```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [01:41<00:00, 3.08s/it,  
train\_loss=0.474]

Validating Epoch 1: 100%| | 24/24 [01:28<00:00, 3.67s/it,  
val\_loss=0.567]

Epoch [0/20], Train Loss: 0.4739, Val Loss: 0.5675, Val Accuracy: 80.54%, Val  
AUROC: 0.7634, Partial AUROC: 0.0691

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [01:56<00:00, 3.53s/it,  
train\_loss=0.351]

Validating Epoch 2: 100%| | 24/24 [01:22<00:00, 3.43s/it,  
val\_loss=0.336]

Epoch [1/20], Train Loss: 0.3508, Val Loss: 0.3355, Val Accuracy: 88.99%, Val  
AUROC: 0.7413, Partial AUROC: 0.0491

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [02:23<00:00, 4.34s/it,  
train\_loss=0.302]

Validating Epoch 3: 100%| | 24/24 [01:28<00:00, 3.67s/it,  
val\_loss=0.398]

Epoch [2/20], Train Loss: 0.3015, Val Loss: 0.3976, Val Accuracy: 79.46%, Val  
AUROC: 0.8102, Partial AUROC: 0.0917

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [01:41<00:00, 3.07s/it,  
train\_loss=0.249]

Validating Epoch 4: 100%| | 24/24 [01:26<00:00, 3.61s/it,  
val\_loss=0.32]

Epoch [3/20], Train Loss: 0.2492, Val Loss: 0.3196, Val Accuracy: 86.31%, Val  
AUROC: 0.7434, Partial AUROC: 0.0684

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [01:41<00:00, 3.07s/it,  
train\_loss=0.228]

Validating Epoch 5: 100%| | 24/24 [01:27<00:00, 3.65s/it,  
val\_loss=0.259]

Epoch [4/20], Train Loss: 0.2279, Val Loss: 0.2589, Val Accuracy: 91.21%, Val  
AUROC: 0.8385, Partial AUROC: 0.0922

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [01:44<00:00, 3.15s/it,  
train\_loss=0.158]

Validating Epoch 6: 100%| | 24/24 [01:33<00:00, 3.90s/it,  
val\_loss=0.366]

Epoch [5/20], Train Loss: 0.1581, Val Loss: 0.3662, Val Accuracy: 86.78%, Val  
AUROC: 0.7717, Partial AUROC: 0.0584  
Epoch 7/20

Training Epoch 7: 100%| | 33/33 [01:47<00:00, 3.26s/it,  
train\_loss=0.17]

Validating Epoch 7: 100%| | 24/24 [01:25<00:00, 3.57s/it,  
val\_loss=0.398]

Epoch [6/20], Train Loss: 0.1701, Val Loss: 0.3978, Val Accuracy: 77.92%, Val  
AUROC: 0.7380, Partial AUROC: 0.0591  
Epoch 8/20

Training Epoch 8: 100%| | 33/33 [01:43<00:00, 3.15s/it,  
train\_loss=0.171]

Validating Epoch 8: 100%| | 24/24 [01:23<00:00, 3.46s/it,  
val\_loss=0.368]

Epoch [7/20], Train Loss: 0.1706, Val Loss: 0.3679, Val Accuracy: 84.77%, Val  
AUROC: 0.7922, Partial AUROC: 0.0746  
Epoch 9/20

Training Epoch 9: 100%| | 33/33 [01:44<00:00, 3.15s/it,  
train\_loss=0.116]

Validating Epoch 9: 100%| | 24/24 [01:31<00:00, 3.83s/it,  
val\_loss=0.456]

Epoch [8/20], Train Loss: 0.1159, Val Loss: 0.4561, Val Accuracy: 80.00%, Val  
AUROC: 0.8201, Partial AUROC: 0.0955  
Epoch 10/20

Training Epoch 10: 100%| | 33/33 [01:42<00:00, 3.12s/it,  
train\_loss=0.118]

Validating Epoch 10: 100%| | 24/24 [01:23<00:00, 3.49s/it,  
val\_loss=0.258]

Epoch [9/20], Train Loss: 0.1177, Val Loss: 0.2578, Val Accuracy: 89.40%, Val  
AUROC: 0.7430, Partial AUROC: 0.0672  
Epoch 11/20

Training Epoch 11: 100%| | 33/33 [01:52<00:00, 3.42s/it,  
train\_loss=0.108]

Validating Epoch 11: 100%| | 24/24 [01:13<00:00, 3.08s/it,  
val\_loss=0.527]

Epoch [10/20], Train Loss: 0.1078, Val Loss: 0.5266, Val Accuracy: 73.29%, Val  
AUROC: 0.7982, Partial AUROC: 0.0709  
Epoch 12/20

Training Epoch 12: 100%| | 33/33 [01:54<00:00, 3.48s/it,  
train\_loss=0.0786]



Validating Epoch 12: 100%|        | 24/24 [01:45<00:00, 4.39s/it,  
val\_loss=0.183]

Epoch [11/20], Train Loss: 0.0786, Val Loss: 0.1828, Val Accuracy: 93.02%, Val  
AUROC: 0.8238, Partial AUROC: 0.0840

Epoch 13/20

Training Epoch 13: 100%|        | 33/33 [03:01<00:00, 5.49s/it,  
train\_loss=0.0714]

Validating Epoch 13: 100%|        | 24/24 [02:20<00:00, 5.87s/it,  
val\_loss=0.292]

Epoch [12/20], Train Loss: 0.0714, Val Loss: 0.2925, Val Accuracy: 90.54%, Val  
AUROC: 0.7638, Partial AUROC: 0.0659

Epoch 14/20

Training Epoch 14: 100%|        | 33/33 [02:38<00:00, 4.81s/it,  
train\_loss=0.0764]

Validating Epoch 14: 100%|        | 24/24 [02:03<00:00, 5.15s/it,  
val\_loss=0.233]

Epoch [13/20], Train Loss: 0.0764, Val Loss: 0.2325, Val Accuracy: 91.74%, Val  
AUROC: 0.7967, Partial AUROC: 0.0725

Epoch 15/20

Training Epoch 15: 100%|        | 33/33 [02:26<00:00, 4.45s/it,  
train\_loss=0.0652]

Validating Epoch 15: 100%|        | 24/24 [02:00<00:00, 5.01s/it,  
val\_loss=0.378]

Epoch [14/20], Train Loss: 0.0652, Val Loss: 0.3783, Val Accuracy: 84.63%, Val  
AUROC: 0.7758, Partial AUROC: 0.0631

Epoch 16/20

Training Epoch 16: 100%|        | 33/33 [02:29<00:00, 4.53s/it,  
train\_loss=0.0704]

Validating Epoch 16: 100%|        | 24/24 [02:06<00:00, 5.27s/it,  
val\_loss=0.18]

Epoch [15/20], Train Loss: 0.0704, Val Loss: 0.1804, Val Accuracy: 93.56%, Val  
AUROC: 0.8074, Partial AUROC: 0.0833

Epoch 17/20

Training Epoch 17: 100%|        | 33/33 [02:20<00:00, 4.27s/it,  
train\_loss=0.0362]

Validating Epoch 17: 100%|        | 24/24 [01:55<00:00, 4.80s/it,  
val\_loss=0.242]

Epoch [16/20], Train Loss: 0.0362, Val Loss: 0.2416, Val Accuracy: 93.22%, Val  
AUROC: 0.7827, Partial AUROC: 0.0742

Epoch 18/20

Training Epoch 18: 100%|        | 33/33 [02:22<00:00, 4.33s/it,  
train\_loss=0.0648]

Validating Epoch 18: 100%| | 24/24 [02:14<00:00, 5.61s/it,  
val\_loss=0.246]

Epoch [17/20], Train Loss: 0.0648, Val Loss: 0.2462, Val Accuracy: 93.15%, Val  
AUROC: 0.7440, Partial AUROC: 0.0623

Epoch 19/20

Training Epoch 19: 100%| | 33/33 [02:14<00:00, 4.09s/it,  
train\_loss=0.0355]

Validating Epoch 19: 100%| | 24/24 [01:29<00:00, 3.73s/it,  
val\_loss=0.416]

Epoch [18/20], Train Loss: 0.0355, Val Loss: 0.4160, Val Accuracy: 82.82%, Val  
AUROC: 0.7396, Partial AUROC: 0.0643

Epoch 20/20


Training Epoch 20: 100%| | 33/33 [02:31<00:00, 4.60s/it,  
train\_loss=0.0653]


Validating Epoch 20: 100%| | 24/24 [02:36<00:00, 6.51s/it,  
val\_loss=0.229]

Epoch [19/20], Train Loss: 0.0653, Val Loss: 0.2292, Val Accuracy: 92.48%, Val  
AUROC: 0.7764, Partial AUROC: 0.0750

Best Epoch: 16, Best Validation Loss: 0.1804

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.95	0.96	1431
Class 1	0.24	0.42	0.31	59
accuracy			0.92	1490
macro avg	0.61	0.68	0.63	1490
weighted avg	0.95	0.92	0.93	1490

```
[20]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(

```

```

        (0): BasicBlock(
          (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    (7): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```

```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (8): AdaptiveAvgPool2d(output_size=(1, 1))
    )
    (fc_image): Linear(in_features=512, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )

```

## 0.9 Model 5

```

[21]: model5 = CustomImageFeatureResNet(feature_input_size=9) # Assuming 9 features
      ↪ for metadata
      model5.to(device)
      # Initialize optimizer
      optimizer = optim.SGD(model5.parameters(), lr=0.001)
      # Define the loss function with the class weights
      criterion = nn.BCELoss() # Binary classification loss
      # Set the number of epochs
      epochs = 20
      best_model_path = "best_model5.pth"

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[22]: train_and_validate(model5, resnet_train_dataloader, resnet_val_dataloader,
      ↪ criterion, optimizer, epochs, device, best_model_path )

```

Epoch 1/20

Training Epoch 1: 100%| | 33/33 [02:46<00:00, 5.05s/it,  
train\_loss=0.671]

Validating Epoch 1: 100%| | 24/24 [01:58<00:00, 4.95s/it,  
val\_loss=0.57]

Epoch [0/20], Train Loss: 0.6705, Val Loss: 0.5696, Val Accuracy: 96.04%, Val  
AUROC: 0.5256, Partial AUROC: 0.0269

Epoch 2/20

Training Epoch 2: 100%| | 33/33 [03:04<00:00, 5.59s/it,  
train\_loss=0.65]

Validating Epoch 2: 100%| | 24/24 [02:28<00:00, 6.17s/it,  
val\_loss=0.545]

Epoch [1/20], Train Loss: 0.6498, Val Loss: 0.5453, Val Accuracy: 96.04%, Val  
AUROC: 0.5536, Partial AUROC: 0.0261

Epoch 3/20

Training Epoch 3: 100%| | 33/33 [02:30<00:00, 4.55s/it,  
train\_loss=0.629]

Validating Epoch 3: 100%| | 24/24 [01:56<00:00, 4.85s/it,  
val\_loss=0.526]

Epoch [2/20], Train Loss: 0.6292, Val Loss: 0.5257, Val Accuracy: 96.04%, Val  
AUROC: 0.5678, Partial AUROC: 0.0274

Epoch 4/20

Training Epoch 4: 100%| | 33/33 [02:06<00:00, 3.85s/it,  
train\_loss=0.625]

Validating Epoch 4: 100%| | 24/24 [01:27<00:00, 3.65s/it,  
val\_loss=0.522]

Epoch [3/20], Train Loss: 0.6248, Val Loss: 0.5224, Val Accuracy: 96.04%, Val  
AUROC: 0.5870, Partial AUROC: 0.0279

Epoch 5/20

Training Epoch 5: 100%| | 33/33 [01:43<00:00, 3.13s/it,  
train\_loss=0.608]

Validating Epoch 5: 100%| | 24/24 [01:28<00:00, 3.68s/it,  
val\_loss=0.516]

Epoch [4/20], Train Loss: 0.6084, Val Loss: 0.5164, Val Accuracy: 96.04%, Val  
AUROC: 0.6024, Partial AUROC: 0.0290

Epoch 6/20

Training Epoch 6: 100%| | 33/33 [01:48<00:00, 3.28s/it,  
train\_loss=0.61]

Validating Epoch 6: 100%| | 24/24 [01:31<00:00, 3.79s/it,  
val\_loss=0.511]

Epoch [5/20], Train Loss: 0.6098, Val Loss: 0.5107, Val Accuracy: 96.04%, Val  
AUROC: 0.6146, Partial AUROC: 0.0297

Epoch 7/20

Training Epoch 7: 100%| | 33/33 [01:38<00:00, 2.99s/it,  
train\_loss=0.602]

Validating Epoch 7: 100%| | 24/24 [01:26<00:00, 3.58s/it,  
val\_loss=0.505]

Epoch [6/20], Train Loss: 0.6022, Val Loss: 0.5051, Val Accuracy: 96.04%, Val  
AUROC: 0.6275, Partial AUROC: 0.0301

Epoch 8/20

Training Epoch 8: 100%| | 33/33 [01:37<00:00, 2.94s/it,  
train\_loss=0.601]

Validating Epoch 8: 100%| | 24/24 [01:27<00:00, 3.64s/it,  
val\_loss=0.511]

Epoch [7/20], Train Loss: 0.6014, Val Loss: 0.5106, Val Accuracy: 96.04%, Val  
AUROC: 0.6324, Partial AUROC: 0.0306

Epoch 9/20

Training Epoch 9: 100%| | 33/33 [01:42<00:00, 3.11s/it,  
train\_loss=0.578]

Validating Epoch 9: 100%| | 24/24 [01:26<00:00, 3.60s/it,  
val\_loss=0.506]

Epoch [8/20], Train Loss: 0.5781, Val Loss: 0.5055, Val Accuracy: 96.04%, Val  
AUROC: 0.6483, Partial AUROC: 0.0316

Epoch 10/20

Training Epoch 10: 100%| | 33/33 [01:36<00:00, 2.93s/it,  
train\_loss=0.575]

Validating Epoch 10: 100%| | 24/24 [01:20<00:00, 3.36s/it,  
val\_loss=0.502]

Epoch [9/20], Train Loss: 0.5752, Val Loss: 0.5022, Val Accuracy: 95.97%, Val  
AUROC: 0.6594, Partial AUROC: 0.0322

Epoch 11/20

Training Epoch 11: 100%| | 33/33 [01:37<00:00, 2.96s/it,  
train\_loss=0.574]

Validating Epoch 11: 100%| | 24/24 [01:27<00:00, 3.64s/it,  
val\_loss=0.499]

Epoch [10/20], Train Loss: 0.5743, Val Loss: 0.4994, Val Accuracy: 95.97%, Val  
AUROC: 0.6636, Partial AUROC: 0.0331

Epoch 12/20

Training Epoch 12: 100%| | 33/33 [01:39<00:00, 3.00s/it,  
train\_loss=0.566]

Validating Epoch 12: 100%| | 24/24 [01:31<00:00, 3.82s/it,  
val\_loss=0.501]

Epoch [11/20], Train Loss: 0.5661, Val Loss: 0.5010, Val Accuracy: 95.97%, Val  
AUROC: 0.6720, Partial AUROC: 0.0331



Epoch 13/20

Training Epoch 13: 100%| | 33/33 [01:36<00:00, 2.91s/it,  
train\_loss=0.553]

Validating Epoch 13: 100%| | 24/24 [01:20<00:00, 3.34s/it,  
val\_loss=0.502]

Epoch [12/20], Train Loss: 0.5532, Val Loss: 0.5025, Val Accuracy: 95.91%, Val  
AUROC: 0.6800, Partial AUROC: 0.0363

Epoch 14/20

Training Epoch 14: 100%| | 33/33 [01:34<00:00, 2.87s/it,  
train\_loss=0.552]

Validating Epoch 14: 100%| | 24/24 [01:24<00:00, 3.51s/it,  
val\_loss=0.495]

Epoch [13/20], Train Loss: 0.5518, Val Loss: 0.4952, Val Accuracy: 95.97%, Val  
AUROC: 0.6843, Partial AUROC: 0.0357

Epoch 15/20

Training Epoch 15: 100%| | 33/33 [01:36<00:00, 2.92s/it,  
train\_loss=0.541]

Validating Epoch 15: 100%| | 24/24 [01:31<00:00, 3.82s/it,  
val\_loss=0.494]

Epoch [14/20], Train Loss: 0.5409, Val Loss: 0.4942, Val Accuracy: 95.84%, Val  
AUROC: 0.6906, Partial AUROC: 0.0363

Epoch 16/20

Training Epoch 16: 100%| | 33/33 [01:36<00:00, 2.93s/it,  
train\_loss=0.531]

Validating Epoch 16: 100%| | 24/24 [01:26<00:00, 3.62s/it,  
val\_loss=0.488]

Epoch [15/20], Train Loss: 0.5306, Val Loss: 0.4885, Val Accuracy: 95.64%, Val  
AUROC: 0.6979, Partial AUROC: 0.0383

Epoch 17/20

Training Epoch 17: 100%| | 33/33 [01:42<00:00, 3.10s/it,  
train\_loss=0.528]

Validating Epoch 17: 100%| | 24/24 [01:19<00:00, 3.30s/it,  
val\_loss=0.484]

Epoch [16/20], Train Loss: 0.5278, Val Loss: 0.4835, Val Accuracy: 95.64%, Val  
AUROC: 0.7056, Partial AUROC: 0.0387

Epoch 18/20

Training Epoch 18: 100%| | 33/33 [01:43<00:00, 3.14s/it,  
train\_loss=0.522]

Validating Epoch 18: 100%| | 24/24 [01:26<00:00, 3.62s/it,  
val\_loss=0.482]

Epoch [17/20], Train Loss: 0.5221, Val Loss: 0.4823, Val Accuracy: 95.50%, Val  
AUROC: 0.7090, Partial AUROC: 0.0380

Epoch 19/20

Training Epoch 19: 100%| | 33/33 [02:06<00:00, 3.84s/it,  
train\_loss=0.513]

Validating Epoch 19: 100%| | 24/24 [02:38<00:00, 6.59s/it,  
val\_loss=0.48]

Epoch [18/20], Train Loss: 0.5129, Val Loss: 0.4799, Val Accuracy: 95.44%, Val  
AUROC: 0.7179, Partial AUROC: 0.0388

Epoch 20/20


Training Epoch 20: 100%| | 33/33 [04:20<00:00, 7.89s/it,  
train\_loss=0.504]


Validating Epoch 20: 100%| | 24/24 [03:41<00:00, 9.23s/it,  
val\_loss=0.477]

Epoch [19/20], Train Loss: 0.5041, Val Loss: 0.4769, Val Accuracy: 95.50%, Val  
AUROC: 0.7221, Partial AUROC: 0.0398

Best Epoch: 20, Best Validation Loss: 0.4769

Training Complete





#### Classification Report:

	precision	recall	f1-score	support
Class 0	0.96	0.99	0.98	1431
Class 1	0.30	0.10	0.15	59
accuracy			0.96	1490
macro avg	0.63	0.55	0.56	1490
weighted avg	0.94	0.96	0.94	1490

```
[22]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
  bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
  track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
  ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
  1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
  track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(

```

```

        (0): BasicBlock(
          (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (1): BasicBlock(
          (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (7): Sequential(
        (0): BasicBlock(
          (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (1): BasicBlock(
          (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```

```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (8): AdaptiveAvgPool2d(output_size=(1, 1))
    )
    (fc_image): Linear(in_features=512, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )

```

## 0.10 Model 6

```

[23]: model6 = CustomImageFeatureResNet(feature_input_size=9) # Assuming 9 features
      ↪ for metadata
model6.to(device)
# Initialize optimizer
optimizer = optim.SGD(model6.parameters(), lr=0.0001, weight_decay=1e-4)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
batch_size = 32
best_model_path = "best_model6.pth"

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[24]: resnet_train_dataloader = DataLoader(resnet_train_dataset,
      ↪ batch_size=batch_size, shuffle=True)

```

```
resnet_val_dataloader = DataLoader(resnet_val_dataset, batch_size=batch_size,
↪shuffle=True)
```

```
[25]: train_and_validate(model6, resnet_train_dataloader, resnet_val_dataloader,
↪criterion, optimizer, epochs, device, best_model_path )
```

Epoch 1/20

Training Epoch 1: 100%| | 66/66 [04:24<00:00, 4.01s/it,  
train\_loss=0.736]

Validating Epoch 1: 100%| | 47/47 [03:12<00:00, 4.10s/it,  
val\_loss=0.725]

Epoch [0/20], Train Loss: 0.7365, Val Loss: 0.7252, Val Accuracy: 25.10%, Val  
AUROC: 0.5898, Partial AUROC: 0.0350

Epoch 2/20

Training Epoch 2: 100%| | 66/66 [03:58<00:00, 3.61s/it,  
train\_loss=0.707]

Validating Epoch 2: 100%| | 47/47 [01:51<00:00, 2.37s/it,  
val\_loss=0.685]

Epoch [1/20], Train Loss: 0.7070, Val Loss: 0.6855, Val Accuracy: 58.93%, Val  
AUROC: 0.5779, Partial AUROC: 0.0277

Epoch 3/20

Training Epoch 3: 100%| | 66/66 [02:02<00:00, 1.86s/it,  
train\_loss=0.677]

Validating Epoch 3: 100%| | 47/47 [01:38<00:00, 2.09s/it,  
val\_loss=0.662]

Epoch [2/20], Train Loss: 0.6774, Val Loss: 0.6620, Val Accuracy: 78.19%, Val  
AUROC: 0.5870, Partial AUROC: 0.0325

Epoch 4/20

Training Epoch 4: 100%| | 66/66 [03:15<00:00, 2.96s/it,  
train\_loss=0.671]

Validating Epoch 4: 100%| | 47/47 [02:23<00:00, 3.05s/it,  
val\_loss=0.641]

Epoch [3/20], Train Loss: 0.6712, Val Loss: 0.6412, Val Accuracy: 89.06%, Val  
AUROC: 0.5699, Partial AUROC: 0.0250

Epoch 5/20

Training Epoch 5: 100%| | 66/66 [02:03<00:00, 1.87s/it,  
train\_loss=0.658]

Validating Epoch 5: 100%| | 47/47 [02:26<00:00, 3.11s/it,  
val\_loss=0.623]

Epoch [4/20], Train Loss: 0.6582, Val Loss: 0.6230, Val Accuracy: 93.29%, Val  
AUROC: 0.5709, Partial AUROC: 0.0260

Epoch 6/20


Training Epoch 6: 100%|        | 66/66 [01:52<00:00, 1.70s/it,  
 train\_loss=0.654]  
 Validating Epoch 6: 100%|        | 47/47 [01:51<00:00, 2.37s/it,  
 val\_loss=0.615]  
  
 Epoch [5/20], Train Loss: 0.6540, Val Loss: 0.6150, Val Accuracy: 94.16%, Val  
 AUROC: 0.5798, Partial AUROC: 0.0272  
 Epoch 7/20  
  
 Training Epoch 7: 100%|        | 66/66 [03:23<00:00, 3.08s/it,  
 train\_loss=0.646]  
 Validating Epoch 7: 100%|        | 47/47 [02:47<00:00, 3.56s/it,  
 val\_loss=0.604]  
  
 Epoch [6/20], Train Loss: 0.6458, Val Loss: 0.6037, Val Accuracy: 95.23%, Val  
 AUROC: 0.5714, Partial AUROC: 0.0252  
 Epoch 8/20  
  
 Training Epoch 8: 100%|        | 66/66 [03:55<00:00, 3.57s/it,  
 train\_loss=0.641]  
 Validating Epoch 8: 100%|        | 47/47 [02:46<00:00, 3.54s/it,  
 val\_loss=0.595]  
  
 Epoch [7/20], Train Loss: 0.6410, Val Loss: 0.5946, Val Accuracy: 95.70%, Val  
 AUROC: 0.5800, Partial AUROC: 0.0254  
 Epoch 9/20  
  
 Training Epoch 9: 100%|        | 66/66 [03:54<00:00, 3.56s/it,  
 train\_loss=0.639]  
 Validating Epoch 9: 100%|        | 47/47 [02:10<00:00, 2.77s/it,  
 val\_loss=0.588]  
  
 Epoch [8/20], Train Loss: 0.6388, Val Loss: 0.5883, Val Accuracy: 95.84%, Val  
 AUROC: 0.5818, Partial AUROC: 0.0252  
 Epoch 10/20  
  
 Training Epoch 10: 100%|        | 66/66 [01:42<00:00, 1.56s/it,  
 train\_loss=0.63]  
 Validating Epoch 10: 100%|        | 47/47 [01:23<00:00, 1.79s/it,  
 val\_loss=0.58]  
  
 Epoch [9/20], Train Loss: 0.6303, Val Loss: 0.5796, Val Accuracy: 96.04%, Val  
 AUROC: 0.5932, Partial AUROC: 0.0282  
 Epoch 11/20  
  
 Training Epoch 11: 100%|        | 66/66 [02:21<00:00, 2.15s/it,  
 train\_loss=0.632]  
 Validating Epoch 11: 100%|        | 47/47 [02:06<00:00, 2.70s/it,  
 val\_loss=0.572]  
  
 Epoch [10/20], Train Loss: 0.6322, Val Loss: 0.5721, Val Accuracy: 96.11%, Val  
 AUROC: 0.5961, Partial AUROC: 0.0276  
 Epoch 12/20



Training Epoch 12: 100%| | 66/66 [02:24<00:00, 2.18s/it,  
 train\_loss=0.627]  
 Validating Epoch 12: 100%| | 47/47 [01:39<00:00, 2.12s/it,  
 val\_loss=0.562]  
  
 Epoch [11/20], Train Loss: 0.6275, Val Loss: 0.5623, Val Accuracy: 96.11%, Val  
 AUROC: 0.5920, Partial AUROC: 0.0263  
 Epoch 13/20  
  
 Training Epoch 13: 100%| | 66/66 [02:33<00:00, 2.32s/it,  
 train\_loss=0.625]  
 Validating Epoch 13: 100%| | 47/47 [02:05<00:00, 2.66s/it,  
 val\_loss=0.562]  
  
 Epoch [12/20], Train Loss: 0.6254, Val Loss: 0.5616, Val Accuracy: 96.17%, Val  
 AUROC: 0.6024, Partial AUROC: 0.0301  
 Epoch 14/20  
  
 Training Epoch 14: 100%| | 66/66 [02:30<00:00, 2.28s/it,  
 train\_loss=0.621]  
 Validating Epoch 14: 100%| | 47/47 [02:00<00:00, 2.57s/it,  
 val\_loss=0.56]  
  
 Epoch [13/20], Train Loss: 0.6207, Val Loss: 0.5602, Val Accuracy: 96.11%, Val  
 AUROC: 0.6045, Partial AUROC: 0.0297  
 Epoch 15/20  
  
 Training Epoch 15: 100%| | 66/66 [01:46<00:00, 1.61s/it,  
 train\_loss=0.623]  
 Validating Epoch 15: 100%| | 47/47 [01:25<00:00, 1.82s/it,  
 val\_loss=0.558]  
  
 Epoch [14/20], Train Loss: 0.6232, Val Loss: 0.5582, Val Accuracy: 96.11%, Val  
 AUROC: 0.6009, Partial AUROC: 0.0285  
 Epoch 16/20  
  
 Training Epoch 16: 100%| | 66/66 [01:38<00:00, 1.49s/it,  
 train\_loss=0.621]  
 Validating Epoch 16: 100%| | 47/47 [01:22<00:00, 1.76s/it,  
 val\_loss=0.556]  
  
 Epoch [15/20], Train Loss: 0.6209, Val Loss: 0.5564, Val Accuracy: 96.11%, Val  
 AUROC: 0.6137, Partial AUROC: 0.0289  
 Epoch 17/20  
  
 Training Epoch 17: 100%| | 66/66 [01:48<00:00, 1.64s/it,  
 train\_loss=0.618]  
 Validating Epoch 17: 100%| | 47/47 [01:18<00:00, 1.67s/it,  
 val\_loss=0.554]  
  
 Epoch [16/20], Train Loss: 0.6179, Val Loss: 0.5540, Val Accuracy: 96.11%, Val  
 AUROC: 0.6156, Partial AUROC: 0.0277  
 Epoch 18/20

Training Epoch 18: 100%| | 66/66 [01:44<00:00, 1.59s/it,  
 train\_loss=0.618]  
 Validating Epoch 18: 100%| | 47/47 [01:22<00:00, 1.76s/it,  
 val\_loss=0.552]  
  
 Epoch [17/20], Train Loss: 0.6176, Val Loss: 0.5519, Val Accuracy: 96.04%, Val  
 AUROC: 0.6092, Partial AUROC: 0.0260  
 Epoch 19/20  
  
 Training Epoch 19: 100%| | 66/66 [01:44<00:00, 1.59s/it,  
 train\_loss=0.61]  
 Validating Epoch 19: 100%| | 47/47 [01:19<00:00, 1.70s/it,  
 val\_loss=0.549]  
  
 Epoch [18/20], Train Loss: 0.6098, Val Loss: 0.5488, Val Accuracy: 96.11%, Val  
 AUROC: 0.6211, Partial AUROC: 0.0298  
 Epoch 20/20  
  
 Training Epoch 20: 100%| | 66/66 [01:59<00:00, 1.82s/it,  
 train\_loss=0.616]  
 Validating Epoch 20: 100%| | 47/47 [01:23<00:00, 1.78s/it,  
 val\_loss=0.549]  
  
 Epoch [19/20], Train Loss: 0.6163, Val Loss: 0.5489, Val Accuracy: 96.11%, Val  
 AUROC: 0.6246, Partial AUROC: 0.0306  
 Best Epoch: 19, Best Validation Loss: 0.5488  
 Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.96	1.00	0.98	1431
Class 1	1.00	0.02	0.03	59
accuracy			0.96	1490
macro avg	0.98	0.51	0.51	1490
weighted avg	0.96	0.96	0.94	1490

```
[25]: CustomImageFeatureResNet(
  (resnet): Sequential(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
    ceil_mode=False)
    (4): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
        1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(

```

```

        (0): BasicBlock(
          (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
          (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
          (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (downsample): Sequential(
            (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (1): BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    (7): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```

```

1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (8): AdaptiveAvgPool2d(output_size=(1, 1))
    )
    (fc_image): Linear(in_features=512, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
    )

```

## 0.11 Model 7

```

[26]: model7 = CustomImageFeatureEfficientNet(feature_input_size=9)  # Assuming 9
    ↪ features for metadata
model7.to(device)
# Initialize optimizer
optimizer = optim.Adam(model7.parameters(), lr= 1.1621608010269284e-05)
# Define the loss function with the class weights
criterion = nn.BCELoss()  # Binary classification loss
# Set the number of epochs
epochs = 20
batch_size = 16
best_model_path = "best_model7.pth"

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[27]: effnet_train_dataloader = DataLoader(effnet_train_dataset,
    ↪ batch_size=batch_size, shuffle=True)

```

```
effnet_val_dataloader = DataLoader(effnet_val_dataset, batch_size=batch_size,
    ↪shuffle=True)
```

```
[28]: train_and_validate(model7, effnet_train_dataloader, effnet_val_dataloader,
    ↪criterion, optimizer, epochs, device, best_model_path )
```

Epoch 1/20

Training Epoch 1: 100%| | 131/131 [01:47<00:00, 1.21it/s,  
train\_loss=0.665]

Validating Epoch 1: 100%| | 94/94 [01:39<00:00, 1.06s/it,  
val\_loss=0.624]

Epoch [0/20], Train Loss: 0.6653, Val Loss: 0.6238, Val Accuracy: 85.03%, Val  
AUROC: 0.5190, Partial AUROC: 0.0248

Epoch 2/20

Training Epoch 2: 100%| | 131/131 [01:46<00:00, 1.23it/s,  
train\_loss=0.608]

Validating Epoch 2: 100%| | 94/94 [01:15<00:00, 1.25it/s,  
val\_loss=0.561]

Epoch [1/20], Train Loss: 0.6077, Val Loss: 0.5613, Val Accuracy: 91.28%, Val  
AUROC: 0.6415, Partial AUROC: 0.0523

Epoch 3/20

Training Epoch 3: 100%| | 131/131 [02:12<00:00, 1.01s/it,  
train\_loss=0.545]

Validating Epoch 3: 100%| | 94/94 [01:15<00:00, 1.24it/s,  
val\_loss=0.526]

Epoch [2/20], Train Loss: 0.5450, Val Loss: 0.5263, Val Accuracy: 89.33%, Val  
AUROC: 0.7150, Partial AUROC: 0.0620

Epoch 4/20

Training Epoch 4: 100%| | 131/131 [01:50<00:00, 1.19it/s,  
train\_loss=0.496]

Validating Epoch 4: 100%| | 94/94 [01:14<00:00, 1.27it/s,  
val\_loss=0.448]

Epoch [3/20], Train Loss: 0.4959, Val Loss: 0.4483, Val Accuracy: 90.40%, Val  
AUROC: 0.7604, Partial AUROC: 0.0666

Epoch 5/20

Training Epoch 5: 100%| | 131/131 [01:44<00:00, 1.26it/s,  
train\_loss=0.447]

Validating Epoch 5: 100%| | 94/94 [01:13<00:00, 1.28it/s,  
val\_loss=0.414]

Epoch [4/20], Train Loss: 0.4469, Val Loss: 0.4141, Val Accuracy: 89.26%, Val  
AUROC: 0.7891, Partial AUROC: 0.0740

Epoch 6/20

Training Epoch 6: 100%| | 131/131 [01:46<00:00, 1.23it/s,  
 train\_loss=0.413]  
 Validating Epoch 6: 100%| | 94/94 [01:17<00:00, 1.21it/s,  
 val\_loss=0.44]  
  
 Epoch [5/20], Train Loss: 0.4128, Val Loss: 0.4404, Val Accuracy: 84.70%, Val  
 AUROC: 0.8196, Partial AUROC: 0.0824  
 Epoch 7/20  
  
 Training Epoch 7: 100%| | 131/131 [01:45<00:00, 1.24it/s,  
 train\_loss=0.373]  
 Validating Epoch 7: 100%| | 94/94 [00:52<00:00, 1.79it/s,  
 val\_loss=0.412]  
  
 Epoch [6/20], Train Loss: 0.3726, Val Loss: 0.4124, Val Accuracy: 85.84%, Val  
 AUROC: 0.8262, Partial AUROC: 0.0850  
 Epoch 8/20  
  
 Training Epoch 8: 100%| | 131/131 [01:43<00:00, 1.27it/s,  
 train\_loss=0.365]  
 Validating Epoch 8: 100%| | 94/94 [01:13<00:00, 1.27it/s,  
 val\_loss=0.368]  
  
 Epoch [7/20], Train Loss: 0.3649, Val Loss: 0.3682, Val Accuracy: 87.38%, Val  
 AUROC: 0.8312, Partial AUROC: 0.0893  
 Epoch 9/20  
  
 Training Epoch 9: 100%| | 131/131 [01:40<00:00, 1.30it/s,  
 train\_loss=0.334]  
 Validating Epoch 9: 100%| | 94/94 [01:14<00:00, 1.25it/s,  
 val\_loss=0.342]  
  
 Epoch [8/20], Train Loss: 0.3337, Val Loss: 0.3420, Val Accuracy: 87.65%, Val  
 AUROC: 0.8349, Partial AUROC: 0.0849  
 Epoch 10/20  
  
 Training Epoch 10: 100%| | 131/131 [01:47<00:00, 1.22it/s,  
 train\_loss=0.307]  
 Validating Epoch 10: 100%| | 94/94 [01:04<00:00, 1.47it/s,  
 val\_loss=0.356]  
  
 Epoch [9/20], Train Loss: 0.3074, Val Loss: 0.3557, Val Accuracy: 85.17%, Val  
 AUROC: 0.8424, Partial AUROC: 0.0945  
 Epoch 11/20  
  
 Training Epoch 11: 100%| | 131/131 [01:42<00:00, 1.27it/s,  
 train\_loss=0.312]  
 Validating Epoch 11: 100%| | 94/94 [01:17<00:00, 1.21it/s,  
 val\_loss=0.339]  
  
 Epoch [10/20], Train Loss: 0.3117, Val Loss: 0.3386, Val Accuracy: 86.38%, Val  
 AUROC: 0.8504, Partial AUROC: 0.0984  
 Epoch 12/20



Training Epoch 12: 100%| | 131/131 [01:47<00:00, 1.22it/s,  
 train\_loss=0.288]  
 Validating Epoch 12: 100%| | 94/94 [01:00<00:00, 1.54it/s,  
 val\_loss=0.348]  
  
 Epoch [11/20], Train Loss: 0.2876, Val Loss: 0.3480, Val Accuracy: 84.77%, Val  
 AUROC: 0.8519, Partial AUROC: 0.1023  
 Epoch 13/20  
  
 Training Epoch 13: 100%| | 131/131 [02:07<00:00, 1.03it/s,  
 train\_loss=0.26]  
 Validating Epoch 13: 100%| | 94/94 [01:12<00:00, 1.29it/s,  
 val\_loss=0.312]  
  
 Epoch [12/20], Train Loss: 0.2600, Val Loss: 0.3124, Val Accuracy: 85.64%, Val  
 AUROC: 0.8514, Partial AUROC: 0.0994  
 Epoch 14/20  
  
 Training Epoch 14: 100%| | 131/131 [02:17<00:00, 1.05s/it,  
 train\_loss=0.257]  
 Validating Epoch 14: 100%| | 94/94 [00:53<00:00, 1.75it/s,  
 val\_loss=0.283]  
  
 Epoch [13/20], Train Loss: 0.2573, Val Loss: 0.2826, Val Accuracy: 89.19%, Val  
 AUROC: 0.8584, Partial AUROC: 0.1038  
 Epoch 15/20  
  
 Training Epoch 15: 100%| | 131/131 [01:52<00:00, 1.17it/s,  
 train\_loss=0.258]  
 Validating Epoch 15: 100%| | 94/94 [01:06<00:00, 1.42it/s,  
 val\_loss=0.274]  
  
 Epoch [14/20], Train Loss: 0.2578, Val Loss: 0.2743, Val Accuracy: 88.99%, Val  
 AUROC: 0.8578, Partial AUROC: 0.1049  
 Epoch 16/20  
  
 Training Epoch 16: 100%| | 131/131 [01:47<00:00, 1.21it/s,  
 train\_loss=0.241]  
 Validating Epoch 16: 100%| | 94/94 [01:12<00:00, 1.29it/s,  
 val\_loss=0.267]  
  
 Epoch [15/20], Train Loss: 0.2406, Val Loss: 0.2668, Val Accuracy: 89.06%, Val  
 AUROC: 0.8513, Partial AUROC: 0.1010  
 Epoch 17/20  
  
 Training Epoch 17: 100%| | 131/131 [01:45<00:00, 1.24it/s,  
 train\_loss=0.229]  
 Validating Epoch 17: 100%| | 94/94 [01:02<00:00, 1.51it/s,  
 val\_loss=0.277]  
  
 Epoch [16/20], Train Loss: 0.2286, Val Loss: 0.2768, Val Accuracy: 88.46%, Val  
 AUROC: 0.8439, Partial AUROC: 0.0954  
 Epoch 18/20

Training Epoch 18: 100%| | 131/131 [01:46<00:00, 1.23it/s,  
train\_loss=0.231]

Validating Epoch 18: 100%| | 94/94 [00:58<00:00, 1.61it/s,  
val\_loss=0.241]

Epoch [17/20], Train Loss: 0.2313, Val Loss: 0.2406, Val Accuracy: 90.74%, Val  
AUROC: 0.8417, Partial AUROC: 0.0939

Epoch 19/20

Training Epoch 19: 100%| | 131/131 [01:42<00:00, 1.28it/s,  
train\_loss=0.211]

Validating Epoch 19: 100%| | 94/94 [01:04<00:00, 1.45it/s,  
val\_loss=0.217]

Epoch [18/20], Train Loss: 0.2109, Val Loss: 0.2166, Val Accuracy: 92.08%, Val  
AUROC: 0.8535, Partial AUROC: 0.1034

Epoch 20/20


Training Epoch 20: 100%| | 131/131 [01:47<00:00, 1.22it/s,  
train\_loss=0.204]


Validating Epoch 20: 100%| | 94/94 [00:51<00:00, 1.84it/s,  
val\_loss=0.281]

Epoch [19/20], Train Loss: 0.2044, Val Loss: 0.2810, Val Accuracy: 87.32%, Val  
AUROC: 0.8402, Partial AUROC: 0.0970

Best Epoch: 19, Best Validation Loss: 0.2166

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.88	0.93	1431
Class 1	0.18	0.64	0.29	59
accuracy			0.87	1490
macro avg	0.58	0.76	0.61	1490
weighted avg	0.95	0.87	0.90	1490

```
[28]: CustomImageFeatureEfficientNet(
  (efficientnet): Sequential(
    (0): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Sequential(
        (0): MBConv(
          (block): Sequential(
            (0): Conv2dNormActivation(
              (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
```

```

        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (2): Conv2dNormActivation(
        (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.0, mode=row)
    )
    )
    (2): Sequential(
        (0): MBConv(
            (block): Sequential(
                (0): Conv2dNormActivation(
                    (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(
                (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
                (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.025, mode=row)
)
(3): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )

```

```

        (1): Conv2dNormActivation(
          (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
          (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
  )
  (1): MBCConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(

```

```

        (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
)
(4): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
  )
  (1): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
)
(2): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))

```



```

        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
)
)
(5): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
  )
  (1): MBConv(

```

```

        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
      )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
      )
    )

```

```

        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.125, mode=row)
    )
  )
  (6): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )

```

```

    )
    (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)

```

```

        (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
)
(3): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
    )
    )
    (7): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
track_running_stats=True)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      )
      (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
    )
    )
    (8): Conv2dNormActivation(
      (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
  )
)

```

```

        (1): AdaptiveAvgPool2d(output_size=1)
    )
    (fc_image): Linear(in_features=1280, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

## 0.12 Model 8

```

[29]: model8 = CustomImageFeatureEfficientNet(feature_input_size=9) # Assuming 9
      ↪ features for metadata
model8.to(device)
# Initialize optimizer
optimizer = optim.SGD(model8.parameters(), lr=0.01)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs
epochs = 20
best_model_path = "best_model8.pth"

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)

```

```

[31]: train_and_validate(model8, effnet_train_dataloader, effnet_val_dataloader,
      ↪ criterion, optimizer, epochs, device, best_model_path )

```

Epoch 1/20

```

Training Epoch 1: 100%|          | 131/131 [01:39<00:00,  1.31it/s,
train_loss=0.616]
Validating Epoch 1: 100%|          | 94/94 [00:51<00:00,  1.82it/s,
val_loss=0.46]

```

```

Epoch [0/20], Train Loss: 0.6159, Val Loss: 0.4601, Val Accuracy: 95.37%, Val
AUROC: 0.6985, Partial AUROC: 0.0364

```

Epoch 2/20

Training Epoch 2: 100%| | 131/131 [01:43<00:00, 1.27it/s,  
 train\_loss=0.551]  
 Validating Epoch 2: 100%| | 94/94 [00:50<00:00, 1.85it/s,  
 val\_loss=0.417]  
 Epoch [1/20], Train Loss: 0.5513, Val Loss: 0.4175, Val Accuracy: 93.83%, Val  
 AUROC: 0.7256, Partial AUROC: 0.0514  
 Epoch 3/20  
 Training Epoch 3: 100%| | 131/131 [01:36<00:00, 1.35it/s,  
 train\_loss=0.494]  
 Validating Epoch 3: 100%| | 94/94 [00:50<00:00, 1.85it/s,  
 val\_loss=0.441]  
 Epoch [2/20], Train Loss: 0.4937, Val Loss: 0.4410, Val Accuracy: 90.00%, Val  
 AUROC: 0.7560, Partial AUROC: 0.0557  
 Epoch 4/20  
 Training Epoch 4: 100%| | 131/131 [01:37<00:00, 1.35it/s,  
 train\_loss=0.444]  
 Validating Epoch 4: 100%| | 94/94 [00:52<00:00, 1.80it/s,  
 val\_loss=0.418]  
 Epoch [3/20], Train Loss: 0.4435, Val Loss: 0.4181, Val Accuracy: 88.86%, Val  
 AUROC: 0.8098, Partial AUROC: 0.0727  
 Epoch 5/20  
 Training Epoch 5: 100%| | 131/131 [01:37<00:00, 1.34it/s,  
 train\_loss=0.395]  
 Validating Epoch 5: 100%| | 94/94 [00:52<00:00, 1.79it/s,  
 val\_loss=0.411]  
 Epoch [4/20], Train Loss: 0.3950, Val Loss: 0.4110, Val Accuracy: 88.32%, Val  
 AUROC: 0.8227, Partial AUROC: 0.0769  
 Epoch 6/20  
 Training Epoch 6: 100%| | 131/131 [01:46<00:00, 1.23it/s,  
 train\_loss=0.339]  
 Validating Epoch 6: 100%| | 94/94 [00:51<00:00, 1.83it/s,  
 val\_loss=0.402]  
 Epoch [5/20], Train Loss: 0.3392, Val Loss: 0.4024, Val Accuracy: 84.16%, Val  
 AUROC: 0.8389, Partial AUROC: 0.0866  
 Epoch 7/20  
 Training Epoch 7: 100%| | 131/131 [01:40<00:00, 1.30it/s,  
 train\_loss=0.303]  
 Validating Epoch 7: 100%| | 94/94 [00:52<00:00, 1.80it/s,  
 val\_loss=0.391]  
 Epoch [6/20], Train Loss: 0.3032, Val Loss: 0.3913, Val Accuracy: 84.36%, Val  
 AUROC: 0.8484, Partial AUROC: 0.0918  
 Epoch 8/20





Training Epoch 8: 100%| | 131/131 [01:39<00:00, 1.32it/s,  
 train\_loss=0.293]  
 Validating Epoch 8: 100%| | 94/94 [00:52<00:00, 1.78it/s,  
 val\_loss=0.333]  
 Epoch [7/20], Train Loss: 0.2931, Val Loss: 0.3327, Val Accuracy: 87.99%, Val  
 AUROC: 0.8634, Partial AUROC: 0.1028  
 Epoch 9/20  
 Training Epoch 9: 100%| | 131/131 [01:41<00:00, 1.29it/s,  
 train\_loss=0.261]  
 Validating Epoch 9: 100%| | 94/94 [00:54<00:00, 1.73it/s,  
 val\_loss=0.308]  
 Epoch [8/20], Train Loss: 0.2610, Val Loss: 0.3081, Val Accuracy: 88.26%, Val  
 AUROC: 0.8684, Partial AUROC: 0.1020  
 Epoch 10/20  
 Training Epoch 10: 100%| | 131/131 [01:46<00:00, 1.23it/s,  
 train\_loss=0.243]  
 Validating Epoch 10: 100%| | 94/94 [00:51<00:00, 1.81it/s,  
 val\_loss=0.283]  
 Epoch [9/20], Train Loss: 0.2428, Val Loss: 0.2835, Val Accuracy: 89.66%, Val  
 AUROC: 0.8852, Partial AUROC: 0.1167  
 Epoch 11/20  
 Training Epoch 11: 100%| | 131/131 [01:38<00:00, 1.32it/s,  
 train\_loss=0.23]  
 Validating Epoch 11: 100%| | 94/94 [00:51<00:00, 1.84it/s,  
 val\_loss=0.303]  
 Epoch [10/20], Train Loss: 0.2297, Val Loss: 0.3025, Val Accuracy: 88.05%, Val  
 AUROC: 0.8682, Partial AUROC: 0.1011  
 Epoch 12/20  
 Training Epoch 12: 100%| | 131/131 [01:39<00:00, 1.32it/s,  
 train\_loss=0.194]  
 Validating Epoch 12: 100%| | 94/94 [01:00<00:00, 1.57it/s,  
 val\_loss=0.228]  
 Epoch [11/20], Train Loss: 0.1944, Val Loss: 0.2279, Val Accuracy: 91.28%, Val  
 AUROC: 0.8819, Partial AUROC: 0.1166  
 Epoch 13/20  
 Training Epoch 13: 100%| | 131/131 [01:37<00:00, 1.34it/s,  
 train\_loss=0.168]  
 Validating Epoch 13: 100%| | 94/94 [00:55<00:00, 1.69it/s,  
 val\_loss=0.269]  
 Epoch [12/20], Train Loss: 0.1682, Val Loss: 0.2695, Val Accuracy: 89.53%, Val  
 AUROC: 0.8395, Partial AUROC: 0.0838  
 Epoch 14/20

Training Epoch 14: 100%| | 131/131 [01:38<00:00, 1.33it/s,  
 train\_loss=0.182]  
 Validating Epoch 14: 100%| | 94/94 [00:51<00:00, 1.81it/s,  
 val\_loss=0.232]  
  
 Epoch [13/20], Train Loss: 0.1816, Val Loss: 0.2322, Val Accuracy: 91.01%, Val  
 AUROC: 0.8722, Partial AUROC: 0.1132  
 Epoch 15/20  
  
 Training Epoch 15: 100%| | 131/131 [01:38<00:00, 1.33it/s,  
 train\_loss=0.157]  
 Validating Epoch 15: 100%| | 94/94 [00:50<00:00, 1.84it/s,  
 val\_loss=0.219]  
  
 Epoch [14/20], Train Loss: 0.1574, Val Loss: 0.2194, Val Accuracy: 91.61%, Val  
 AUROC: 0.8603, Partial AUROC: 0.1037  
 Epoch 16/20  
  
 Training Epoch 16: 100%| | 131/131 [01:37<00:00, 1.34it/s,  
 train\_loss=0.131]  
 Validating Epoch 16: 100%| | 94/94 [00:51<00:00, 1.83it/s,  
 val\_loss=0.227]  
  
 Epoch [15/20], Train Loss: 0.1313, Val Loss: 0.2267, Val Accuracy: 90.74%, Val  
 AUROC: 0.8573, Partial AUROC: 0.1118  
 Epoch 17/20  
  
 Training Epoch 17: 100%| | 131/131 [01:38<00:00, 1.33it/s,  
 train\_loss=0.137]  
 Validating Epoch 17: 100%| | 94/94 [00:56<00:00, 1.66it/s,  
 val\_loss=0.21]  
  
 Epoch [16/20], Train Loss: 0.1373, Val Loss: 0.2097, Val Accuracy: 92.01%, Val  
 AUROC: 0.8591, Partial AUROC: 0.1083  
 Epoch 18/20  
  
 Training Epoch 18: 100%| | 131/131 [01:39<00:00, 1.32it/s,  
 train\_loss=0.12]  
 Validating Epoch 18: 100%| | 94/94 [00:51<00:00, 1.82it/s,  
 val\_loss=0.274]  
  
 Epoch [17/20], Train Loss: 0.1205, Val Loss: 0.2736, Val Accuracy: 89.40%, Val  
 AUROC: 0.8619, Partial AUROC: 0.1053  
 Epoch 19/20  
  
 Training Epoch 19: 100%| | 131/131 [01:37<00:00, 1.34it/s,  
 train\_loss=0.12]  
 Validating Epoch 19: 100%| | 94/94 [00:50<00:00, 1.85it/s,  
 val\_loss=0.23]  
  
 Epoch [18/20], Train Loss: 0.1195, Val Loss: 0.2301, Val Accuracy: 90.87%, Val  
 AUROC: 0.8657, Partial AUROC: 0.1158  
 Epoch 20/20

Training Epoch 20: 100%| | 131/131 [01:37<00:00, 1.34it/s,  
train\_loss=0.102]  
Validating Epoch 20: 100%| | 94/94 [00:51<00:00, 1.82it/s,  
val\_loss=0.214]

Epoch [19/20], Train Loss: 0.1019, Val Loss: 0.2142, Val Accuracy: 91.48%, Val  
AUROC: 0.8567, Partial AUROC: 0.1095  
Best Epoch: 17, Best Validation Loss: 0.2097  
Training Complete



Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.93	0.95	1431
Class 1	0.25	0.56	0.34	59
accuracy			0.91	1490
macro avg	0.61	0.74	0.65	1490
weighted avg	0.95	0.91	0.93	1490

```
[31]: CustomImageFeatureEfficientNet(
      (efficientnet): Sequential(
        (0): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Sequential(
            (0): MBConv(
              (block): Sequential(
                (0): Conv2dNormActivation(
                  (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=32, bias=False)
                  (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                  (2): SiLU(inplace=True)
                )
                (1): SqueezeExcitation(
                  (avgpool): AdaptiveAvgPool2d(output_size=1)
                  (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
                  (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
                  (activation): SiLU(inplace=True)
                  (scale_activation): Sigmoid()
                )
                (2): Conv2dNormActivation(
                  (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
                  (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
              )
            )
          )
          (stochastic_depth): StochasticDepth(p=0.0, mode=row)
```

```

    )
)
(2): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
        (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
  (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.025, mode=row)
)
)
(3): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
    )
    (1): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
    )
    (4): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)

```

```

    )
    (1): Conv2dNormActivation(
      (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), groups=240, bias=False)
      (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
      (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
  )
)

```



```

        (3): Conv2dNormActivation(
          (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
    )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
          (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
    )
  )
  (5): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),

```

```

bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
)
    (1): Conv2dNormActivation(
    (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
    (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
)
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
    (3): Conv2dNormActivation(
    (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
)
    (1): MBConv(
    (block): Sequential(
    (0): Conv2dNormActivation(
    (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
    (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
    (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))

```

```

        (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    )
    (stochastic_depth): StochasticDepth(p=0.125, mode=row)
)
(6): Sequential(

```

```

(0): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
      (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
  )
)

```

```

    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
)

```

```

        (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
    )
    (3): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
          (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
    (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
  )
  (7): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),

```

```

padding=(1, 1), groups=1152, bias=False)
    (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
)
    (2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
    (3): Conv2dNormActivation(
    (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
    )
    )
    (8): Conv2dNormActivation(
    (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
    )
    )
    (1): AdaptiveAvgPool2d(output_size=1)
)
(fc_image): Linear(in_features=1280, out_features=512, bias=True)
(fc_metadata): Linear(in_features=9, out_features=128, bias=True)
(dropout): Dropout(p=0.5, inplace=False)
(fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

### 0.13 Model 9

```

[32]: model9 = CustomImageFeatureEfficientNet(feature_input_size=9) # Assuming 9
    ↪ features for metadata
model9.to(device)
# Initialize optimizer
optimizer = optim.Adam(model9.parameters(), lr=0.001)
# Define the loss function with the class weights
criterion = nn.BCELoss() # Binary classification loss
# Set the number of epochs

```

```
epochs = 20
batch_sizes = 16
best_model_path = "best_model9.path"
```

```
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
    warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
    warnings.warn(msg)
```

```
[33]: train_and_validate(model9, effnet_train_dataloader, effnet_val_dataloader,
    ↪ criterion, optimizer, epochs, device, best_model_path )
```

Epoch 1/20

```
Training Epoch 1: 100%|      | 131/131 [01:40<00:00,  1.30it/s,
train_loss=0.432]
Validating Epoch 1: 100%|      | 94/94 [00:56<00:00,  1.68it/s,
val_loss=0.376]
```

Epoch [0/20], Train Loss: 0.4325, Val Loss: 0.3763, Val Accuracy: 83.69%, Val AUROC: 0.8220, Partial AUROC: 0.0917

Epoch 2/20

```
Training Epoch 2: 100%|      | 131/131 [01:40<00:00,  1.31it/s,
train_loss=0.347]
Validating Epoch 2: 100%|      | 94/94 [00:50<00:00,  1.87it/s,
val_loss=0.42]
```

Epoch [1/20], Train Loss: 0.3475, Val Loss: 0.4196, Val Accuracy: 74.50%, Val AUROC: 0.8406, Partial AUROC: 0.0972

Epoch 3/20

```
Training Epoch 3: 100%|      | 131/131 [01:40<00:00,  1.30it/s,
train_loss=0.319]
Validating Epoch 3: 100%|      | 94/94 [00:58<00:00,  1.60it/s,
val_loss=0.392]
```

Epoch [2/20], Train Loss: 0.3192, Val Loss: 0.3923, Val Accuracy: 88.93%, Val AUROC: 0.7979, Partial AUROC: 0.0693

Epoch 4/20

```
Training Epoch 4: 100%|      | 131/131 [01:39<00:00,  1.31it/s,
train_loss=0.322]
```



Validating Epoch 4: 100%| | 94/94 [01:00<00:00, 1.56it/s,  
val\_loss=0.389]

Epoch [3/20], Train Loss: 0.3220, Val Loss: 0.3888, Val Accuracy: 88.79%, Val  
AUROC: 0.8398, Partial AUROC: 0.0894

Epoch 5/20

Training Epoch 5: 100%| | 131/131 [01:43<00:00, 1.26it/s,  
train\_loss=0.24]

Validating Epoch 5: 100%| | 94/94 [00:55<00:00, 1.69it/s,  
val\_loss=0.26]

Epoch [4/20], Train Loss: 0.2397, Val Loss: 0.2596, Val Accuracy: 89.60%, Val  
AUROC: 0.8398, Partial AUROC: 0.0857

Epoch 6/20

Training Epoch 6: 100%| | 131/131 [01:41<00:00, 1.29it/s,  
train\_loss=0.194]

Validating Epoch 6: 100%| | 94/94 [00:51<00:00, 1.82it/s,  
val\_loss=0.242]

Epoch [5/20], Train Loss: 0.1936, Val Loss: 0.2420, Val Accuracy: 89.73%, Val  
AUROC: 0.8615, Partial AUROC: 0.1131

Epoch 7/20

Training Epoch 7: 100%| | 131/131 [01:40<00:00, 1.30it/s,  
train\_loss=0.208]

Validating Epoch 7: 100%| | 94/94 [00:51<00:00, 1.81it/s,  
val\_loss=0.173]

Epoch [6/20], Train Loss: 0.2082, Val Loss: 0.1731, Val Accuracy: 94.30%, Val  
AUROC: 0.8511, Partial AUROC: 0.0976

Epoch 8/20

Training Epoch 8: 100%| | 131/131 [01:42<00:00, 1.28it/s,  
train\_loss=0.157]

Validating Epoch 8: 100%| | 94/94 [00:51<00:00, 1.83it/s,  
val\_loss=0.329]

Epoch [7/20], Train Loss: 0.1571, Val Loss: 0.3290, Val Accuracy: 83.56%, Val  
AUROC: 0.8171, Partial AUROC: 0.0776

Epoch 9/20

Training Epoch 9: 100%| | 131/131 [01:47<00:00, 1.22it/s,  
train\_loss=0.144]

Validating Epoch 9: 100%| | 94/94 [00:50<00:00, 1.86it/s,  
val\_loss=0.413]

Epoch [8/20], Train Loss: 0.1441, Val Loss: 0.4128, Val Accuracy: 80.40%, Val  
AUROC: 0.8423, Partial AUROC: 0.0984

Epoch 10/20

Training Epoch 10: 100%| | 131/131 [01:40<00:00, 1.30it/s,  
train\_loss=0.144]

Validating Epoch 10: 100%| | 94/94 [00:52<00:00, 1.79it/s,  
val\_loss=0.422]

Epoch [9/20], Train Loss: 0.1439, Val Loss: 0.4219, Val Accuracy: 79.46%, Val  
AUROC: 0.8541, Partial AUROC: 0.1038

Epoch 11/20

Training Epoch 11: 100%| | 131/131 [01:50<00:00, 1.19it/s,  
train\_loss=0.105]

Validating Epoch 11: 100%| | 94/94 [00:54<00:00, 1.73it/s,  
val\_loss=0.451]

Epoch [10/20], Train Loss: 0.1050, Val Loss: 0.4508, Val Accuracy: 90.27%, Val  
AUROC: 0.8261, Partial AUROC: 0.1017

Epoch 12/20

Training Epoch 12: 100%| | 131/131 [01:45<00:00, 1.24it/s,  
train\_loss=0.109]


Validating Epoch 12: 100%| | 94/94 [00:52<00:00, 1.80it/s,  
val\_loss=0.211]


Epoch [11/20], Train Loss: 0.1091, Val Loss: 0.2113, Val Accuracy: 91.34%, Val  
AUROC: 0.8501, Partial AUROC: 0.1038

Early stopping triggered at epoch 11

Best Epoch: 7, Best Validation Loss: 0.1731

Training Complete





Classification Report:

	precision	recall	f1-score	support
Class 0	0.98	0.93	0.95	1431
Class 1	0.22	0.46	0.30	59
accuracy			0.91	1490
macro avg	0.60	0.69	0.62	1490
weighted avg	0.95	0.91	0.93	1490

```
[33]: CustomImageFeatureEfficientNet(
      (efficientnet): Sequential(
        (0): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
            bias=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
            track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Sequential(
            (0): MBConv(
              (block): Sequential(
                (0): Conv2dNormActivation(
                  (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
                  1), groups=32, bias=False)
```

```

        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (2): Conv2dNormActivation(
        (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.0, mode=row)
    )
    )
    (2): Sequential(
        (0): MBConv(
            (block): Sequential(
                (0): Conv2dNormActivation(
                    (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                )
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(
                (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
                (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.0125, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=144, bias=False)
      (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.025, mode=row)
)
(3): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
    )
  )

```

```

        (1): Conv2dNormActivation(
          (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=144, bias=False)
          (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.037500000000000006, mode=row)
    )
    (1): MBCConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=240, bias=False)
          (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(

```

```

        (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.05, mode=row)
)
(4): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (stochastic_depth): StochasticDepth(p=0.0625, mode=row)
  )
  (1): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.07500000000000001, mode=row)
)
(2): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))

```



```

        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.08750000000000001, mode=row)
)
)
(5): Sequential(
  (0): MBConv(
    (block): Sequential(
      (0): Conv2dNormActivation(
        (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (1): Conv2dNormActivation(
        (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
      )
      (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
      )
      (3): Conv2dNormActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1, mode=row)
  )
  (1): MBConv(

```

```

        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
        (stochastic_depth): StochasticDepth(p=0.1125, mode=row)
      )
    (2): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
      )
    )
  )
)

```

```

        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (stochastic_depth): StochasticDepth(p=0.125, mode=row)
    )
  )
  (6): Sequential(
    (0): MBConv(
      (block): Sequential(
        (0): Conv2dNormActivation(
          (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
          (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
          (avgpool): AdaptiveAvgPool2d(output_size=1)
          (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
          (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
          (activation): SiLU(inplace=True)
          (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
          (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )

```

```

    )
    (stochastic_depth): StochasticDepth(p=0.1375, mode=row)
)
(1): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (avgpool): AdaptiveAvgPool2d(output_size=1)
      (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
      (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
      (activation): SiLU(inplace=True)
      (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
      (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (stochastic_depth): StochasticDepth(p=0.15000000000000002, mode=row)
)
(2): MBConv(
  (block): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Conv2dNormActivation(
      (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)

```

```

        (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.1625, mode=row)
)
(3): MBCConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    )
    (stochastic_depth): StochasticDepth(p=0.17500000000000002, mode=row)
    )
    )
    (7): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
track_running_stats=True)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
          )
          (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
          )
          (3): Conv2dNormActivation(
            (0): Conv2d(1152, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      )
      (stochastic_depth): StochasticDepth(p=0.1875, mode=row)
    )
    )
    (8): Conv2dNormActivation(
      (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): SiLU(inplace=True)
    )
  )
)

```

```

        (1): AdaptiveAvgPool2d(output_size=1)
    )
    (fc_image): Linear(in_features=1280, out_features=512, bias=True)
    (fc_metadata): Linear(in_features=9, out_features=128, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (fc_combined): Linear(in_features=640, out_features=1, bias=True)
)

```

## 0.14 Winning Model

```

[9]: #create test dataset
effnet_test_dataset = MultiInputDataset(hdf5_file='../data/raw/test_image.
    ↪hdf5', csv_file='../data/processed/processed-test-metadata1.csv',
    ↪transform=get_normal_transform(resize_size=(128,128)))
# Create test DataLoader
effnet_test_dataloader = DataLoader(effnet_test_dataset, batch_size=64,
    ↪shuffle=True)

```

```

[10]: # Initialize the final model with 9 output features
final_model = CustomImageFeatureEfficientNet(9)

# Define the path to the saved model weights
final_model_path = "best_model7.pth"

# Load the model weights into final_model, mapping to CPU if necessary
final_model.load_state_dict(torch.load(final_model_path, map_location=torch.
    ↪device('cpu'))))

# Set the model to evaluation mode, which disables dropout and batch
    ↪normalization updates
final_model.eval()

# Initialize lists to store labels and predicted probabilities for later
    ↪analysis
all_labels, all_probs = [], []

# Disable gradient computation for testing phase to save memory and improve
    ↪performance
with torch.no_grad():
    # Loop through batches in the test dataloader
    for images, metadata, labels in effnet_test_dataloader:
        # Move data to the specified device (e.g., CPU or GPU) and adjust label
        ↪shape
        images, metadata = images.to(device), metadata.to(device)
        labels = labels.float().to(device).unsqueeze(1)

```

```

# Forward pass to get probabilities from the model
probs = final_model(images, metadata)

# Collect labels and predicted probabilities, converting to numpy arrays
all_labels.extend(labels.cpu().numpy())
all_probs.extend(probs.cpu().numpy())

# Generate binary predictions based on a 0.5 threshold
predicted = (probs > 0.5).float()

# Calculate the partial AUROC score (adjusted for your specific function)
partial_auroc = score(np.array(all_labels), np.array(all_probs))
print(f'The partial AUROC of the final model on the test images is_
↳{partial_auroc}')

# Print the classification report, evaluating performance on Class 0 and_
↳Class 1
print(classification_report(all_labels, (np.array(all_probs) >= 0.5).
↳astype(int), target_names=['Class 0', 'Class 1']))

```

```

/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.
  warnings.warn(
/home/jupyter-sohka/.local/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=EfficientNet_B0_Weights.IMAGENET1K_V1`. You can also use
`weights=EfficientNet_B0_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
/tmp/ipykernel_1189445/1883220575.py:3: FutureWarning: You are using
`torch.load` with `weights_only=False` (the current default value), which uses
the default pickle module implicitly. It is possible to construct malicious
pickle data which will execute arbitrary code during unpickling (See
https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models for
more details). In a future release, the default value for `weights_only` will be
flipped to `True`. This limits the functions that could be executed during
unpickling. Arbitrary objects will no longer be allowed to be loaded via this
mode unless they are explicitly allowlisted by the user via
`torch.serialization.add_safe_globals`. We recommend you start setting
`weights_only=True` for any use case where you don't have full control of the

```



loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```
final_model.load_state_dict(torch.load(final_model_path,  
map_location=torch.device('cpu')))
```

The partial auROC of the final model on the test image is 0.13625176183538829

	precision	recall	f1-score	support
Class 0	0.98	0.92	0.95	1431
Class 1	0.24	0.61	0.35	59
accuracy			0.91	1490
macro avg	0.61	0.77	0.65	1490
weighted avg	0.95	0.91	0.93	1490